

# Build with AI

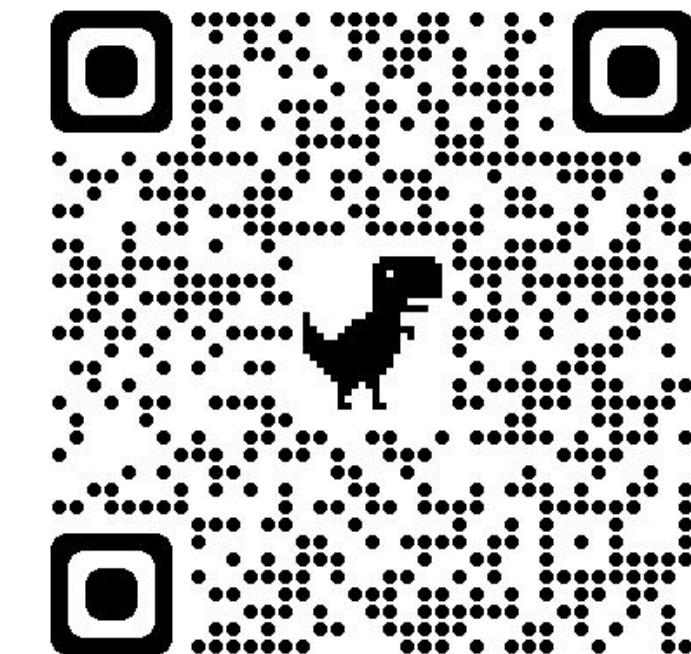
 Google Developer Groups

 Google Developer Experts

## Multimodality with Gemini

Henry Ruiz  
Research Scientist  
Texas A&M University AgriLife Research  
GDE ML  
[@devharuiz](https://devharuiz.github.io)  
[https://devharuiz.github.io/](https://devharuiz.github.io)

Gemini



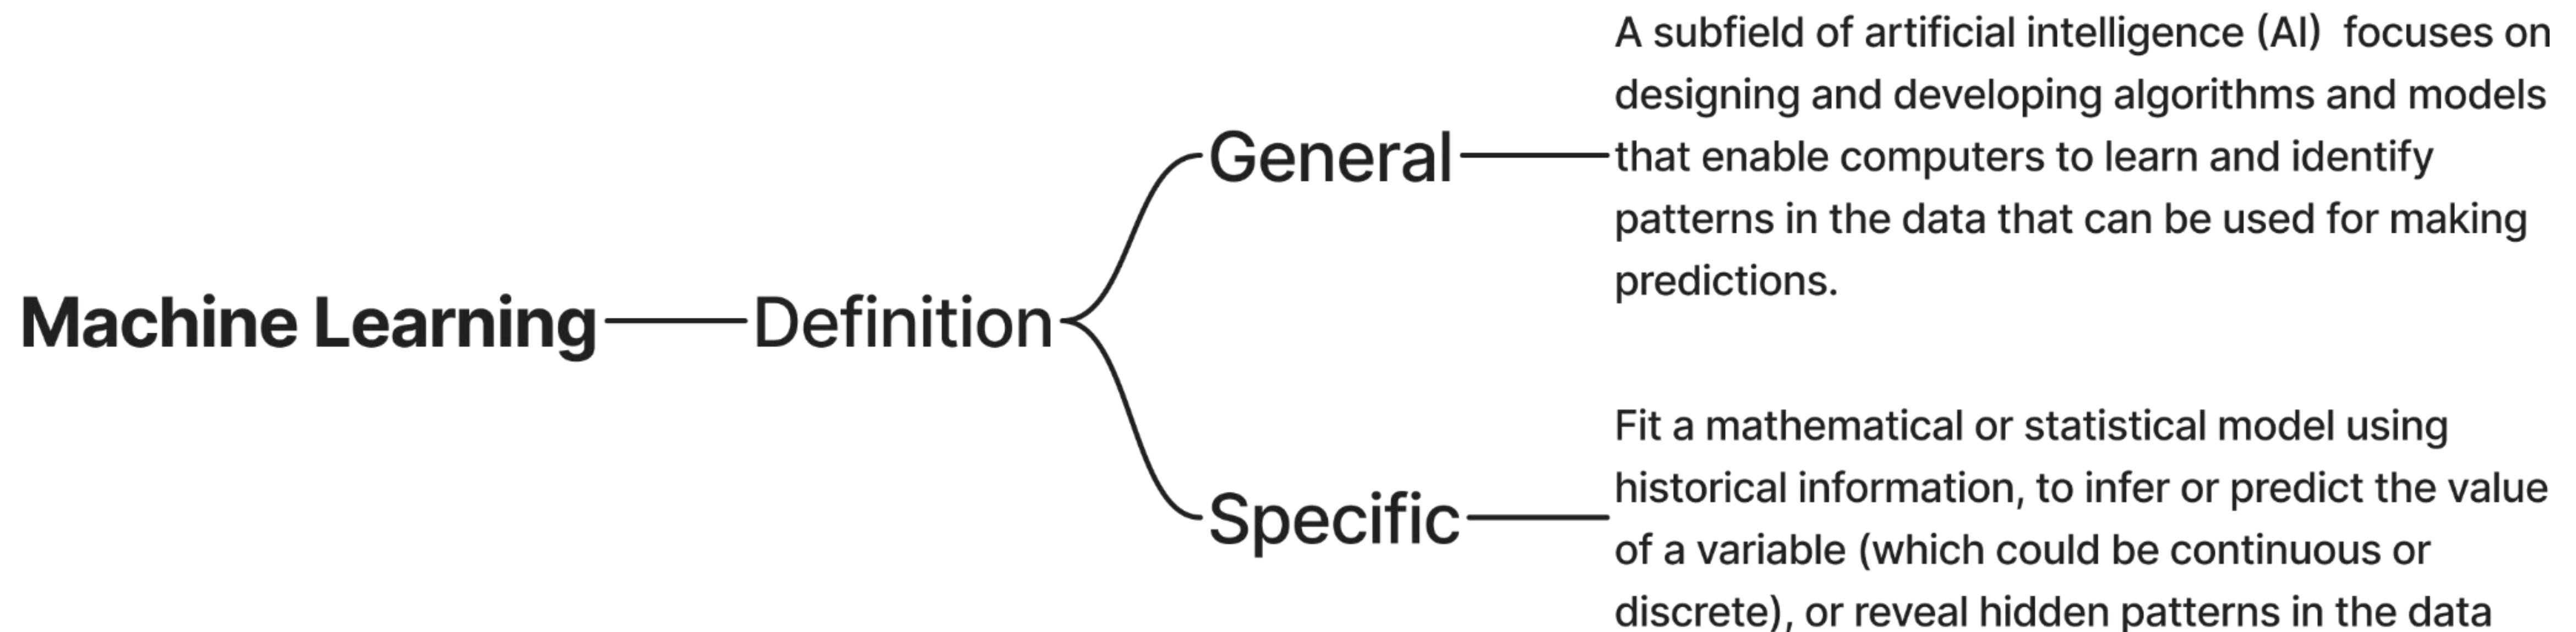
# Table of Contents

1. ML short Intro
2. Generative AI
3. LLMs(Large Language Models)
4. Gemini
5. Models variants and capabilities
6. Multimodality with Gemini Notebook

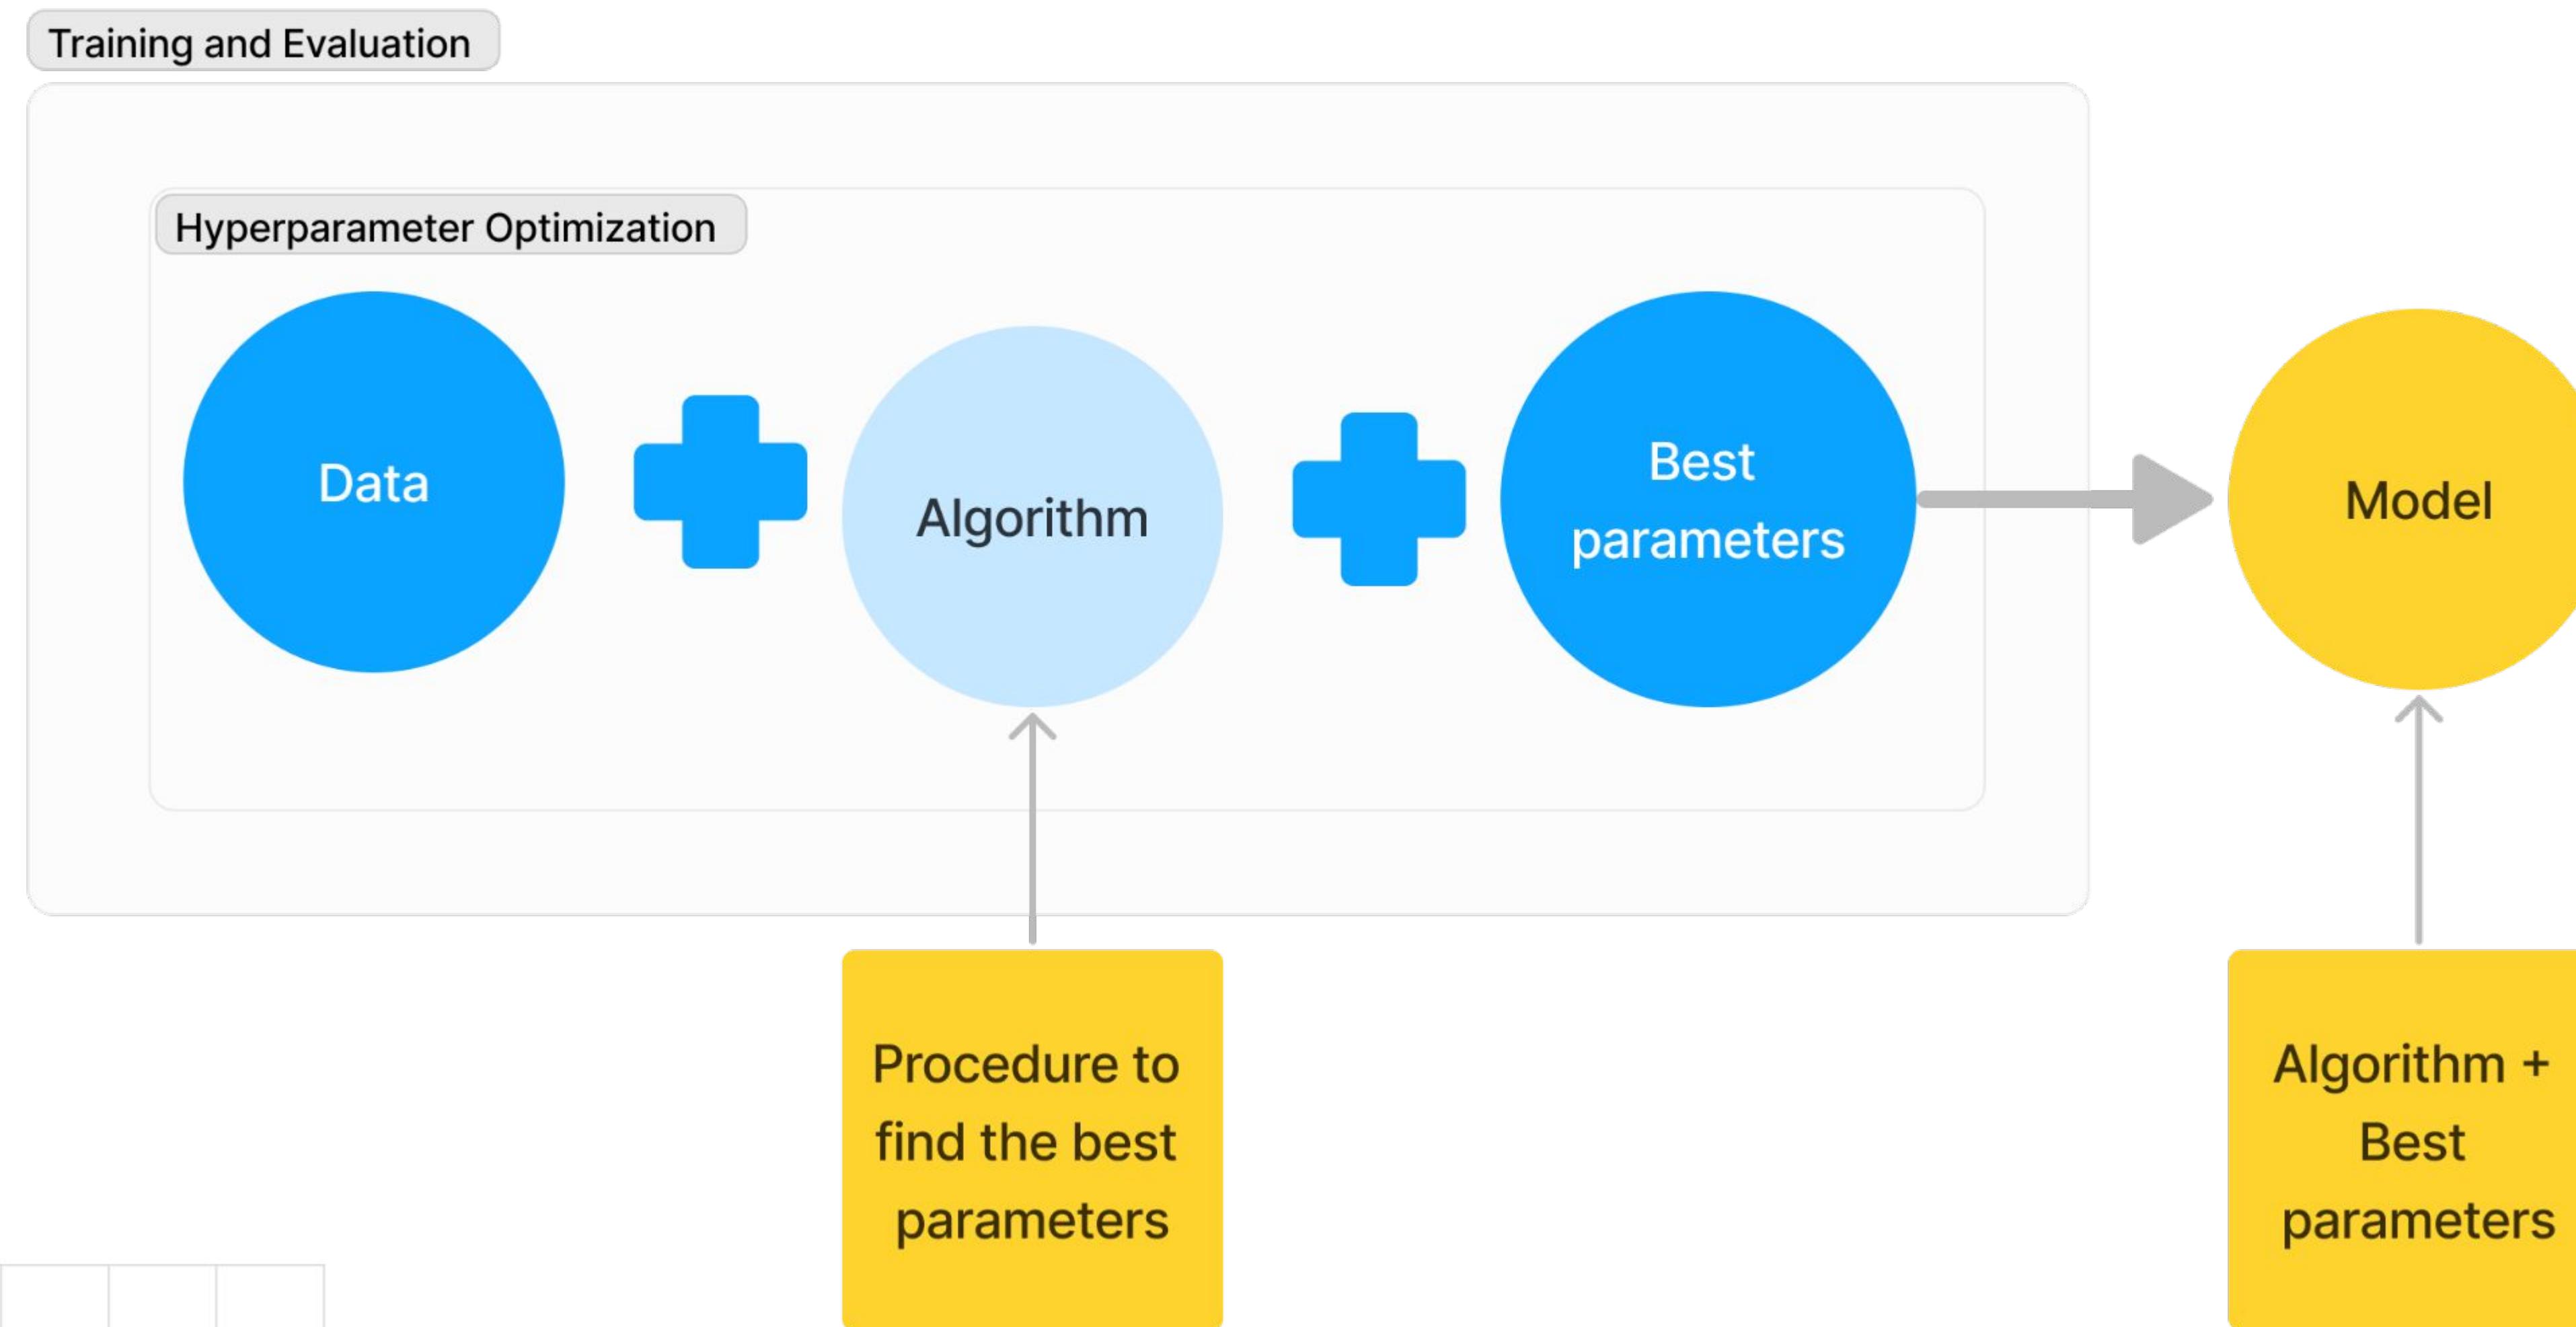
# Access to the notebook



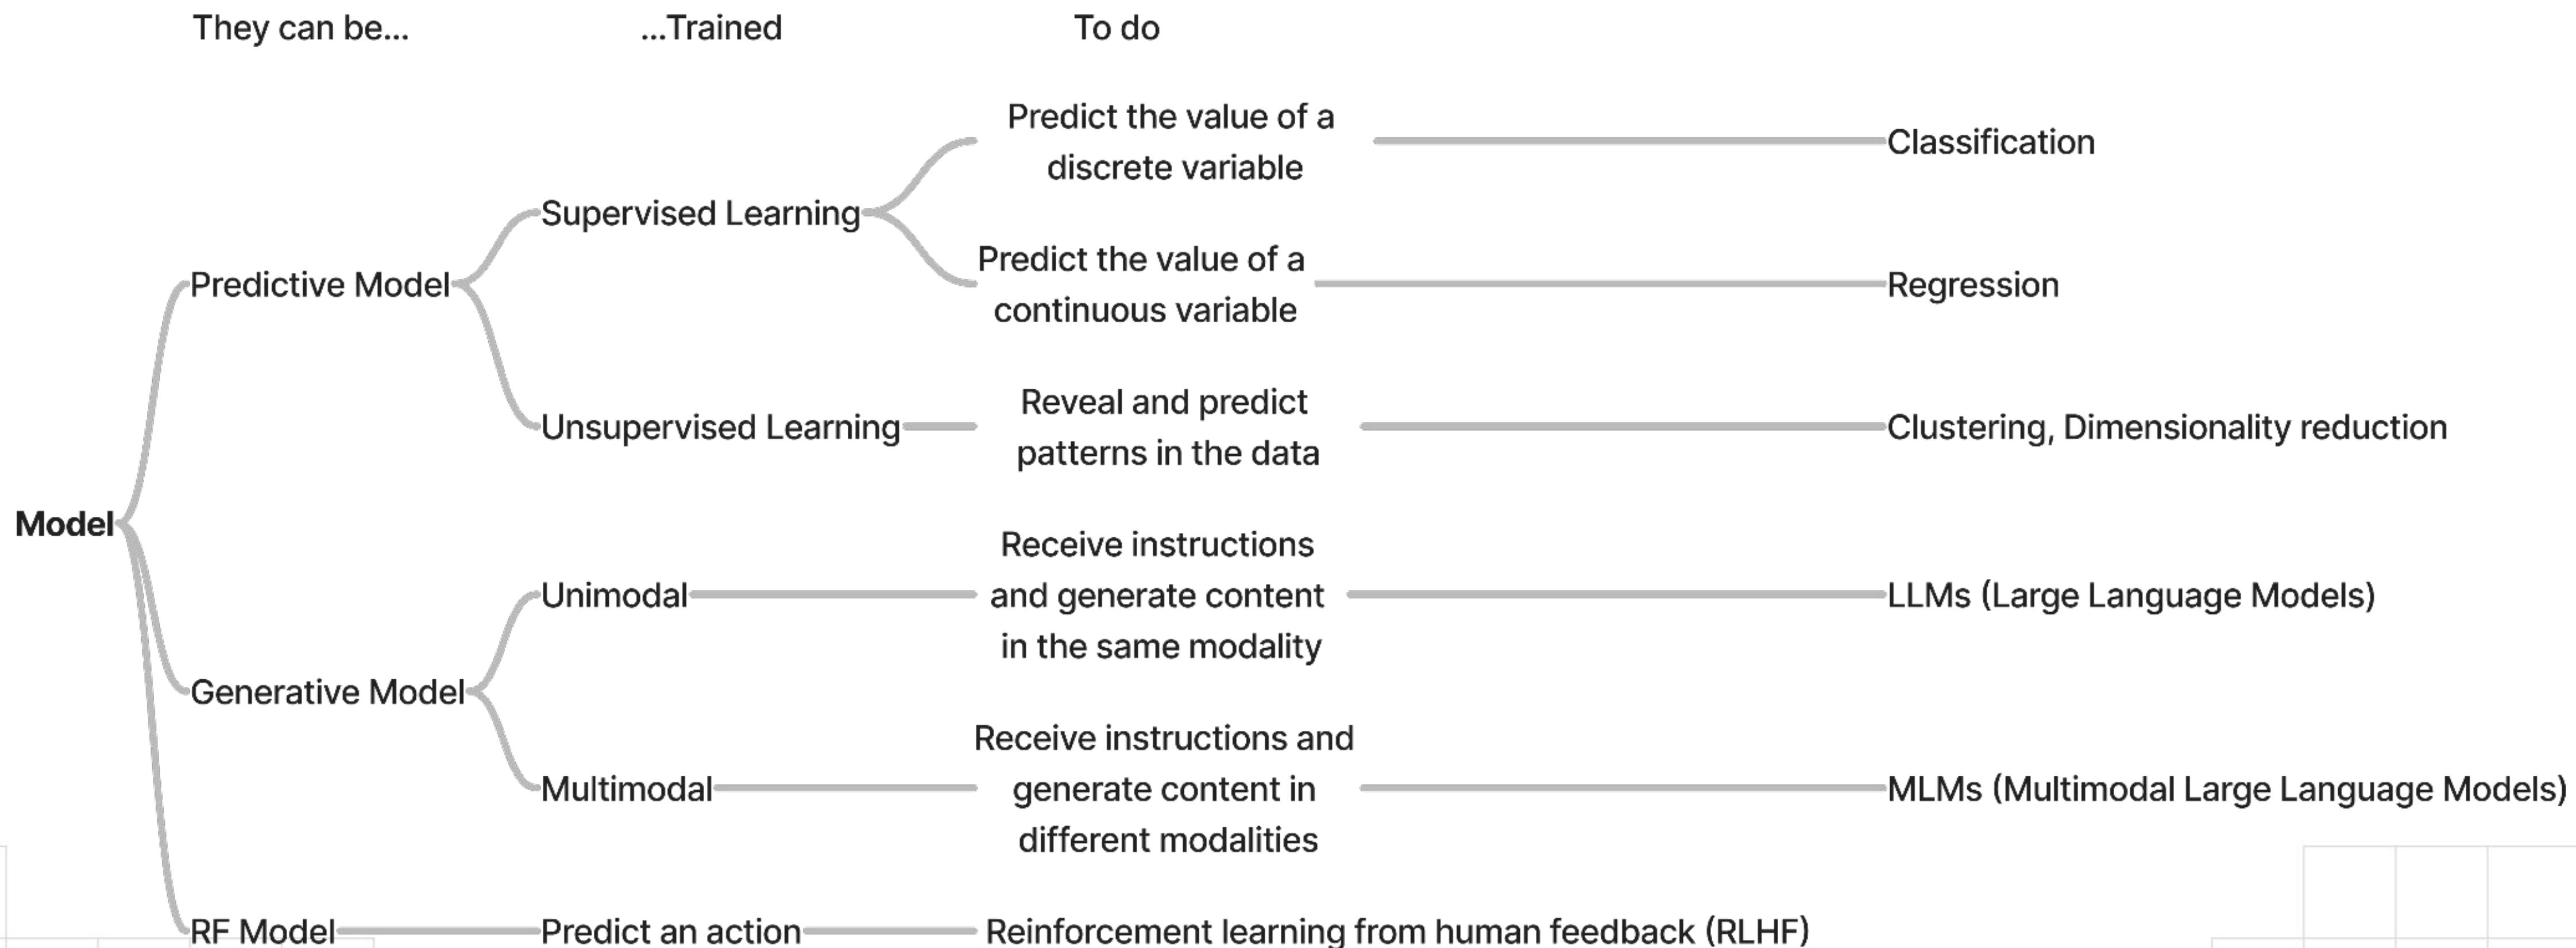
# What is Machine Learning?



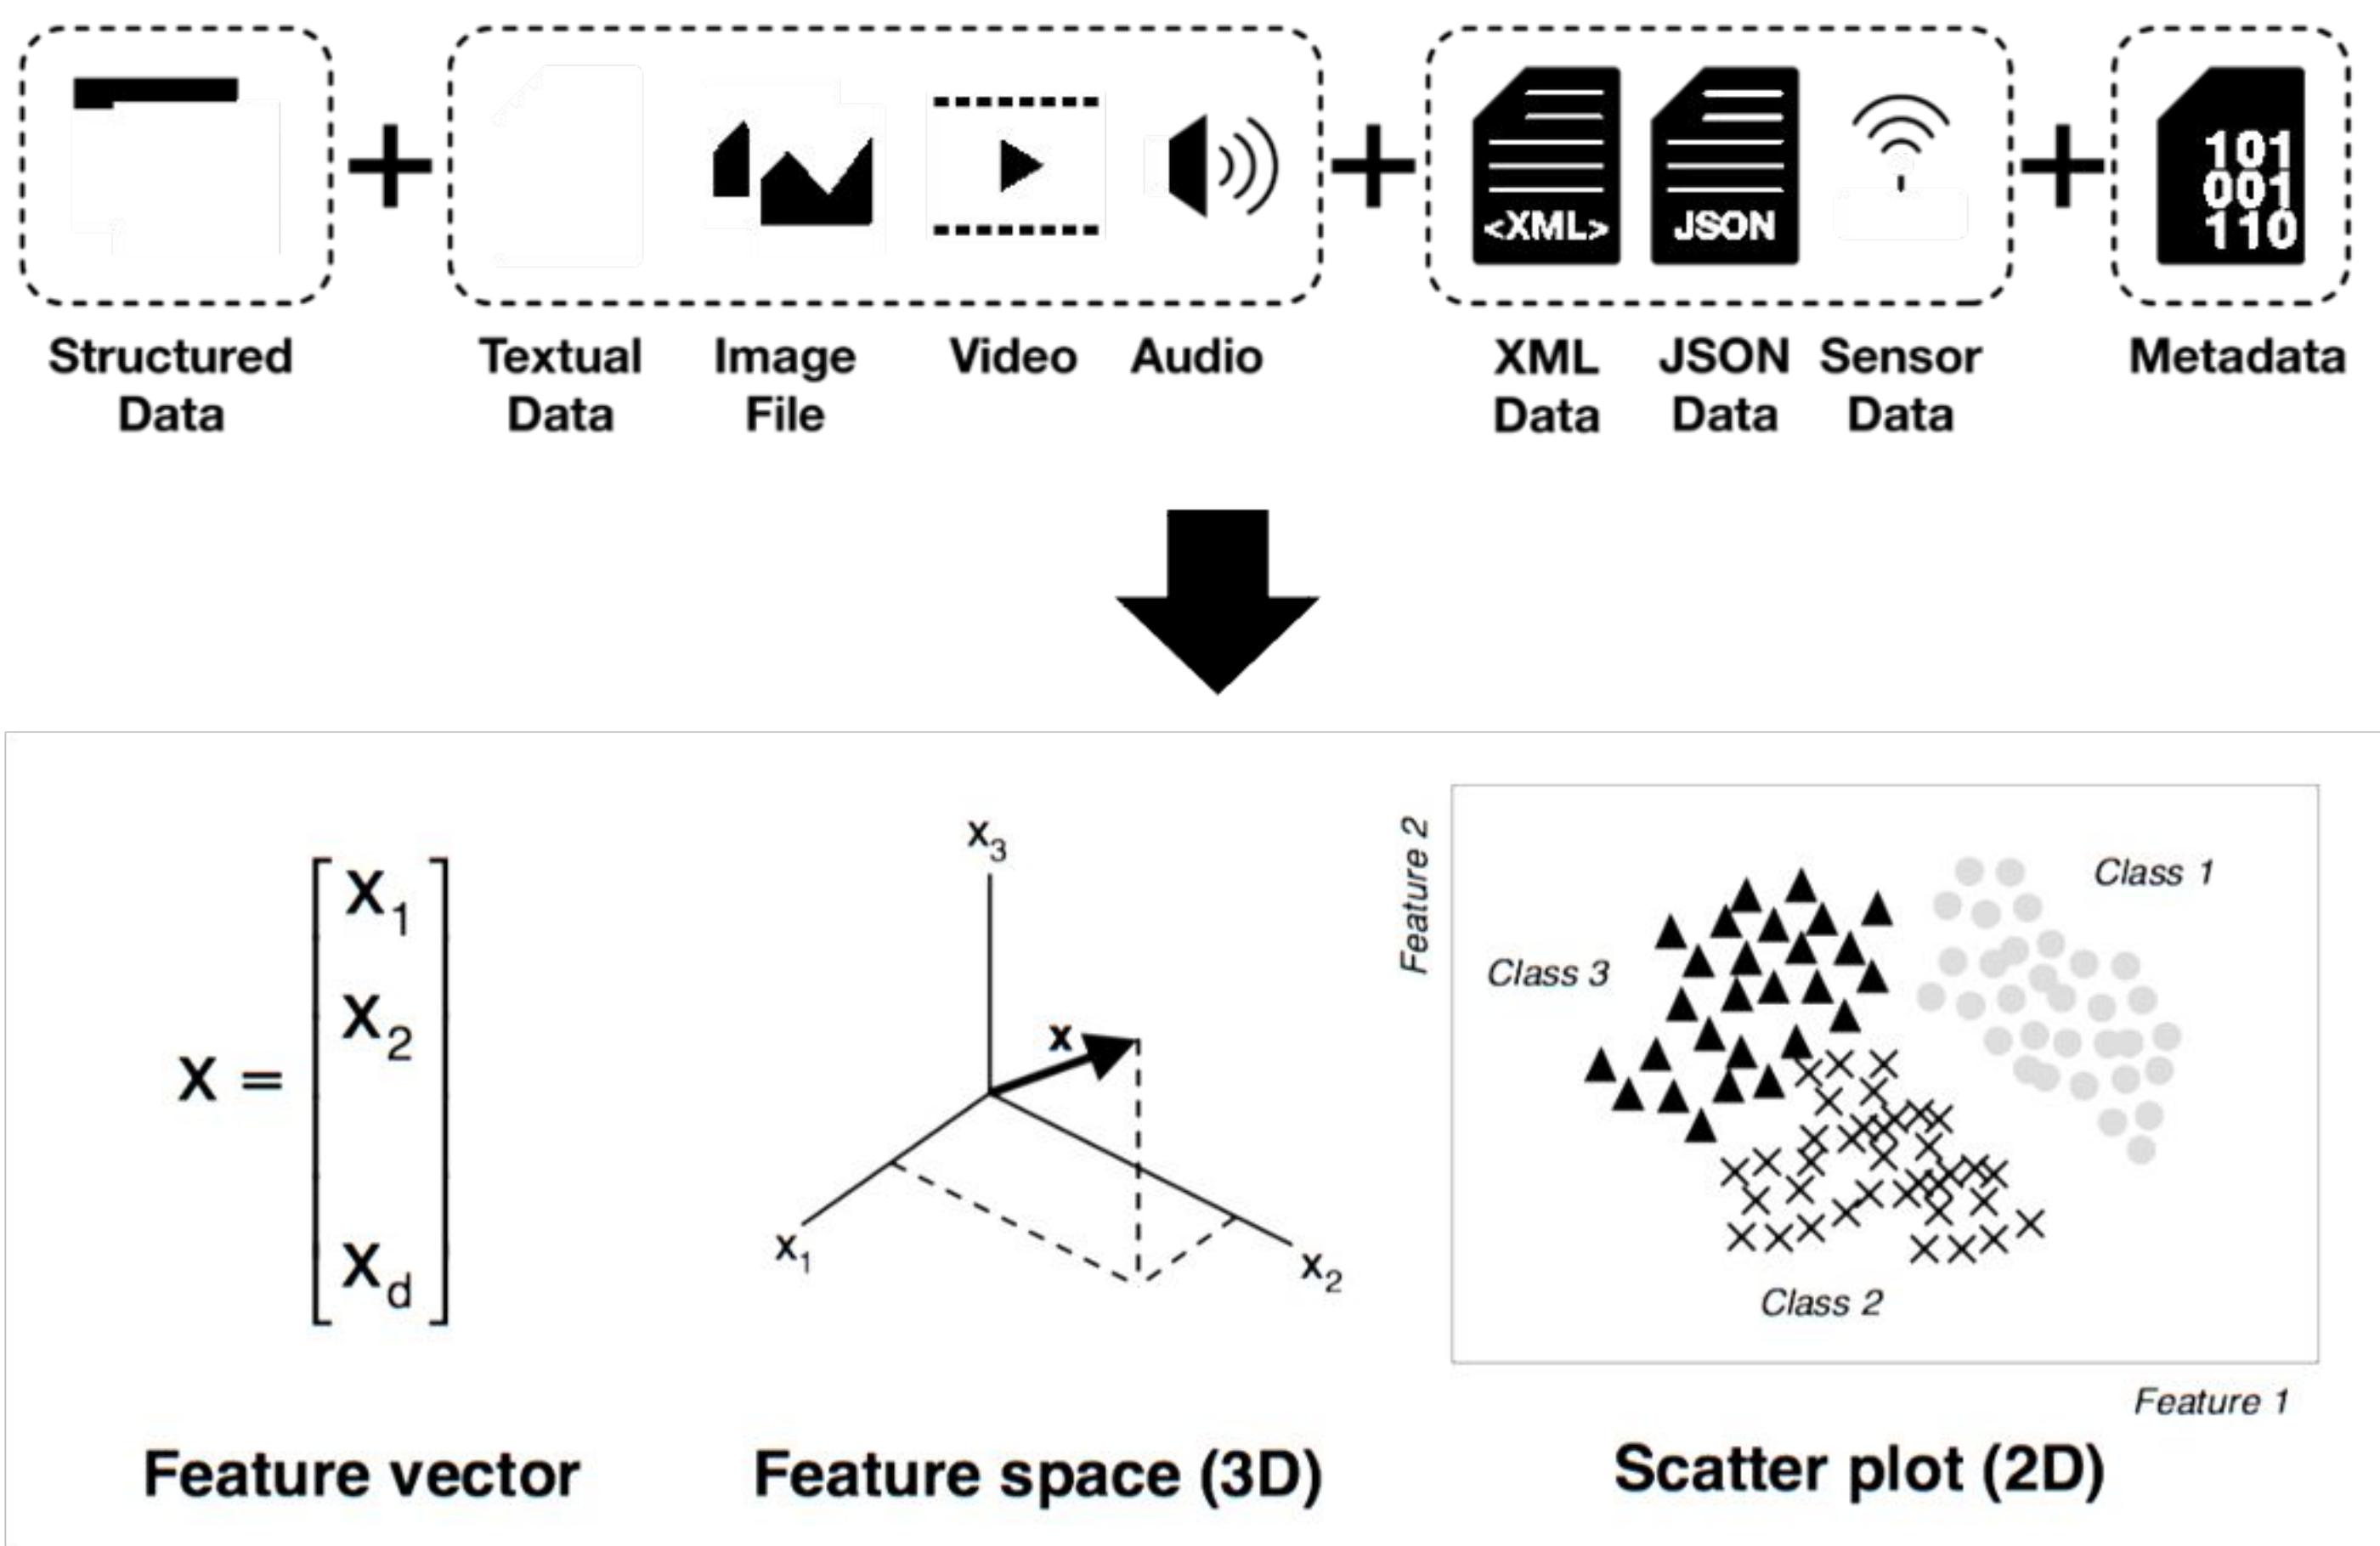
# ....The learning process involves



# The Model can be trained for... (Most common tasks)



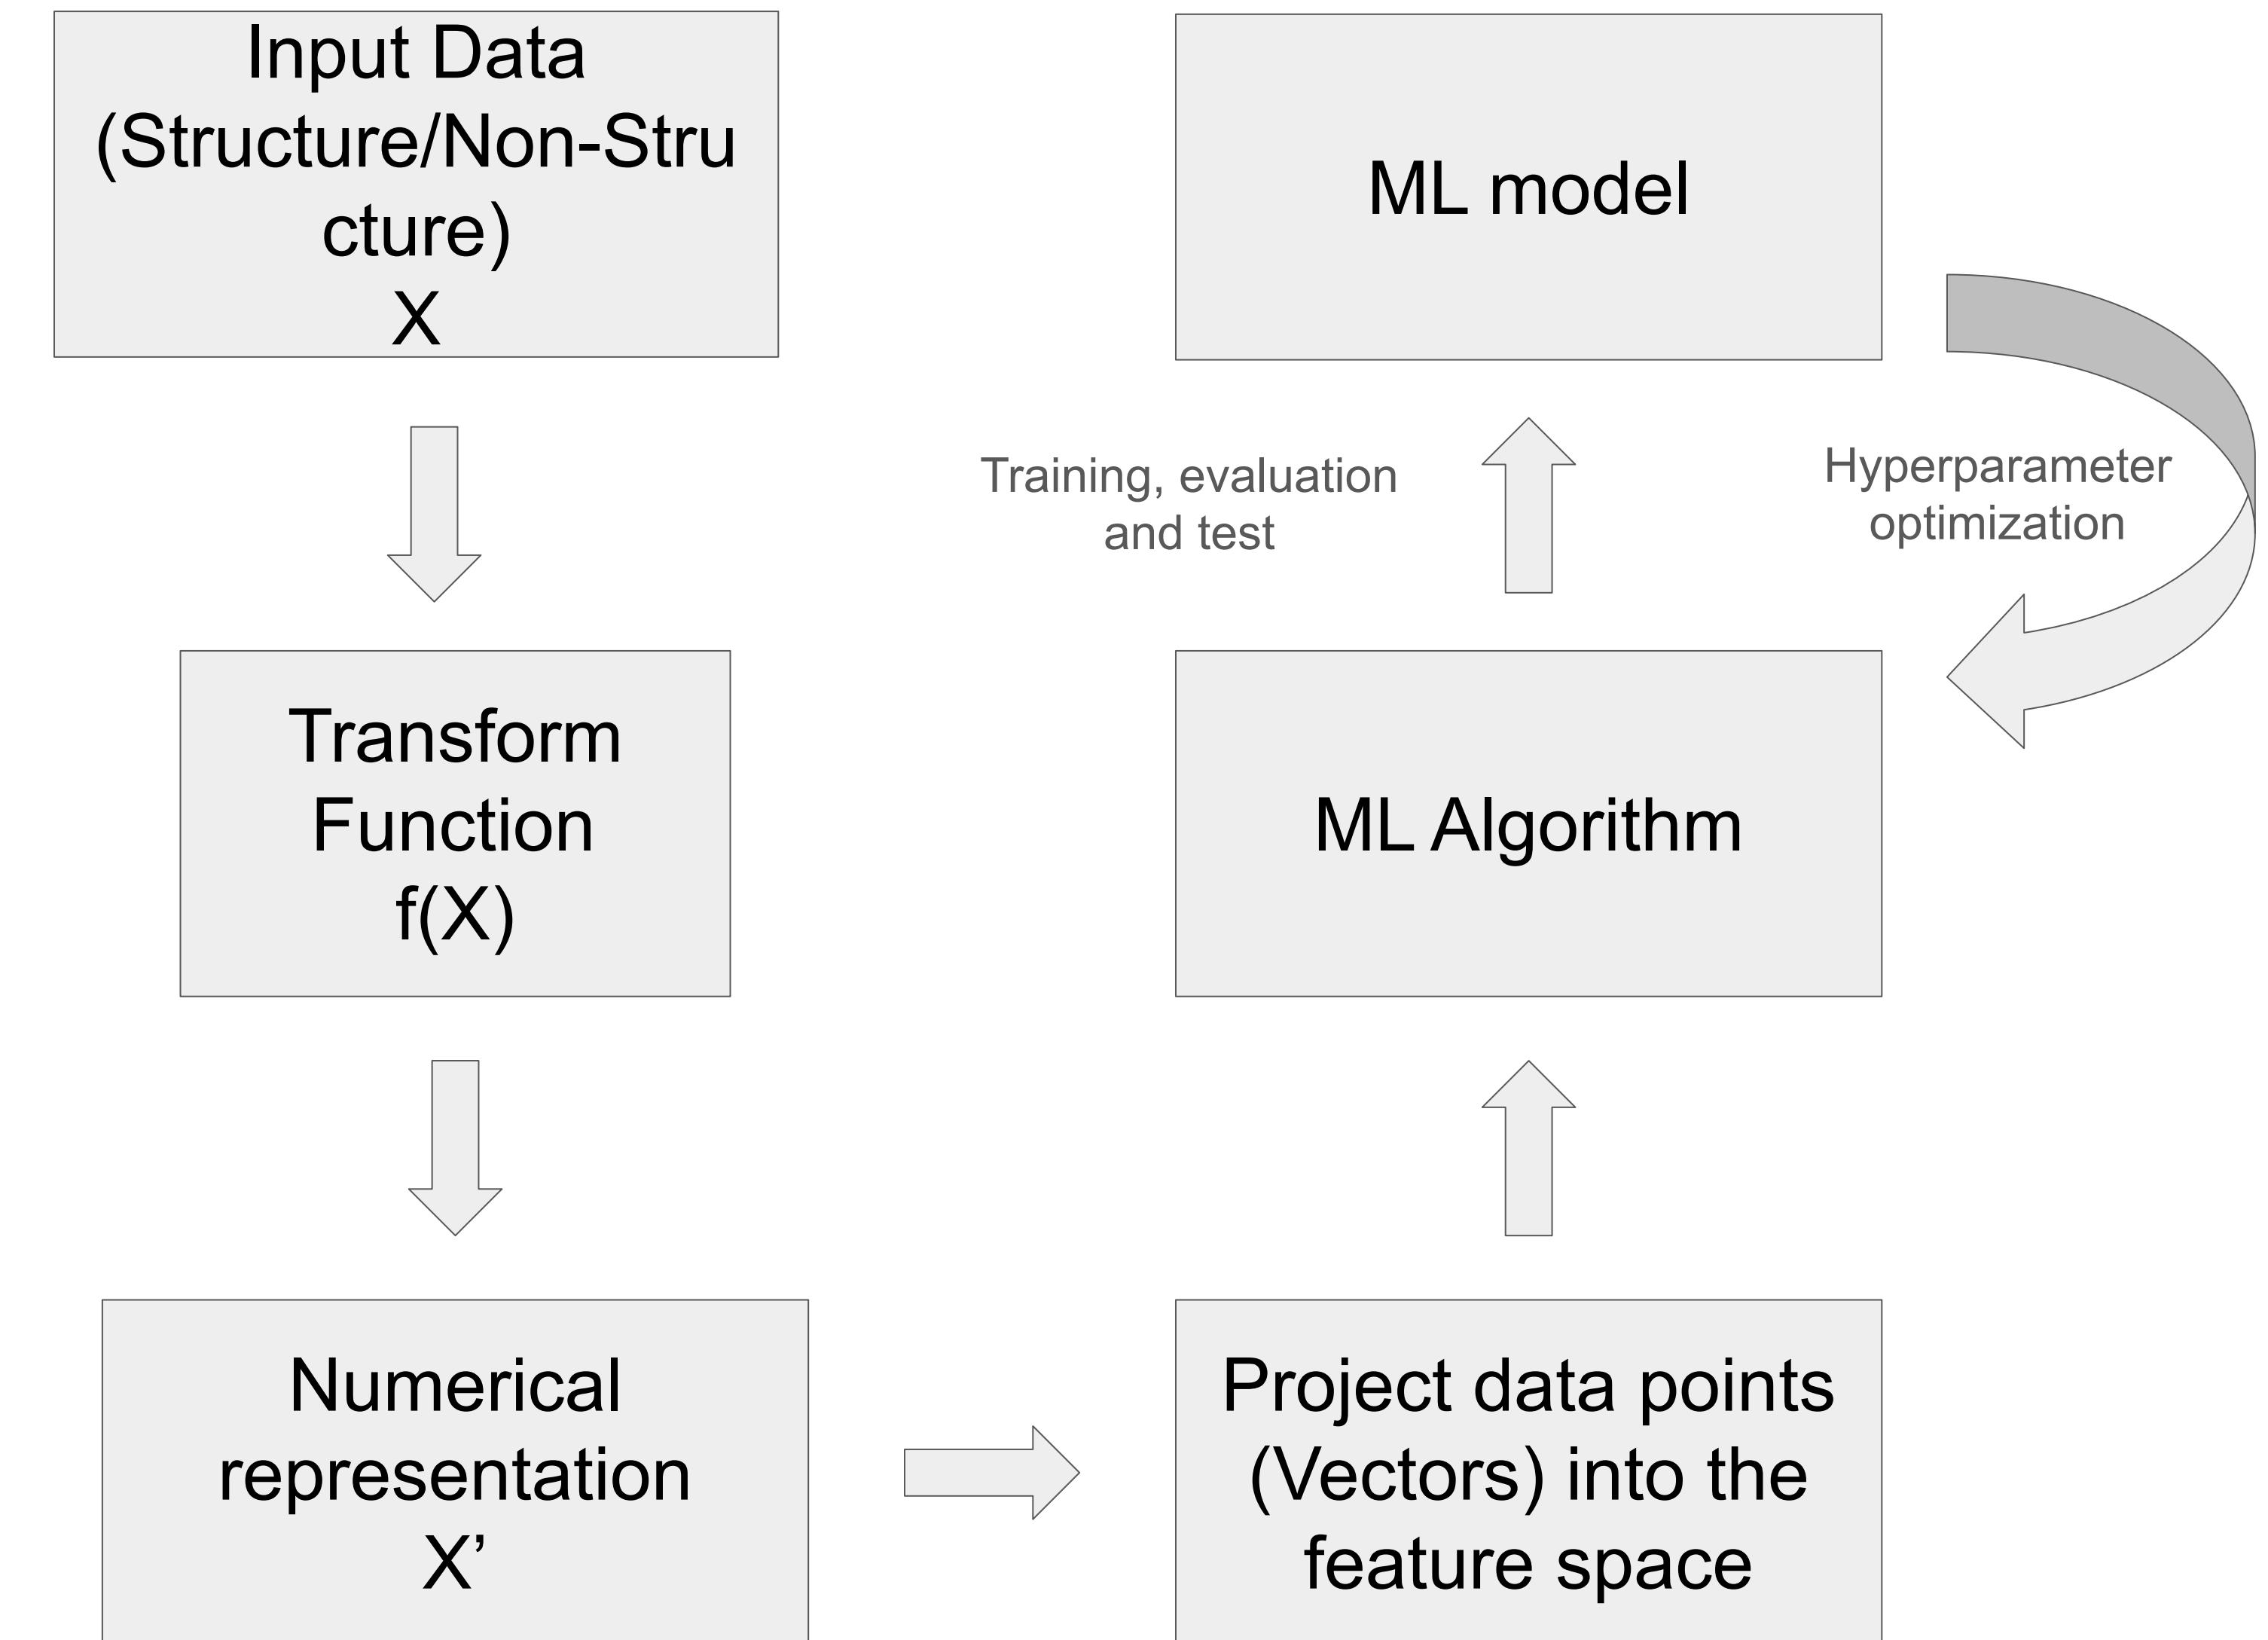
# But, In the end, ML models only understand vectors and numbers



Before or during training the data which could be unstructured or structure must be transformed into a vectorial representation, embedding and projected into the features space.

In the feature space is where the “magic” happens, and the algorithms are applied (as vectors can be mathematically operated).

# Transforming my data into embeddings



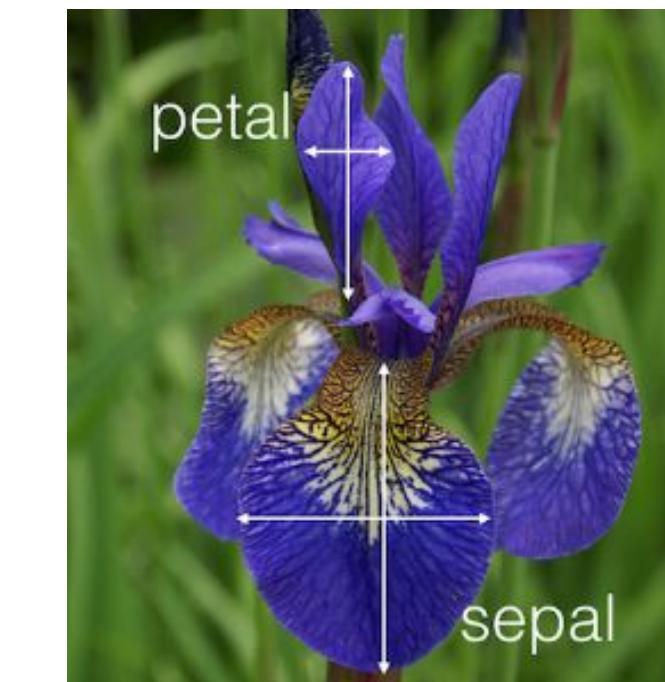
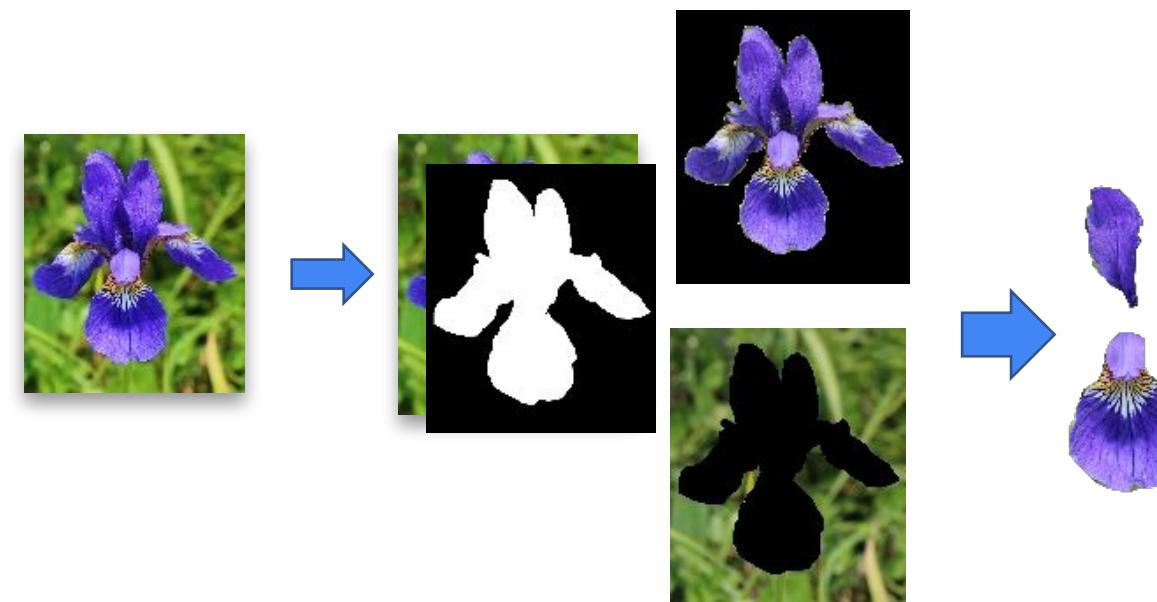
That numerical representation can be hand-crafted,  
applying features engineering  
or learned during training.

Example: You have been asked to develop a mobile app  
in Flutter to classify  
flowers from pictures.

Before deep learning (using traditional ML models and  
feature engineering)

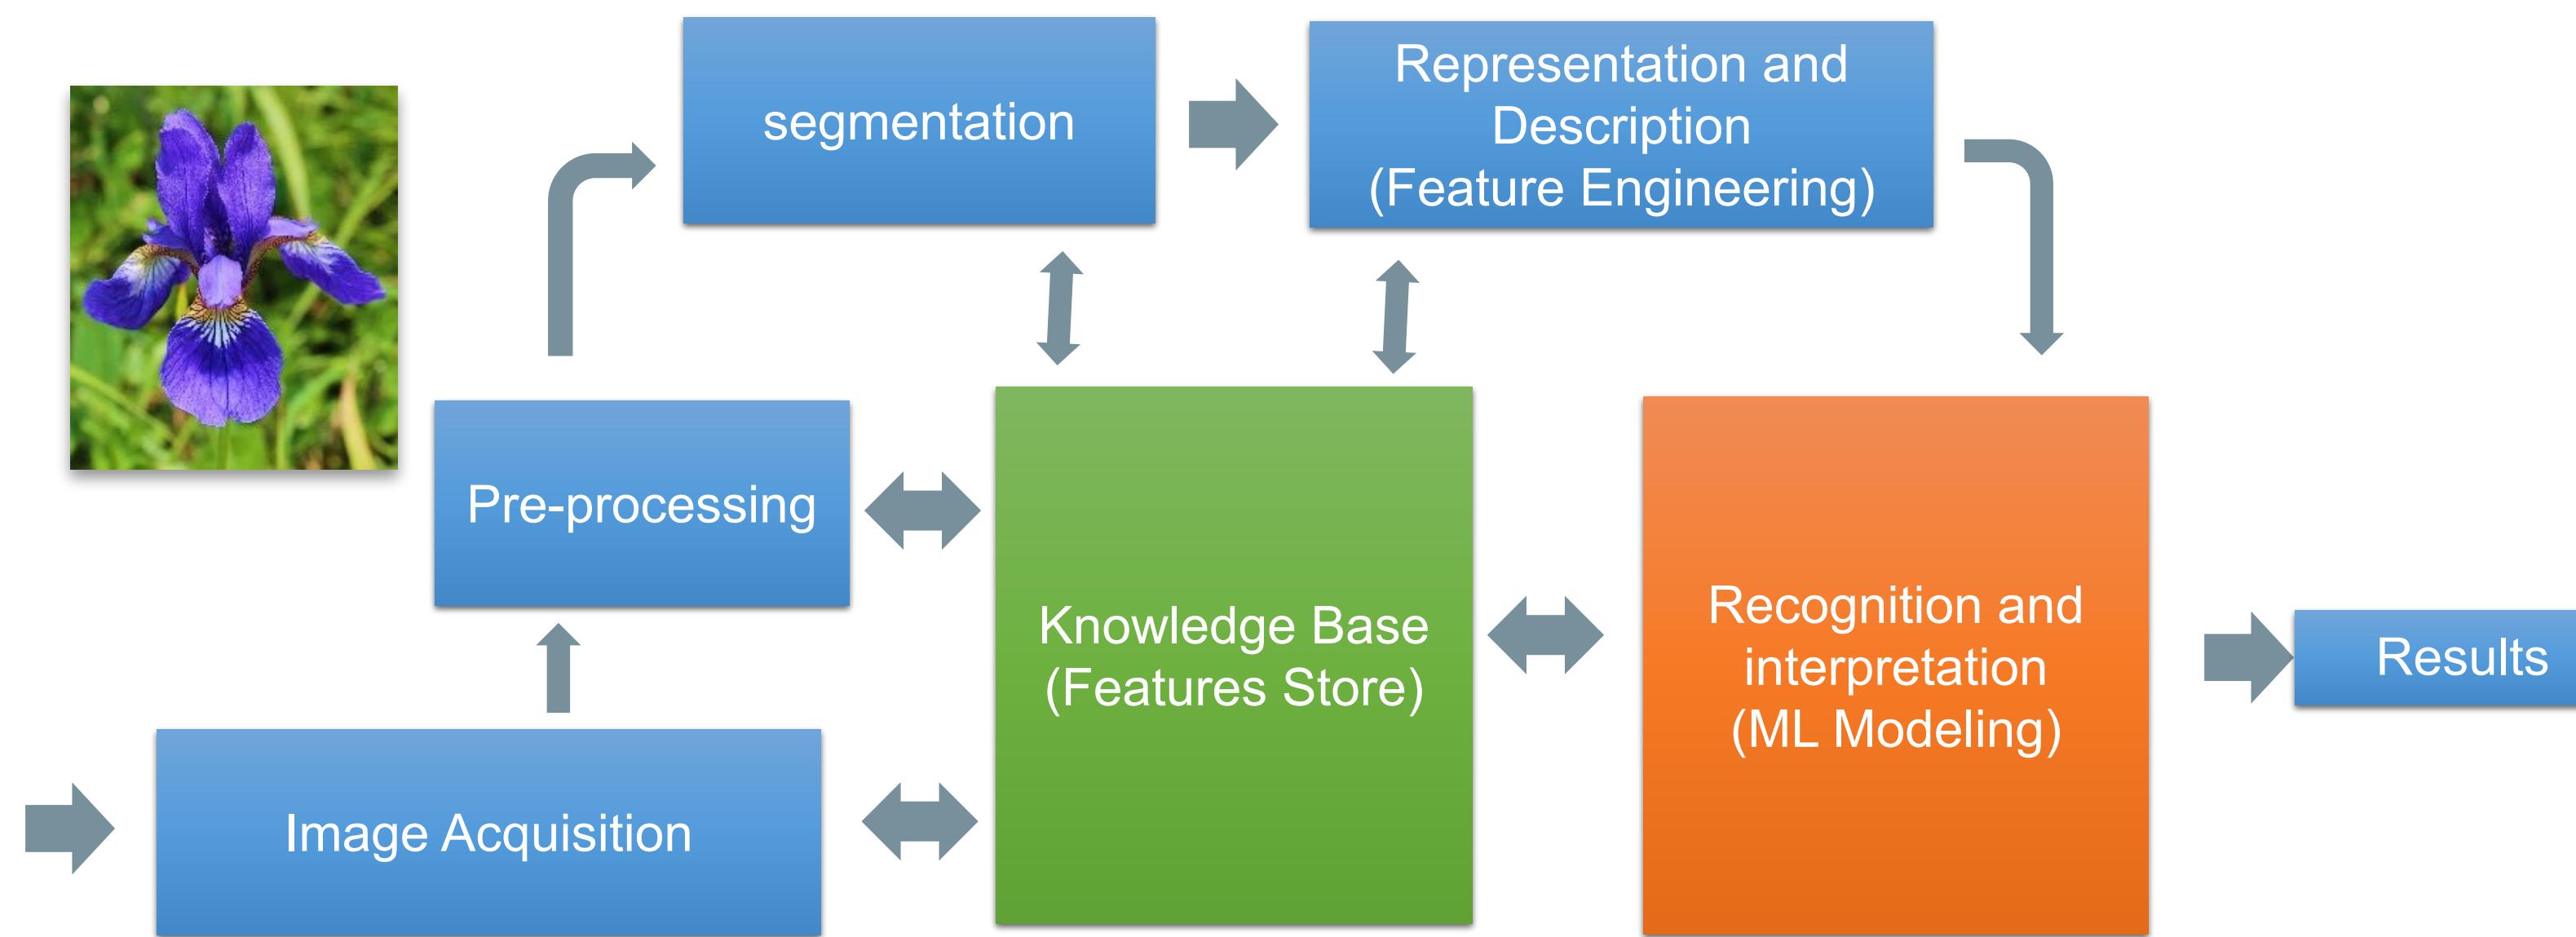
# Features Engineering Approach (Features are known)

Traditional ML relies on your Feature engineering skills to transform your data into numbers/embeddings

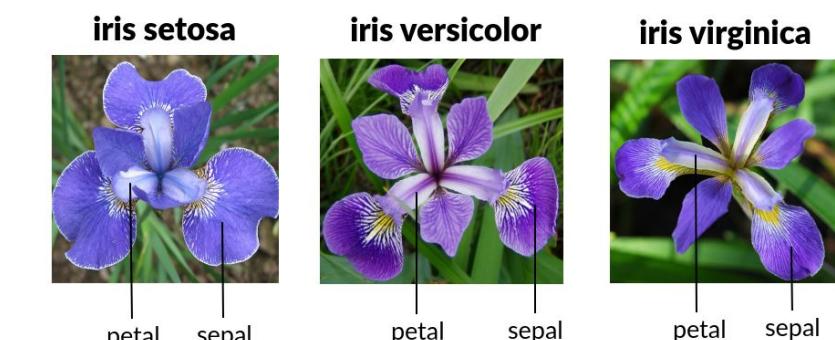


	Sepal.Length	Sepal.Width	Petal.Length	Petal.Width	Species
1	5.1	3.5	1.4	0.2	setosa
2	4.9	3.0	1.4	0.2	setosa
3	4.7	3.2	1.3	0.2	setosa
4	4.6	3.1	1.5	0.2	setosa
5	5.0	3.6	1.4	0.2	setosa
6	5.4	3.9	1.7	0.4	setosa
7	4.6	3.4	1.4	0.3	setosa
8	5.0	3.4	1.5	0.2	setosa

Proficiency issue



Hand Crafted feature extraction  
using image processing



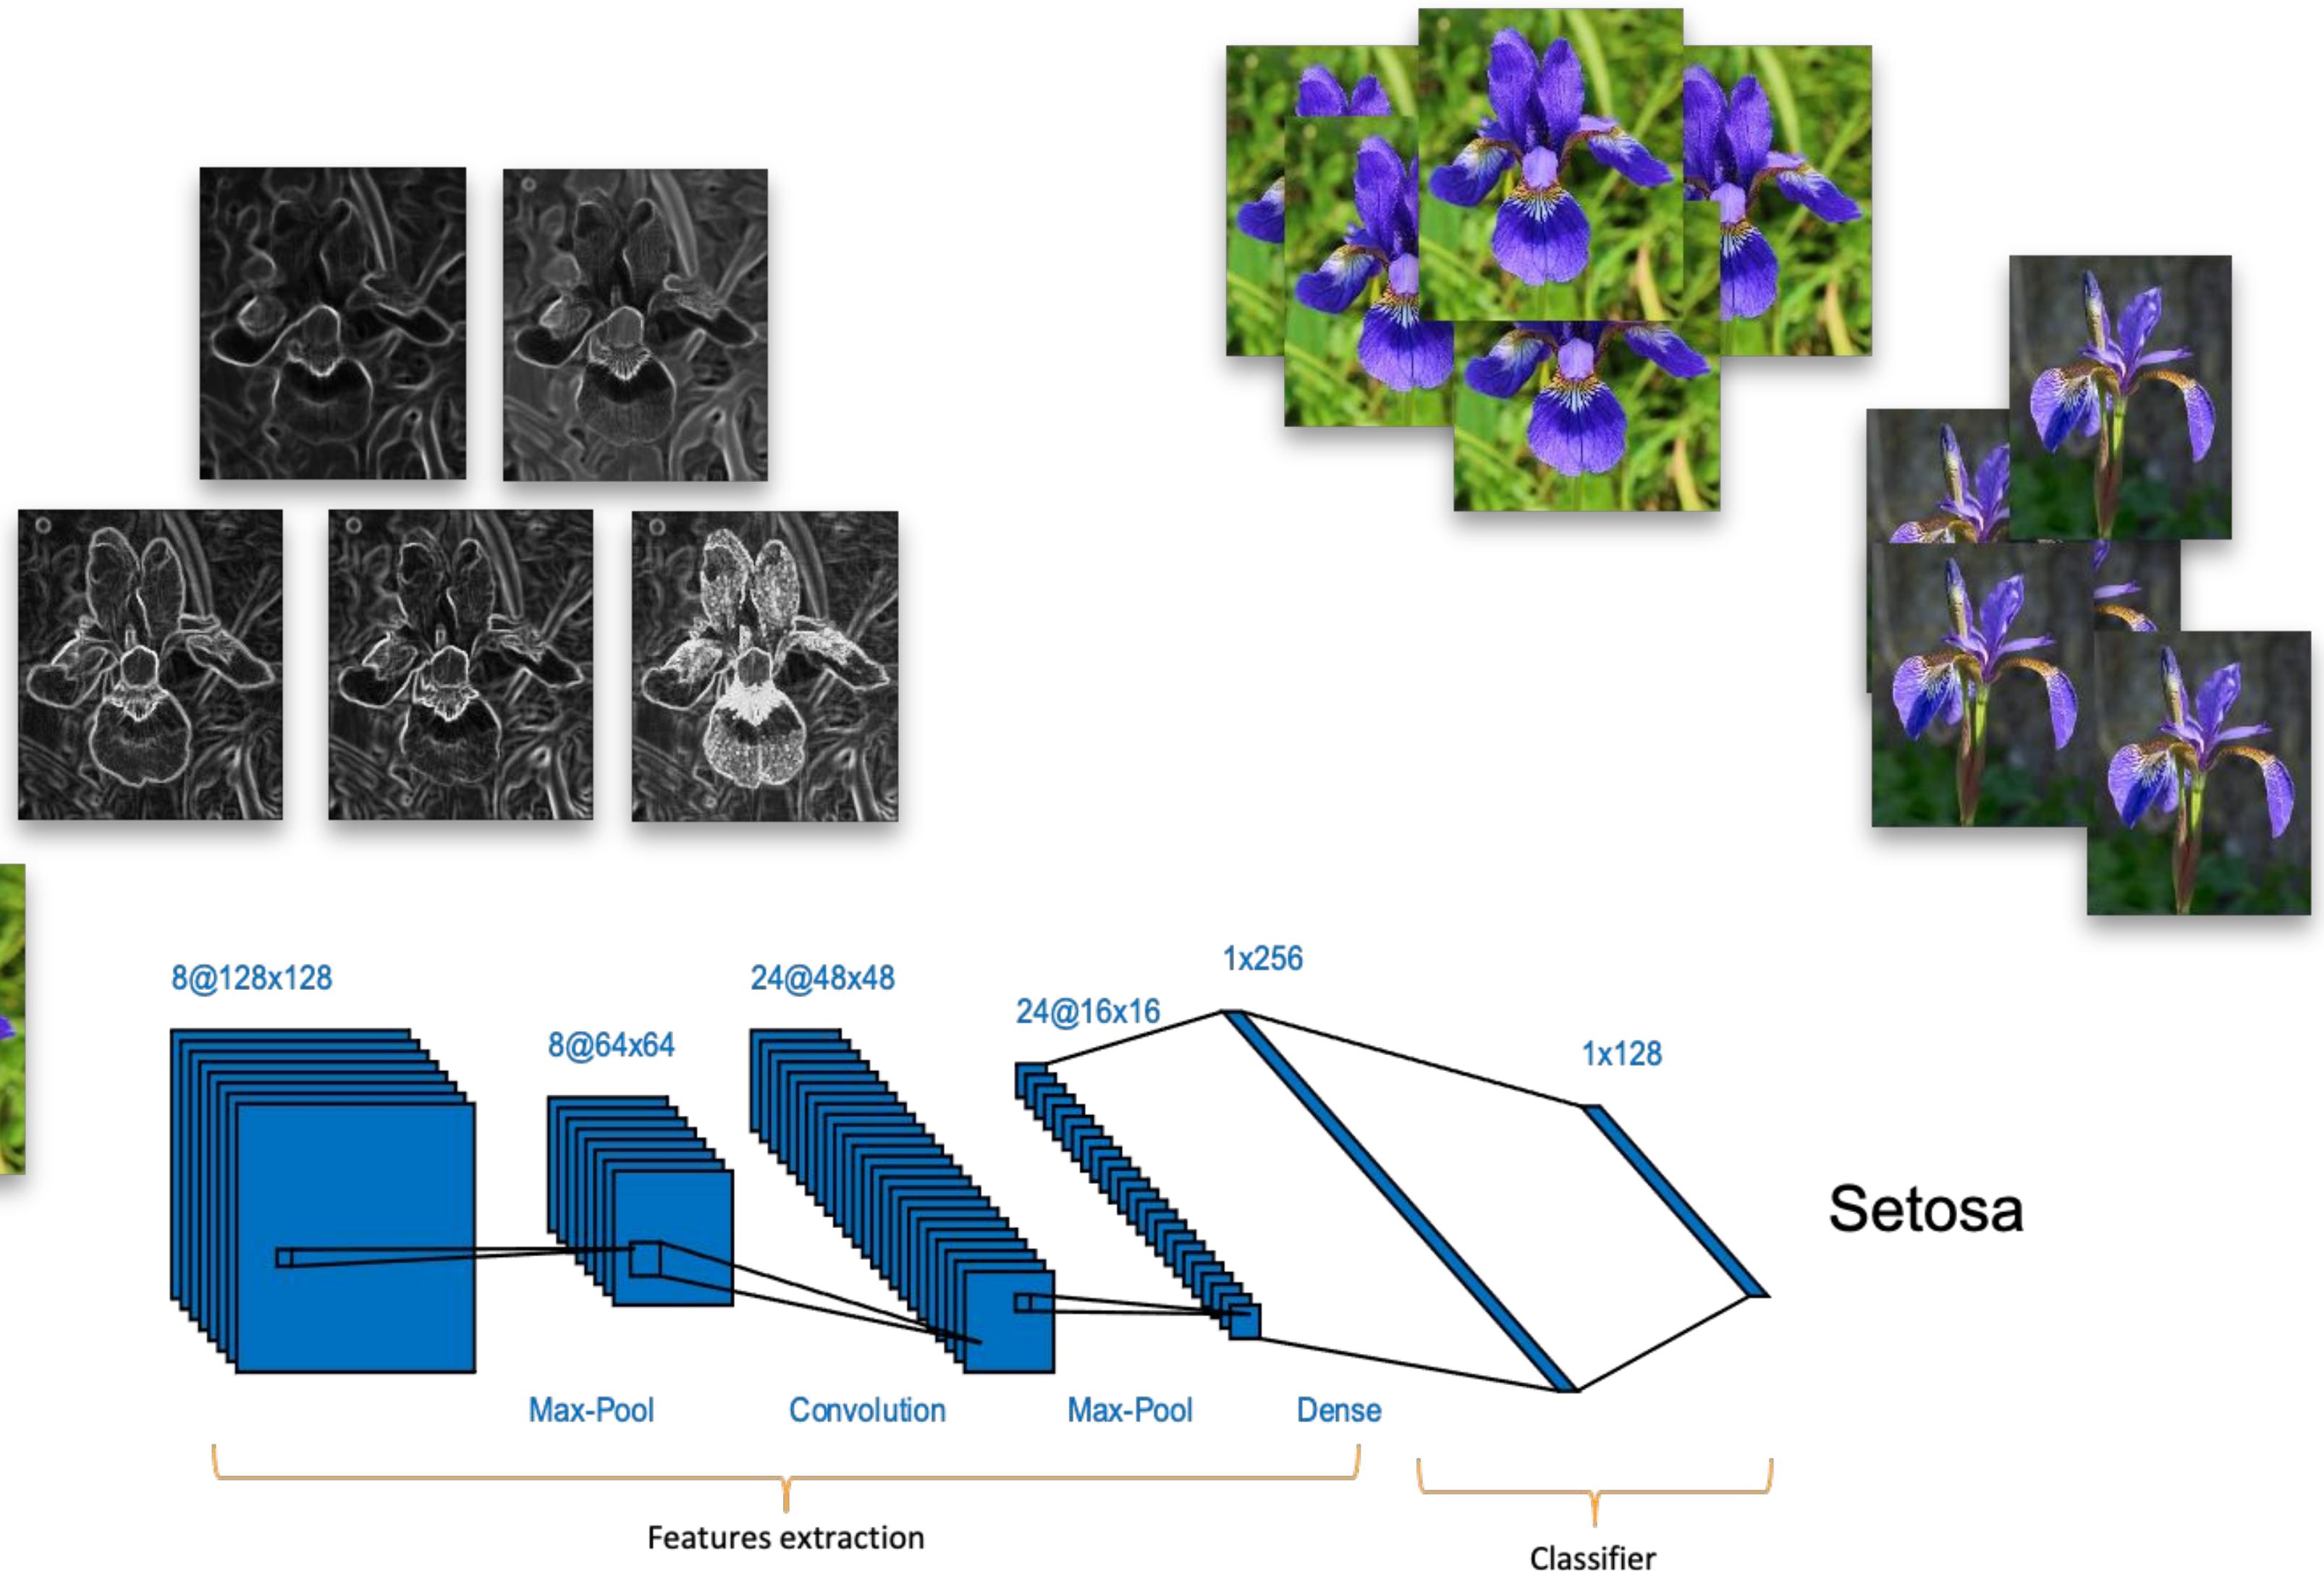
Experts

Google Developers

Using Deep Learning algorithms  
(Automating the feature engineering)

# Deep learning(Features are unknown)

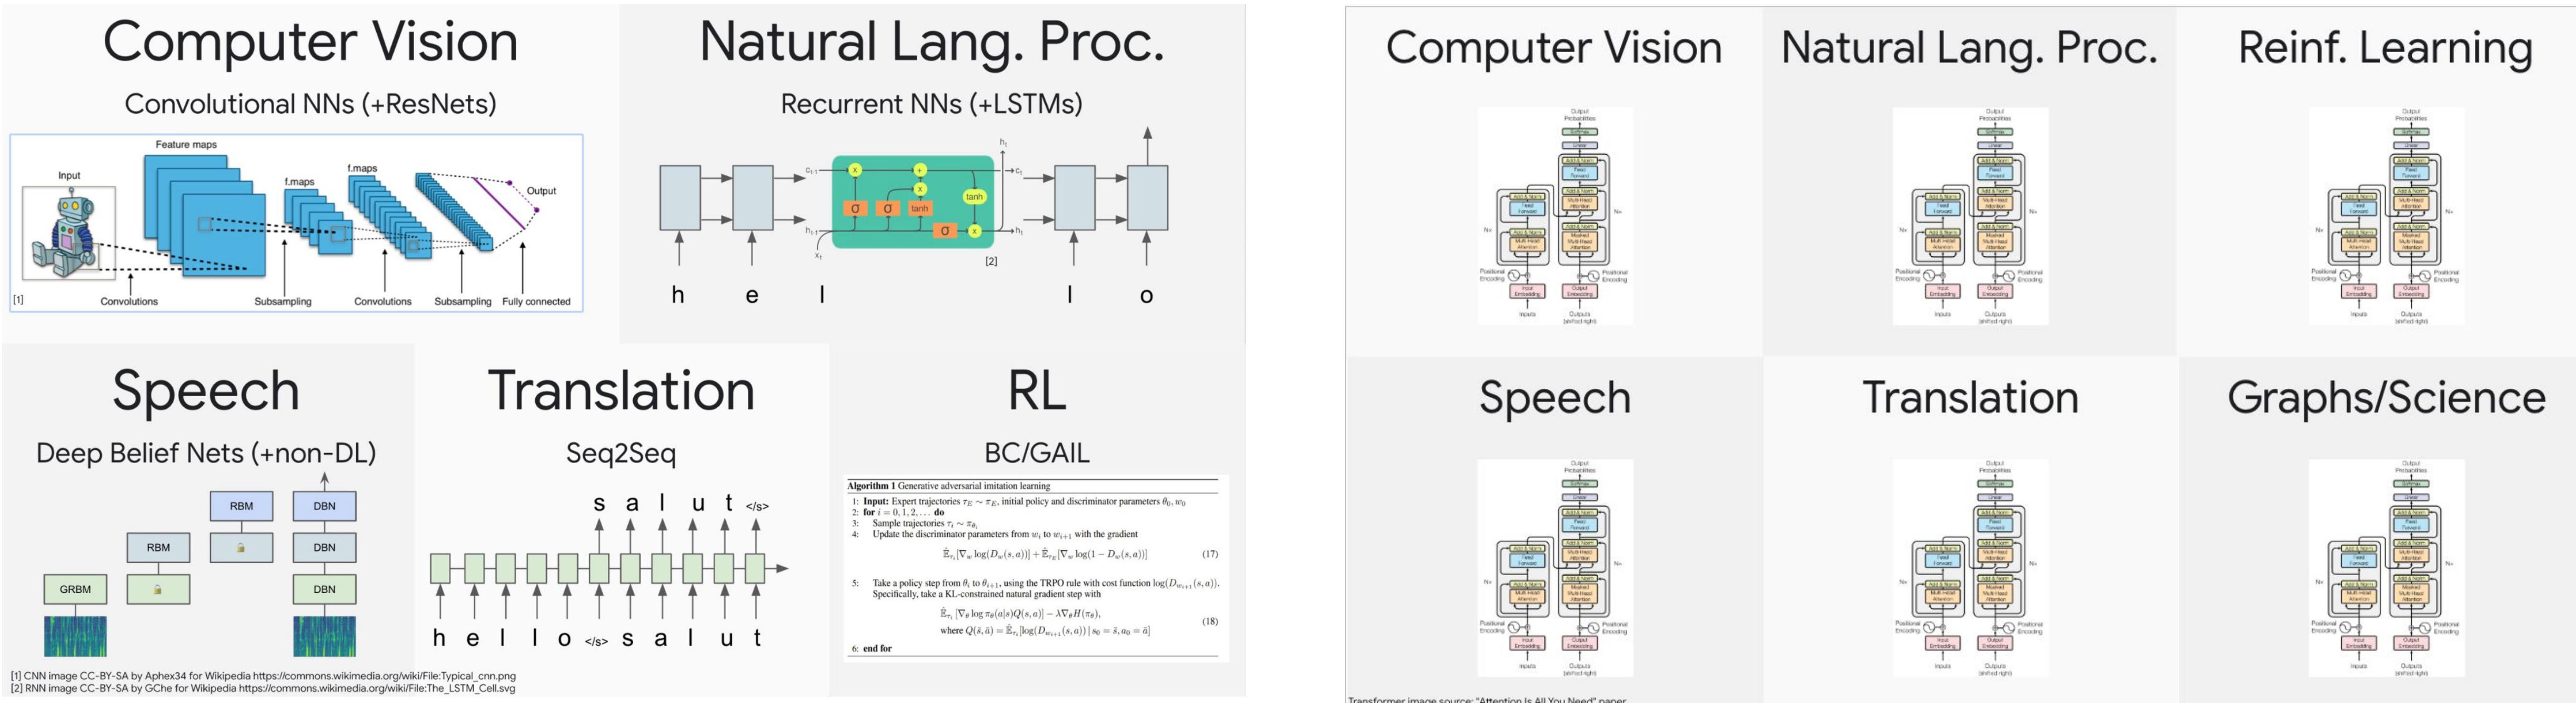
Deep Learning, a type of ML algorithms, automate the feature engineering process by introducing the feature extraction layers.



Automating Feature Extraction through Deep Learning: A Breakthrough in Machine Learning Advances

# Before each task had its own DNN architecture

Therefore, each deep architecture, depending of the task, had its own building blocks (feature extraction layers, loss function etc.)



# Now: Transformers is all you need

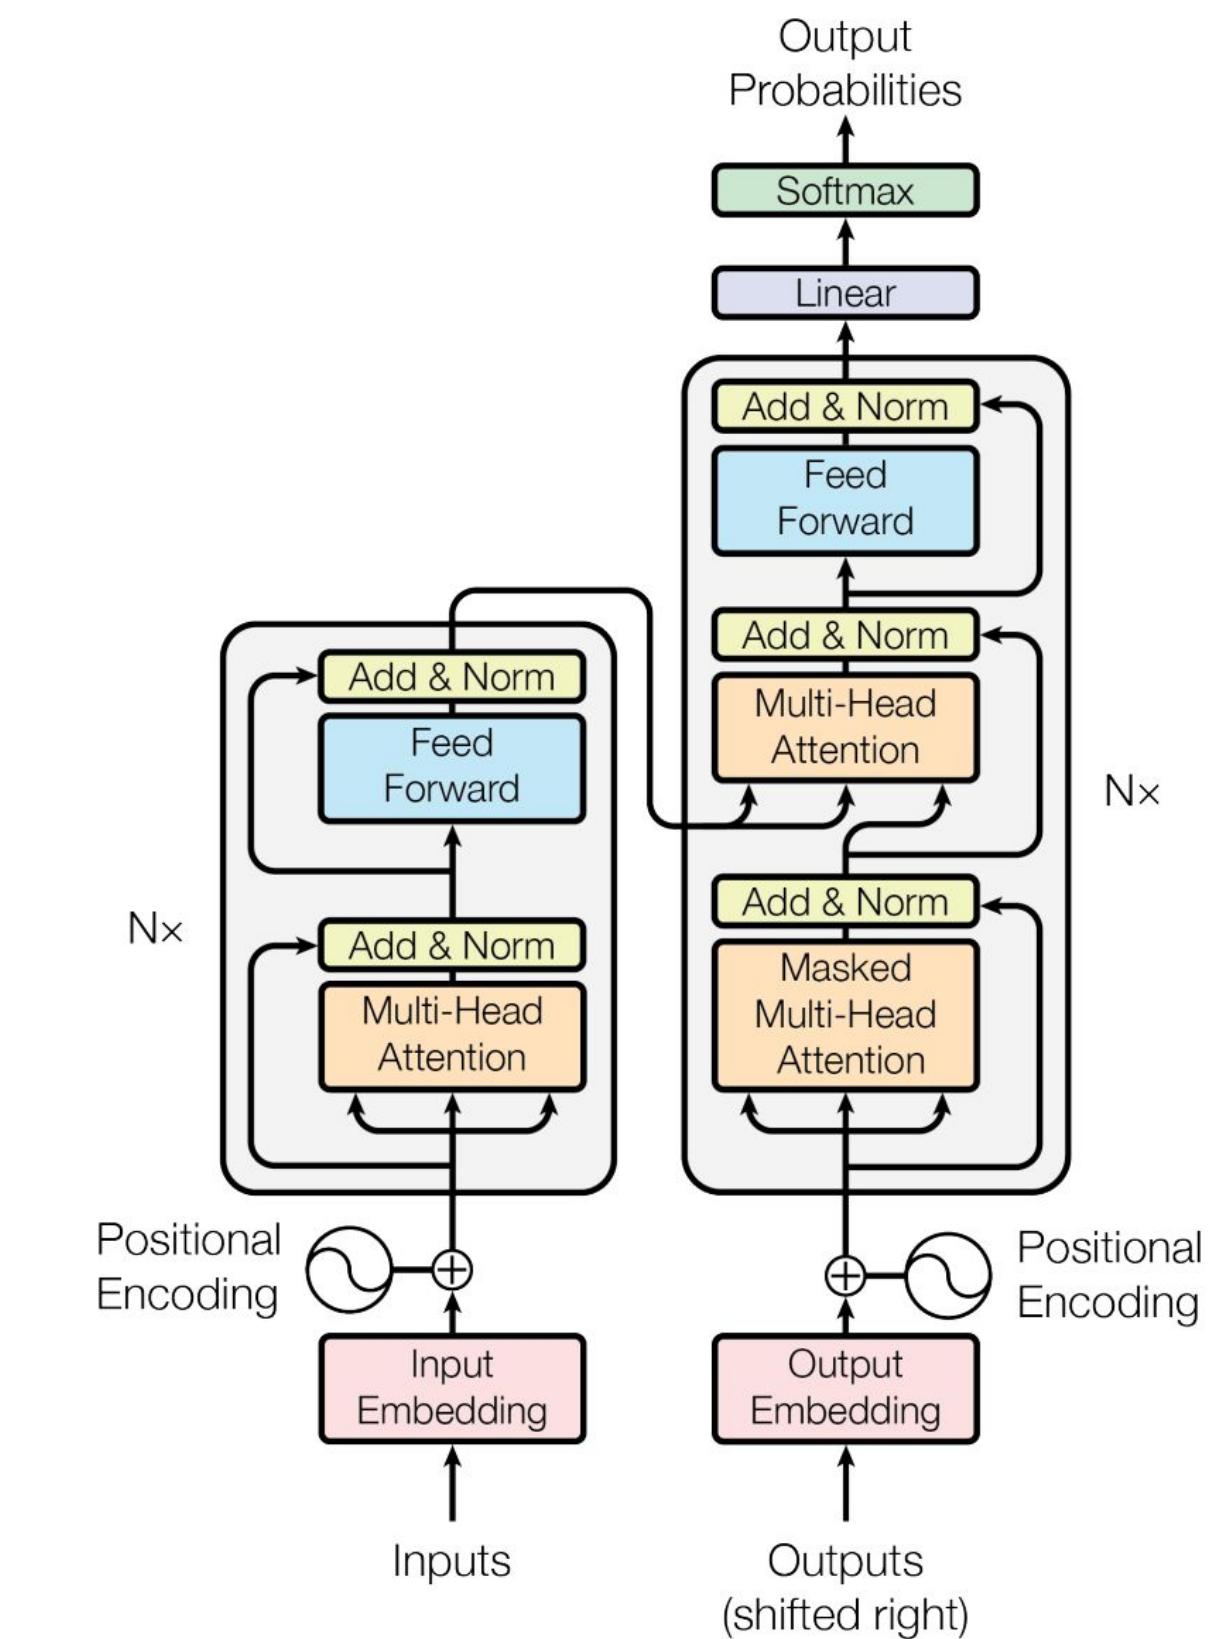
The transformer architecture, introduced in the paper "Attention is All You Need" by Vaswani et al. in 2017, has revolutionized the field of natural language processing (NLP) and beyond

# Transformers novelty

**Attention Mechanism:** Attention mechanisms allows the model to focus on relevant parts of the input sequence during processing.

**Parallelization:** Input sequences can be process in parallel due to their architecture design, leading to faster training times compared to sequential models that process inputs sequentially.

**Scalability:** Transformers are more scalable to longer sequences compared to traditional recurrent models. They can process inputs of arbitrary length without increasing computational complexity significantly, making them suitable for tasks requiring long context understanding, such as machine translation and document summarization.



Andrej Karpathy ✅  
@karpathy

The Transformer is a magnificent neural network architecture because it is a general-purpose differentiable computer. It is simultaneously:  
1) expressive (in the forward pass)  
2) optimizable (via backpropagation+gradient descent)  
3) efficient (high parallelism compute graph)

11:54 AM · Oct 19, 2022

# Anything you can tokenize, you can feed to Transformer

ca 2021 and onwards

Tokenize different modalities each in their own way (some kind of "patching"), and send them all jointly into a Transformer...

Seems to just work...

Currently an explosion of works doing this!

Tokens are vector representations of text, and embeddings are tokens with semantic context.

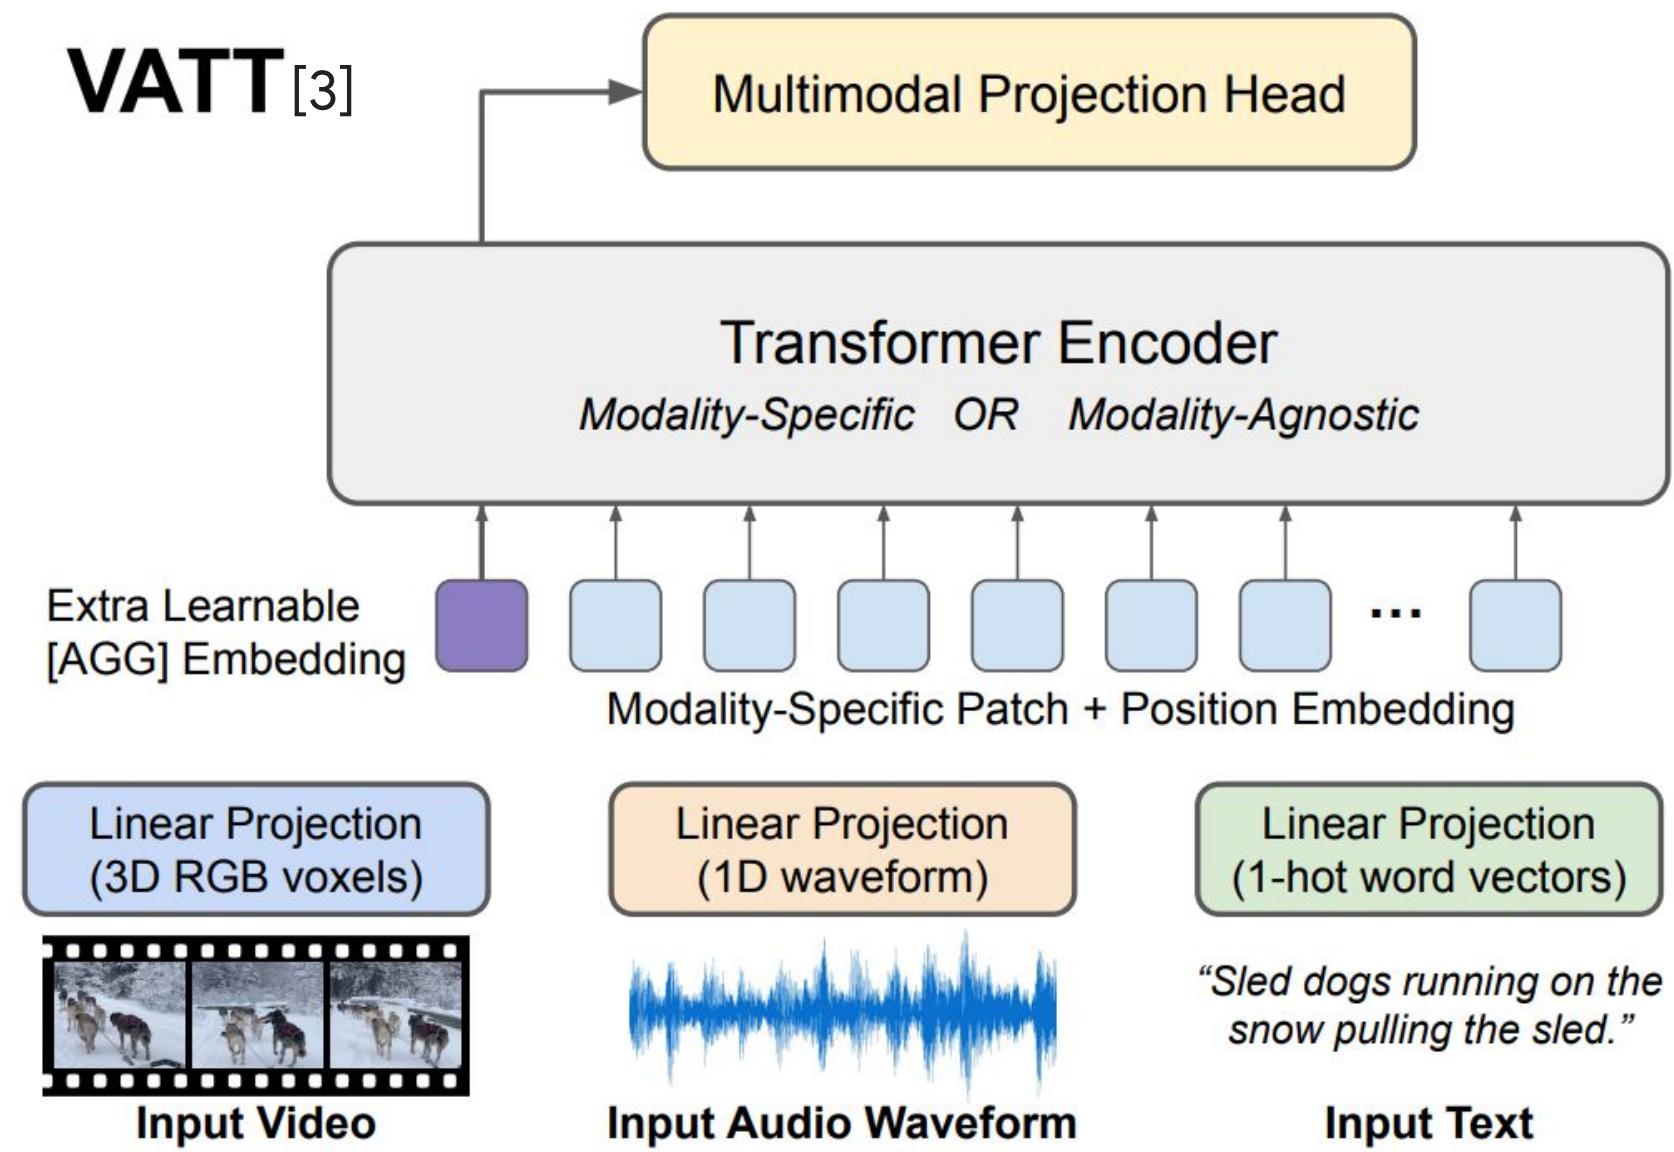
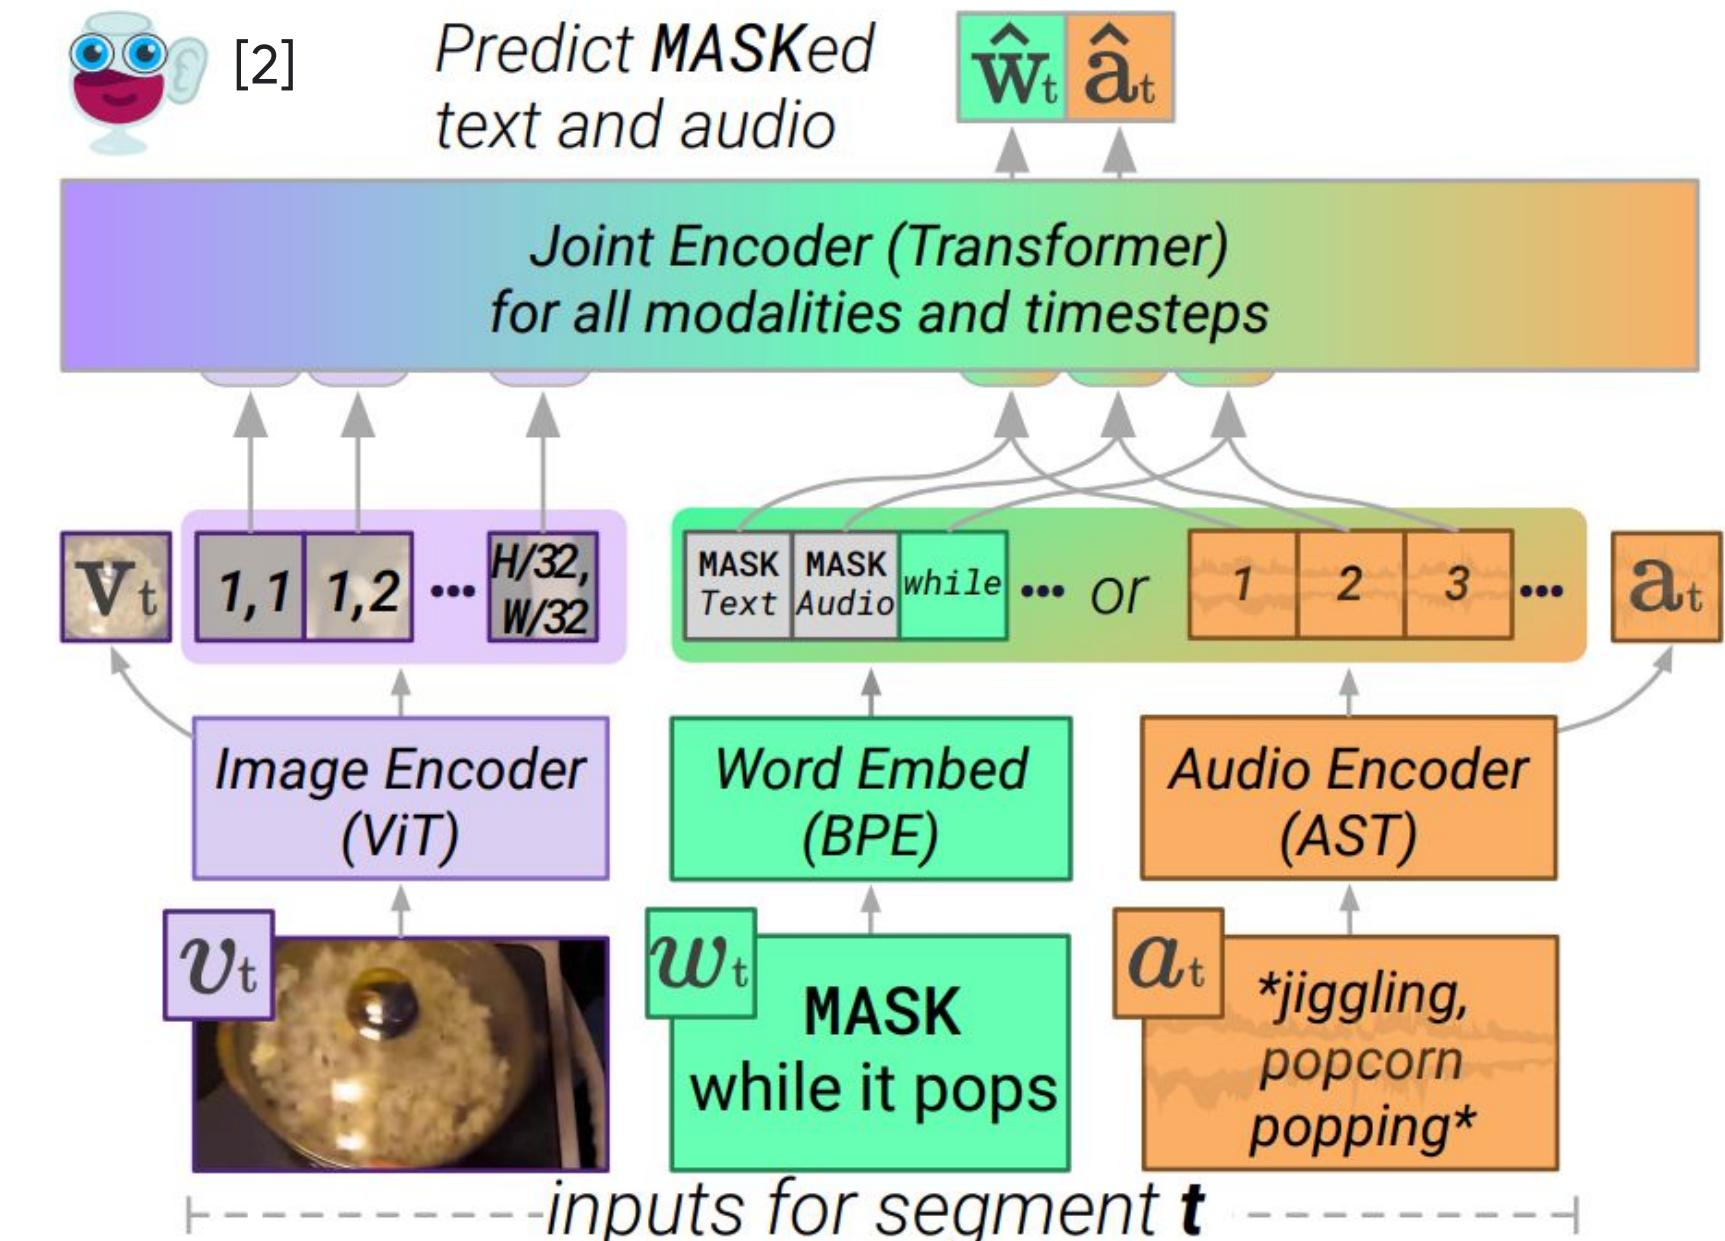
[1]

Images from:

[1] LIMoE by B Mustafa, C Riquelme, J Puigcerver, R Jenatton, N Houlsby

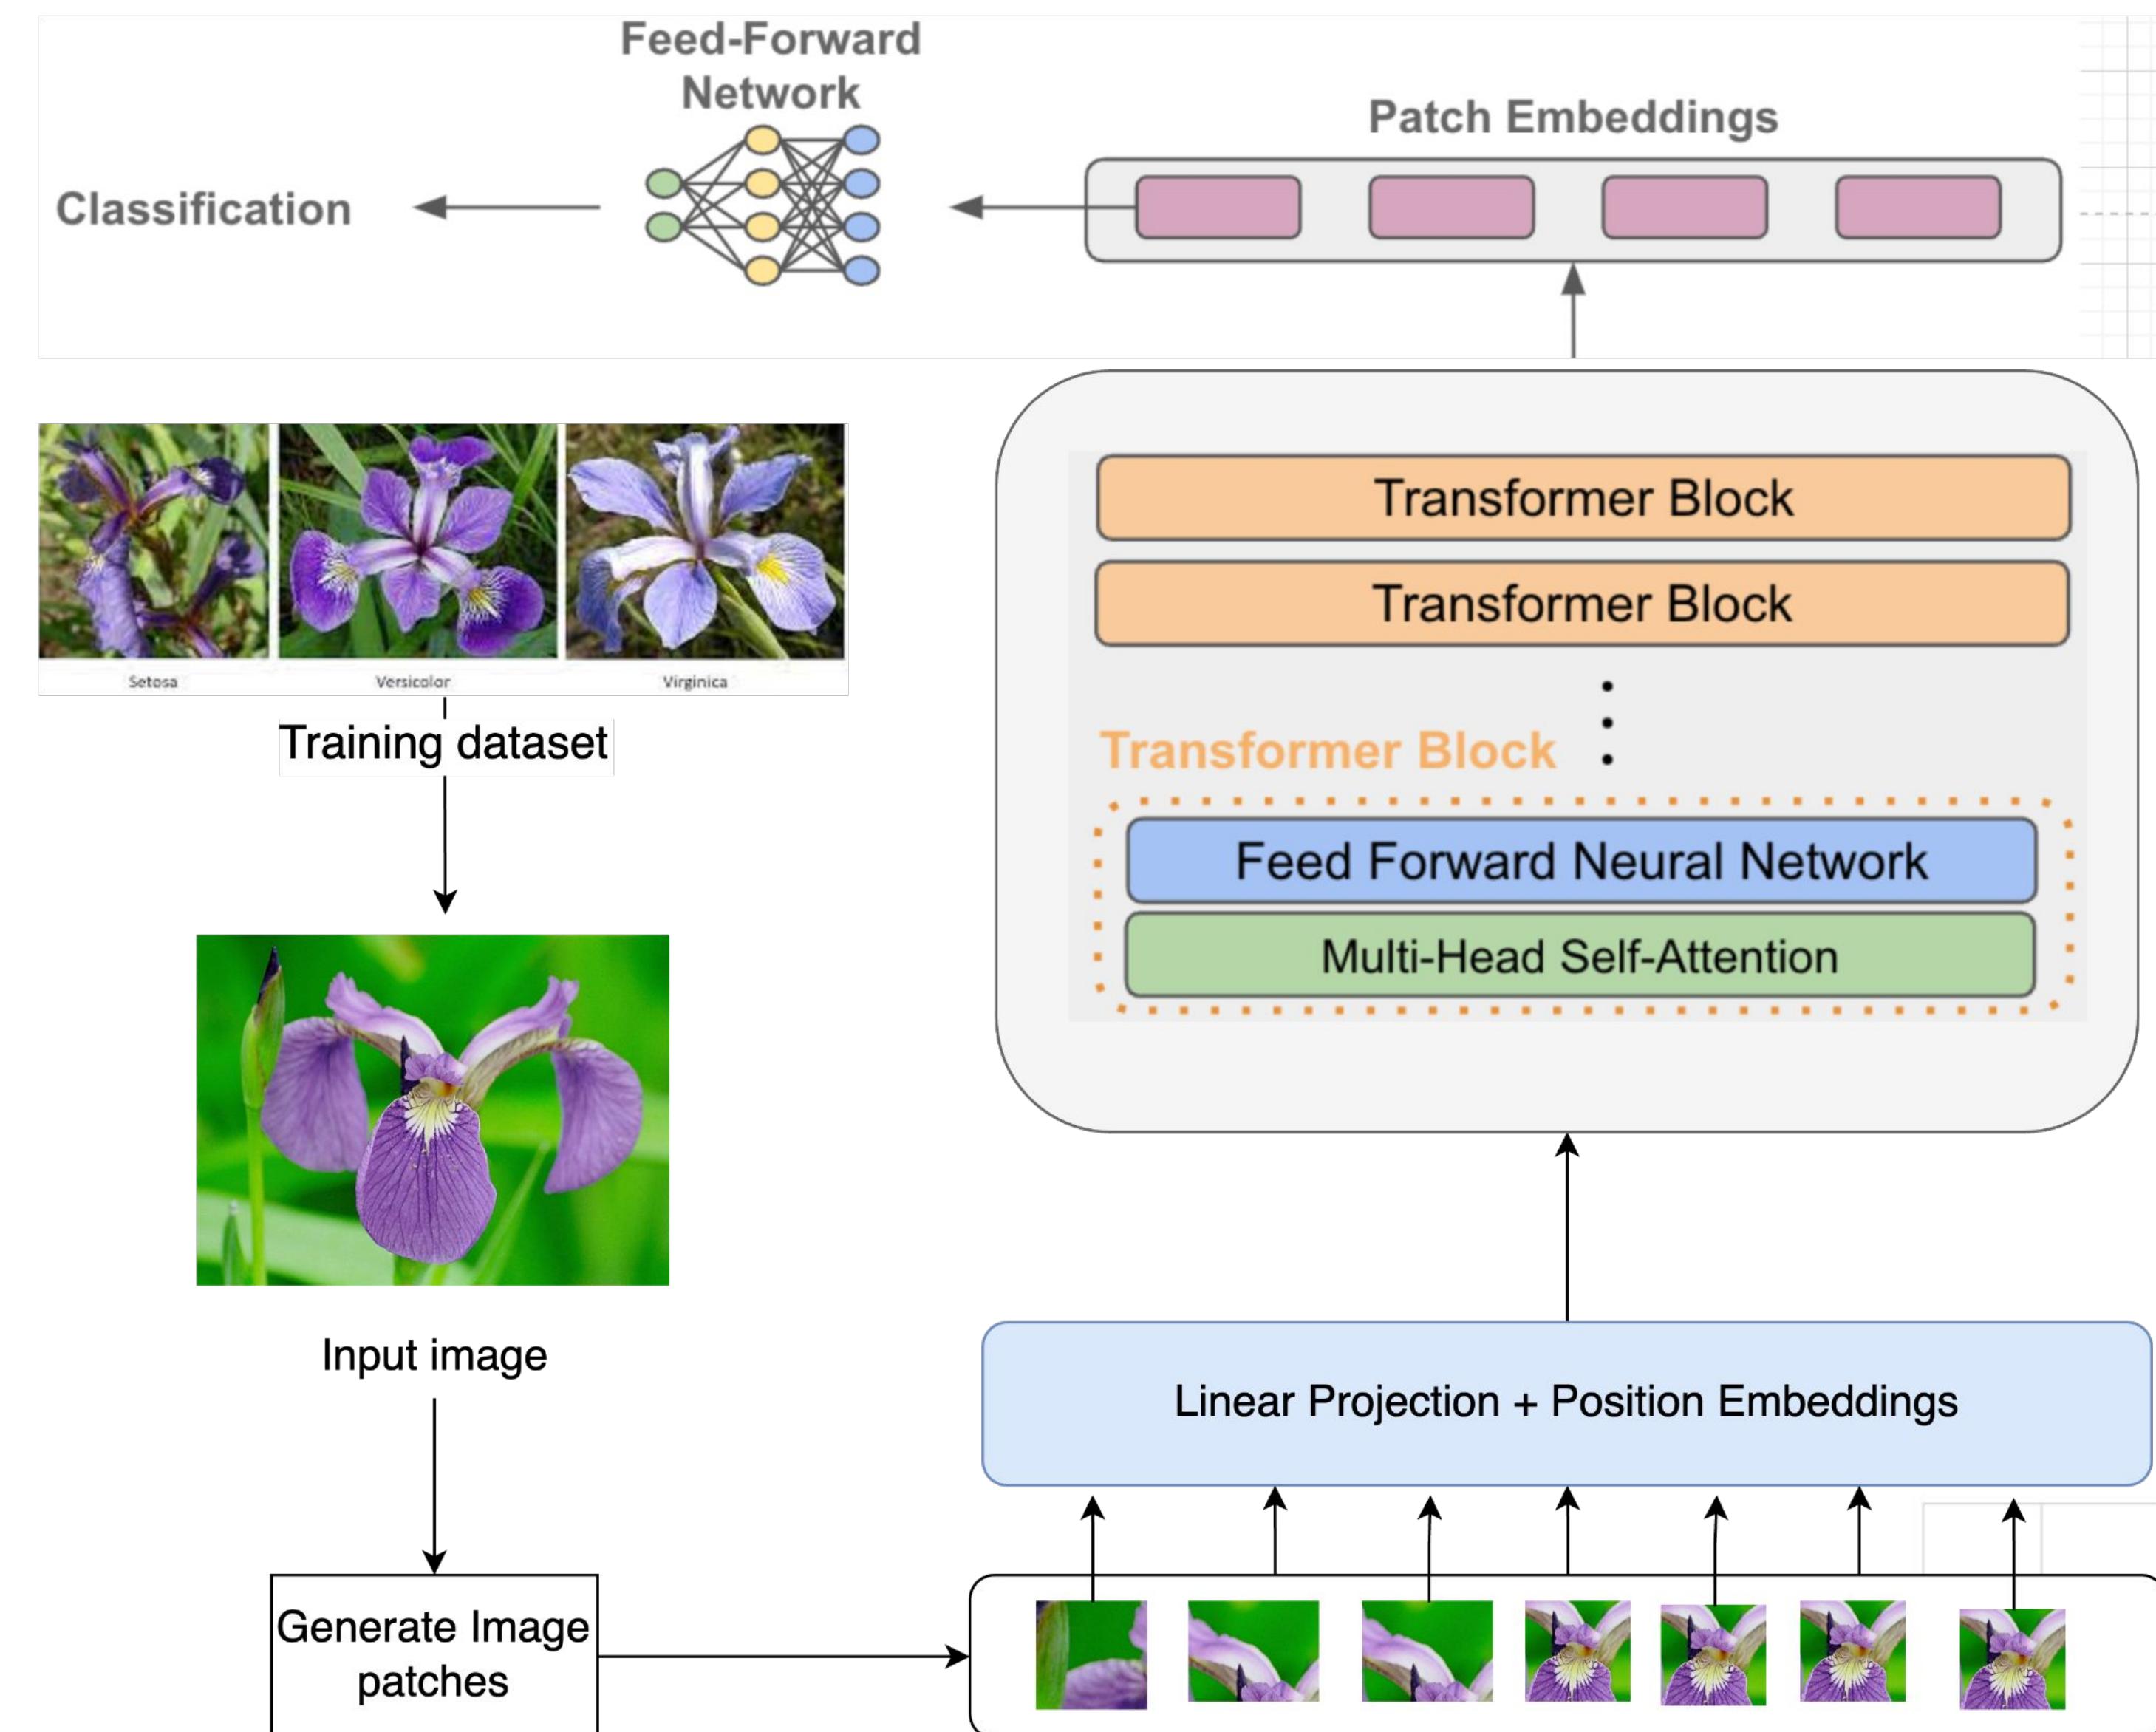
[2] MERLOT Reserve by R Zellers, J Lu, X Lu, Y Yu, Y Zhao, M Salehi, A Kusupati, J Hessel, A Farhadi, Y Choi

[3] VATT by H Akbari, L Yuan, R Qian, W-H Chuang, S-F Chang, Y Cui, B Gong



# Transforming images into sequences

1. Convert an image into a sequence of flattened image patches
2. Pass the image patch sequence through the transformer model
3. Transform the input image into the **embeddings representation**
4. Take the first element of the transformer's output sequence and pass it through a classification module



# Transformers

Transformers, In summary, pushed forward the field of **large language models and generative AI models** by addressing key limitations of traditional recurrent architectures, enabling more efficient processing of long-range dependencies, scalability to large datasets, and superior performance on various NLP tasks through pre-training and fine-tuning strategies.

Transformer is the backbone architecture for many state-of-the-art models, such as Gemini, GPT-4, DALL-E-2, Codex, and Gopher.

**It was first proposed to solve the limitations of traditional models such as RNNs in handling variable-length sequences and context awareness.**



# Generative AI (GAI)

Generative AI is a type of artificial intelligence (AI) that can create new content, such as text, images, and music.

It does this by learning from existing data and then using that data to generate new outputs. The results after training is a probabilistic model that is able to generate new data.

The artificially generated data is expected to have the same statistical properties and structure as real-world data or our training dataset distribution.



# How do these type of ML models work?

In technical terms, the concept of AI-guided content generation (AIGC) **involves utilizing advanced algorithms in generative AI to produce content that aligns with human instructions.**

This process typically entails two key stages: first,

1. extracting the intended information from the given instructions
2. Transform that information into context in form of embedding that can be passed to the foundation model
3. generating content based on the identified intentions.

By employing GAI algorithms, **AIGC enables the model to learn and be guided by human input**, resulting in the creation of content that fulfills the specified criteria.

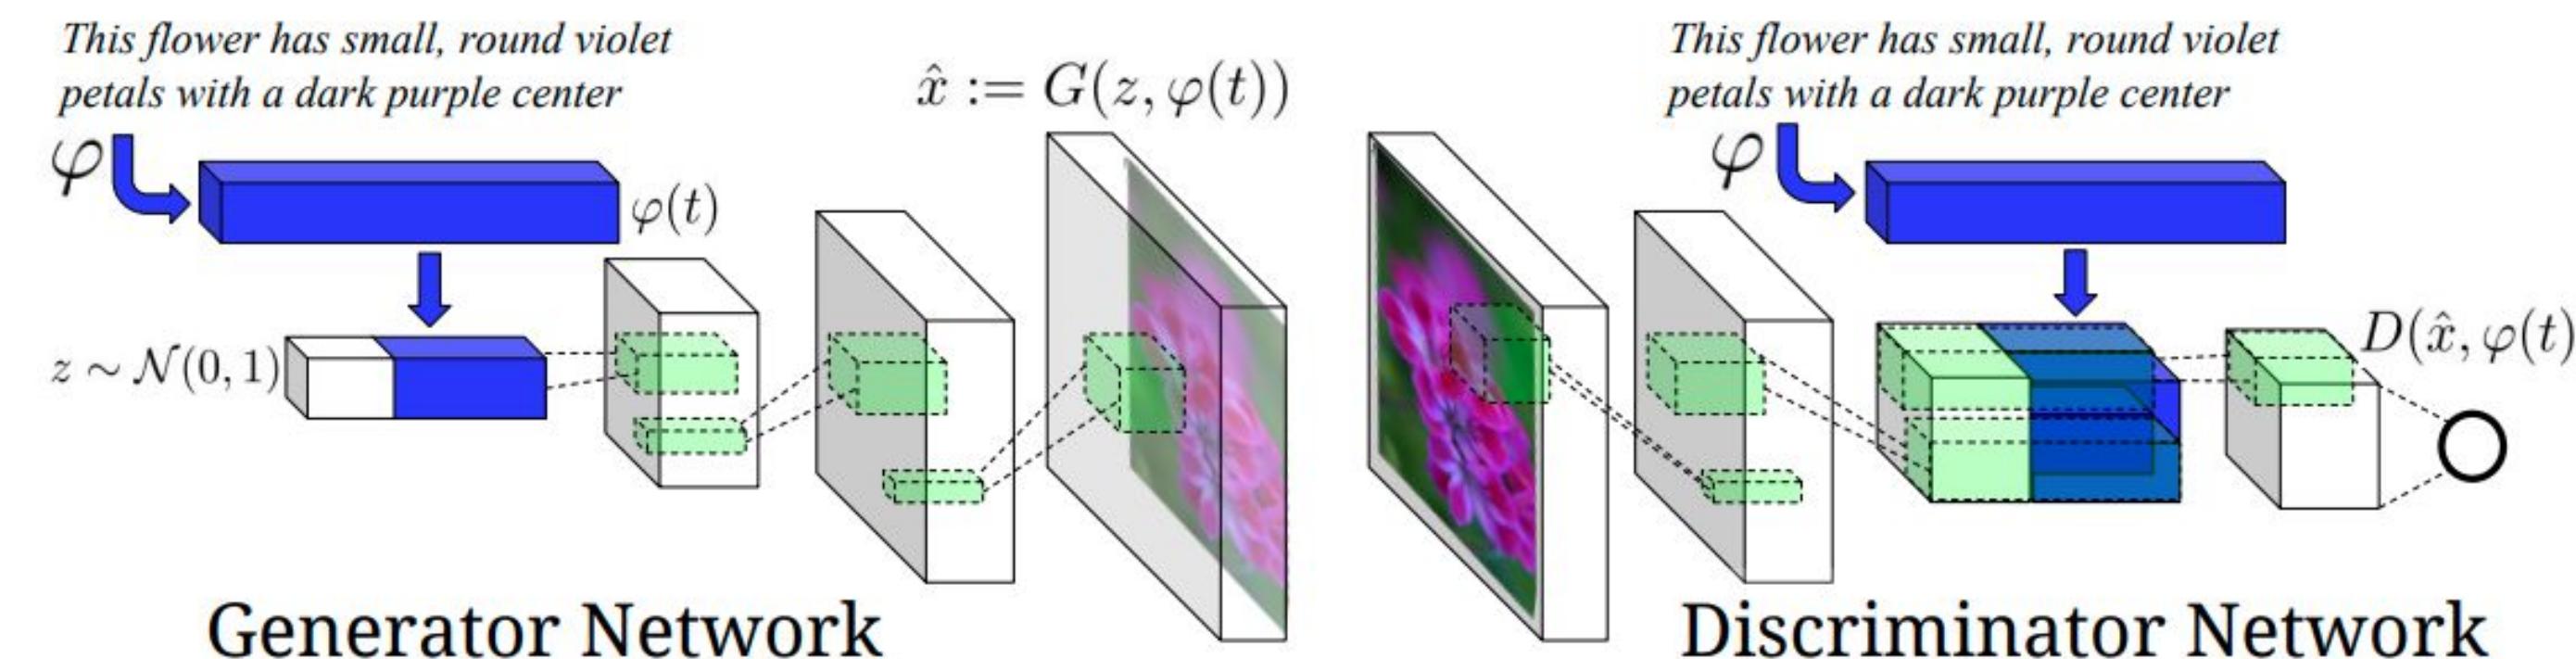


Figure 2. Our text-conditional convolutional GAN architecture. Text encoding  $\varphi(t)$  is used by both generator and discriminator. It is projected to a lower-dimensions and depth concatenated with image feature maps for further stages of convolutional processing.

## Generative Adversarial Text to Image Synthesis

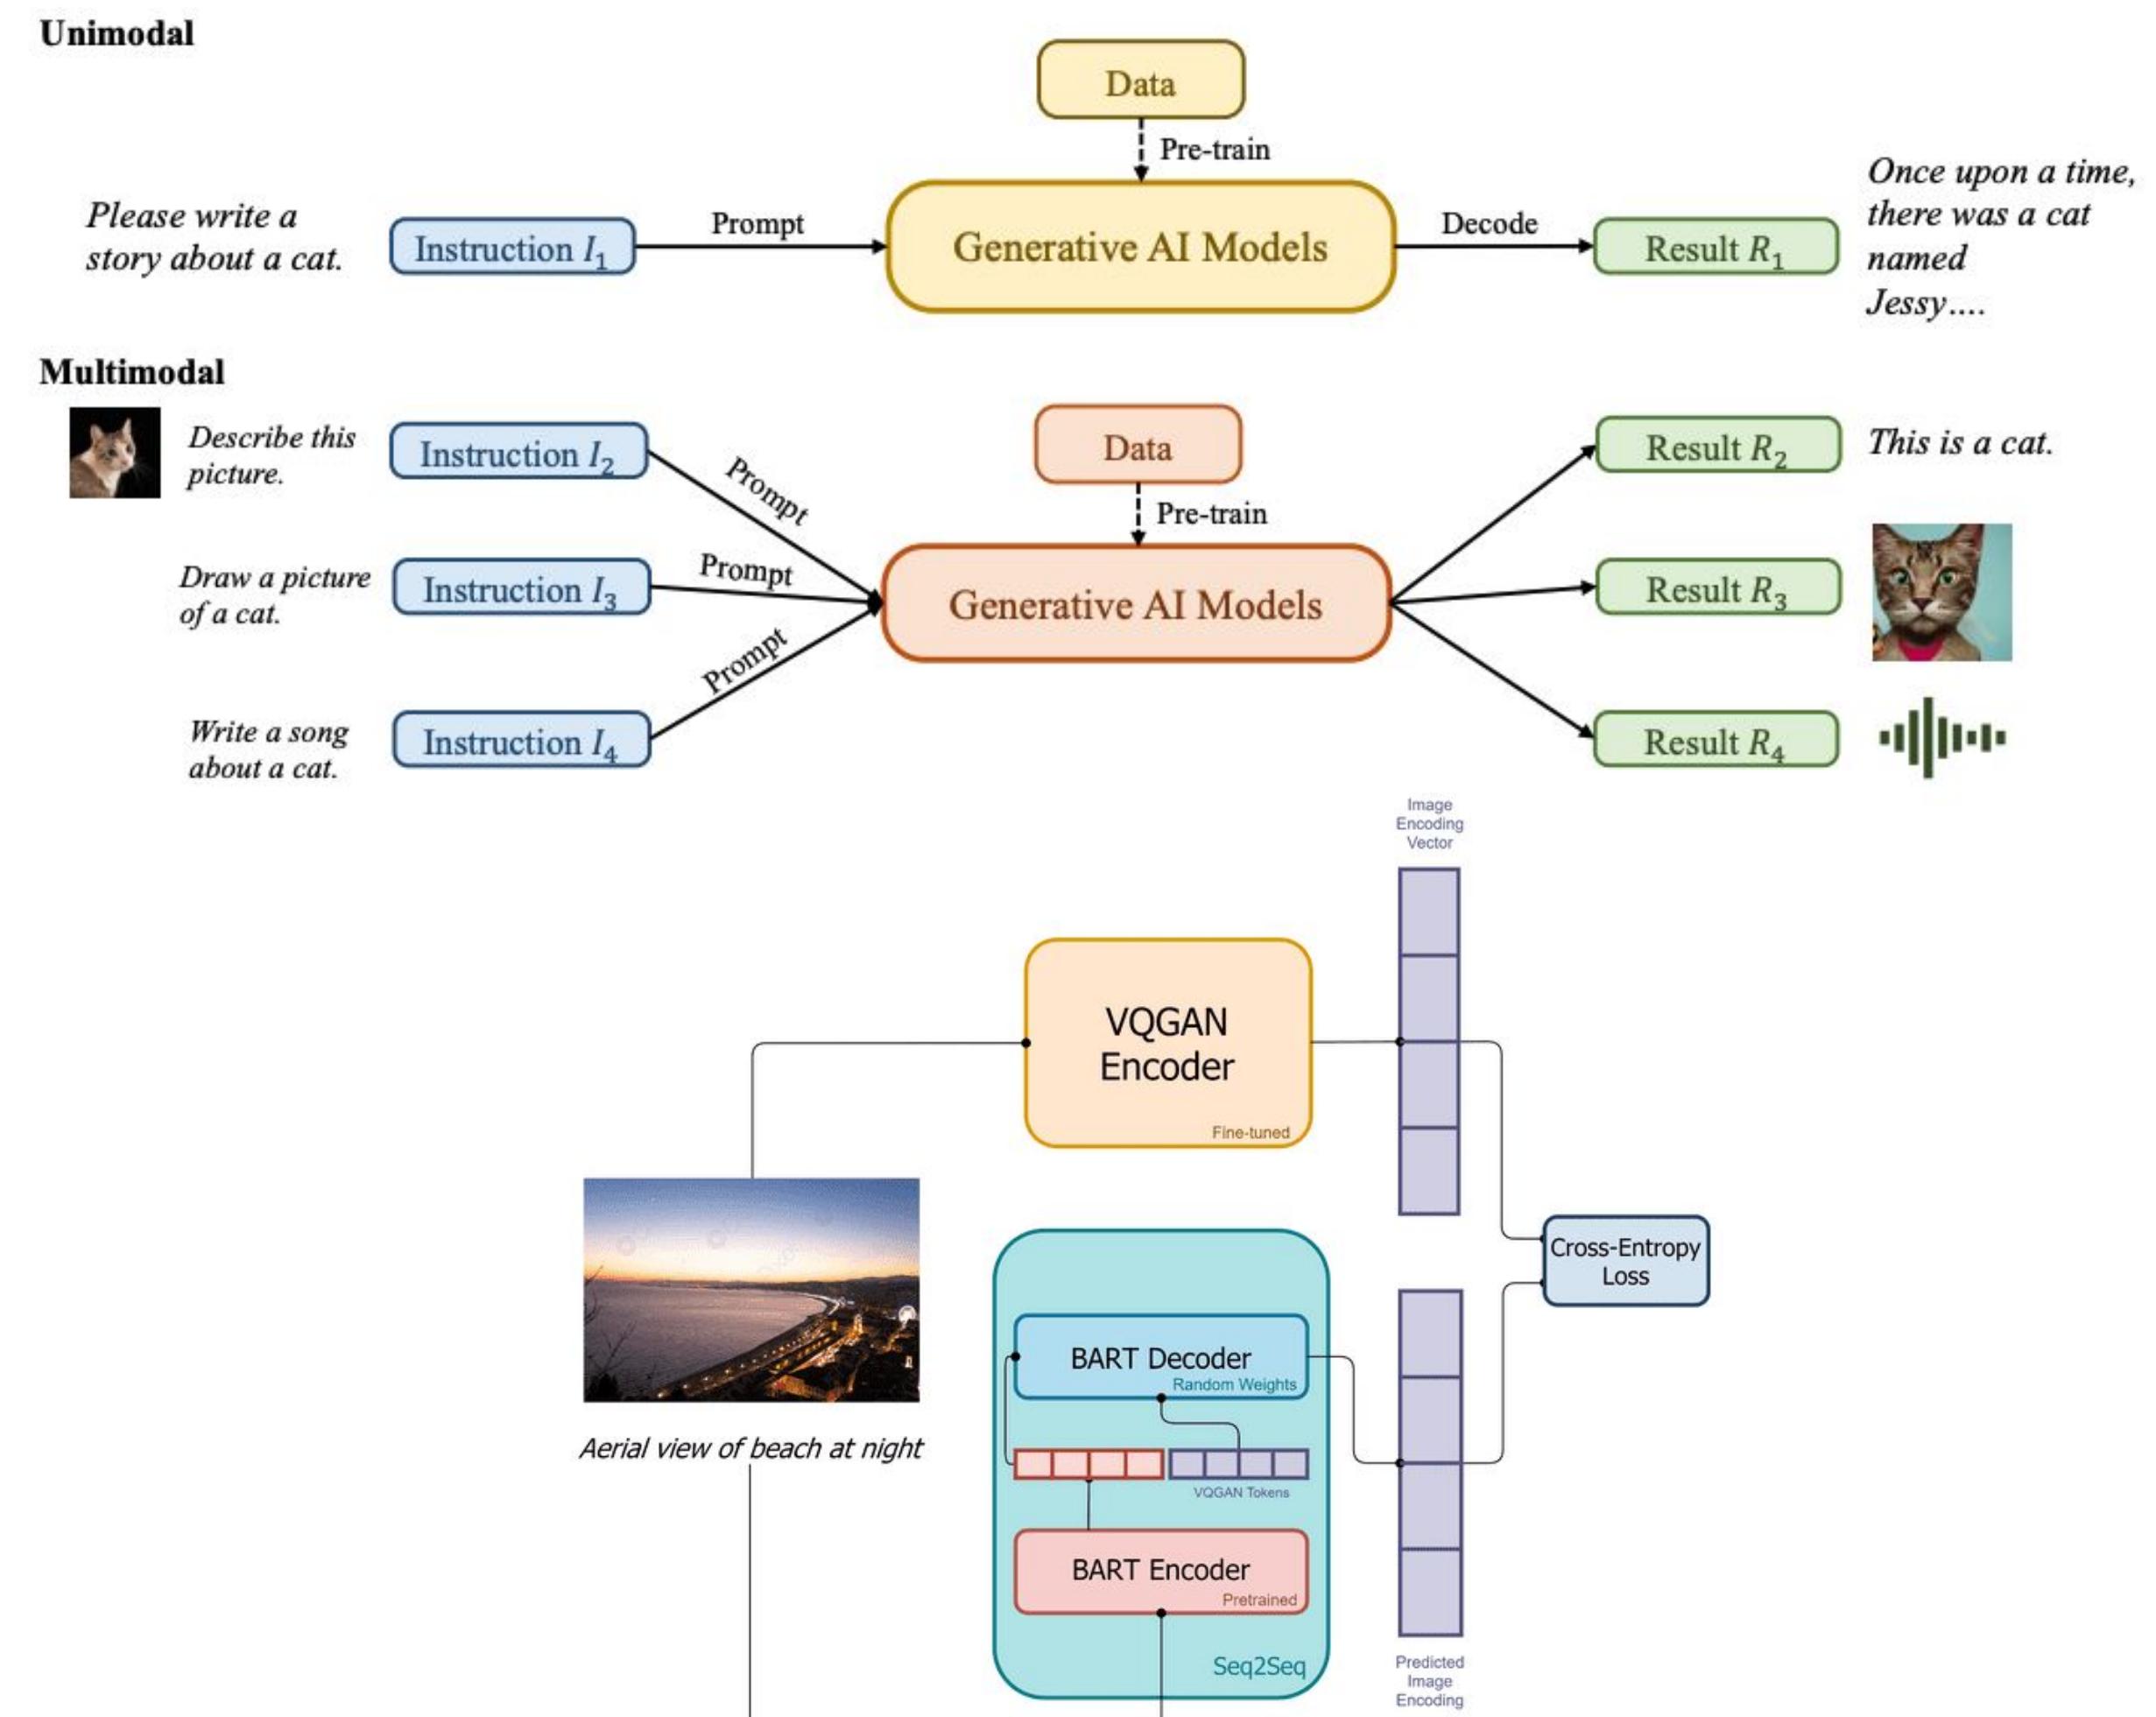
# Unimodal vs. Multimodal

Generative models can be categorized into two types: unimodal models and multimodal models.

**Unimodal models receive instructions and generate content in the same modality.** For example, a unimodal text-to-text model would receive instructions in the form of text and generate text as output.

**Multimodal models receive instructions and generate content in different modalities.** For example, a multimodal image-to-text model would receive instructions in the form of an image and generate text as output.

Multimodal models are more complex than unimodal models, but they can also generate more creative and engaging content.



# What are LLMs

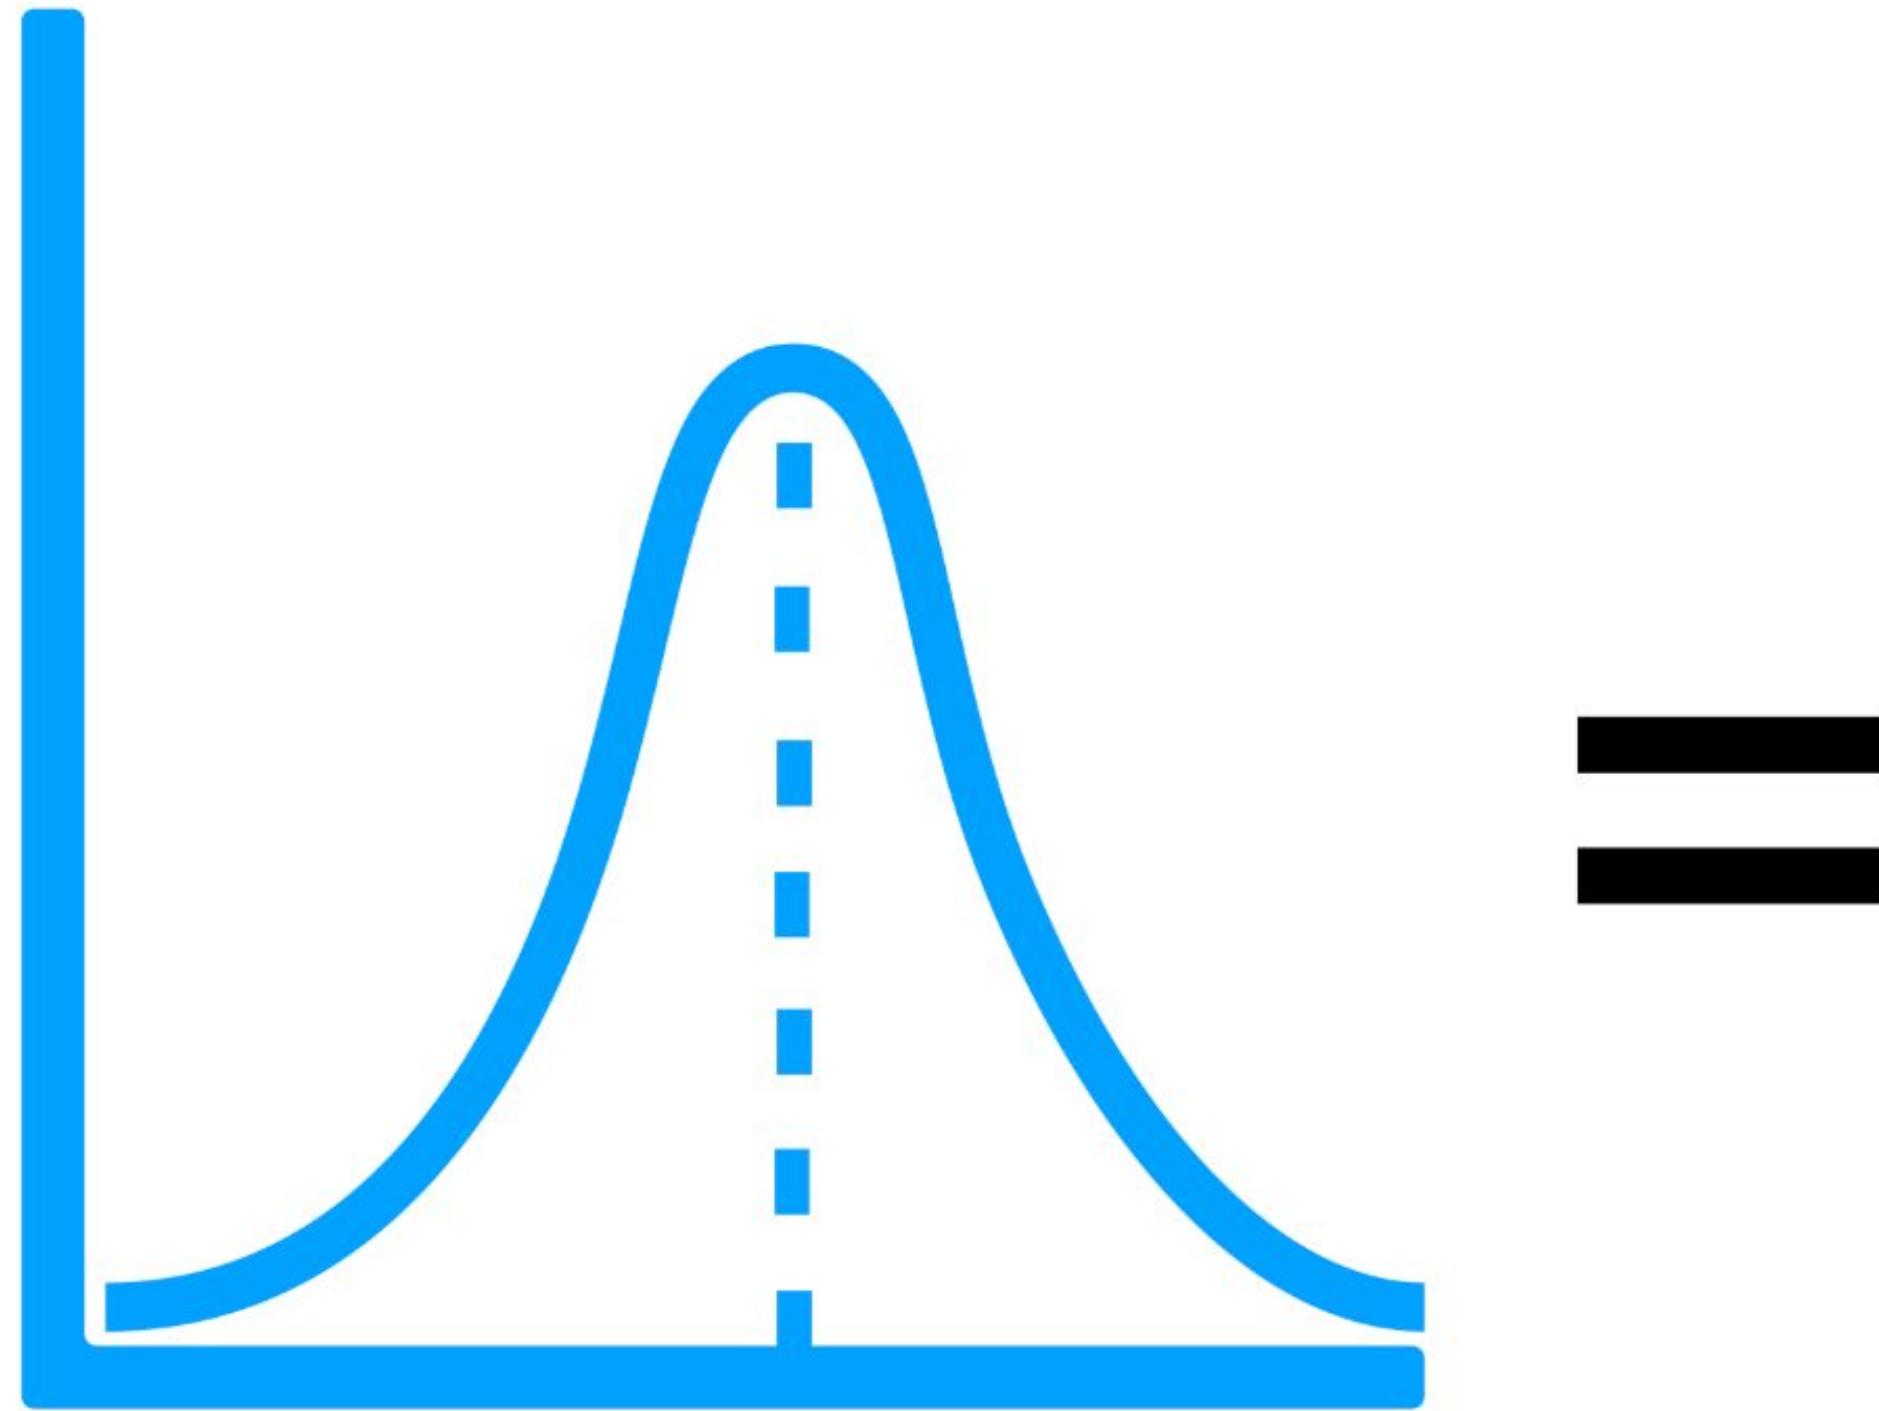
## High Level:

A type of Generative AI architecture, that usually means a >1B parameter **transformer-based model** trained to predict the next token or sequence of tokens in a sequence

## Low Level:

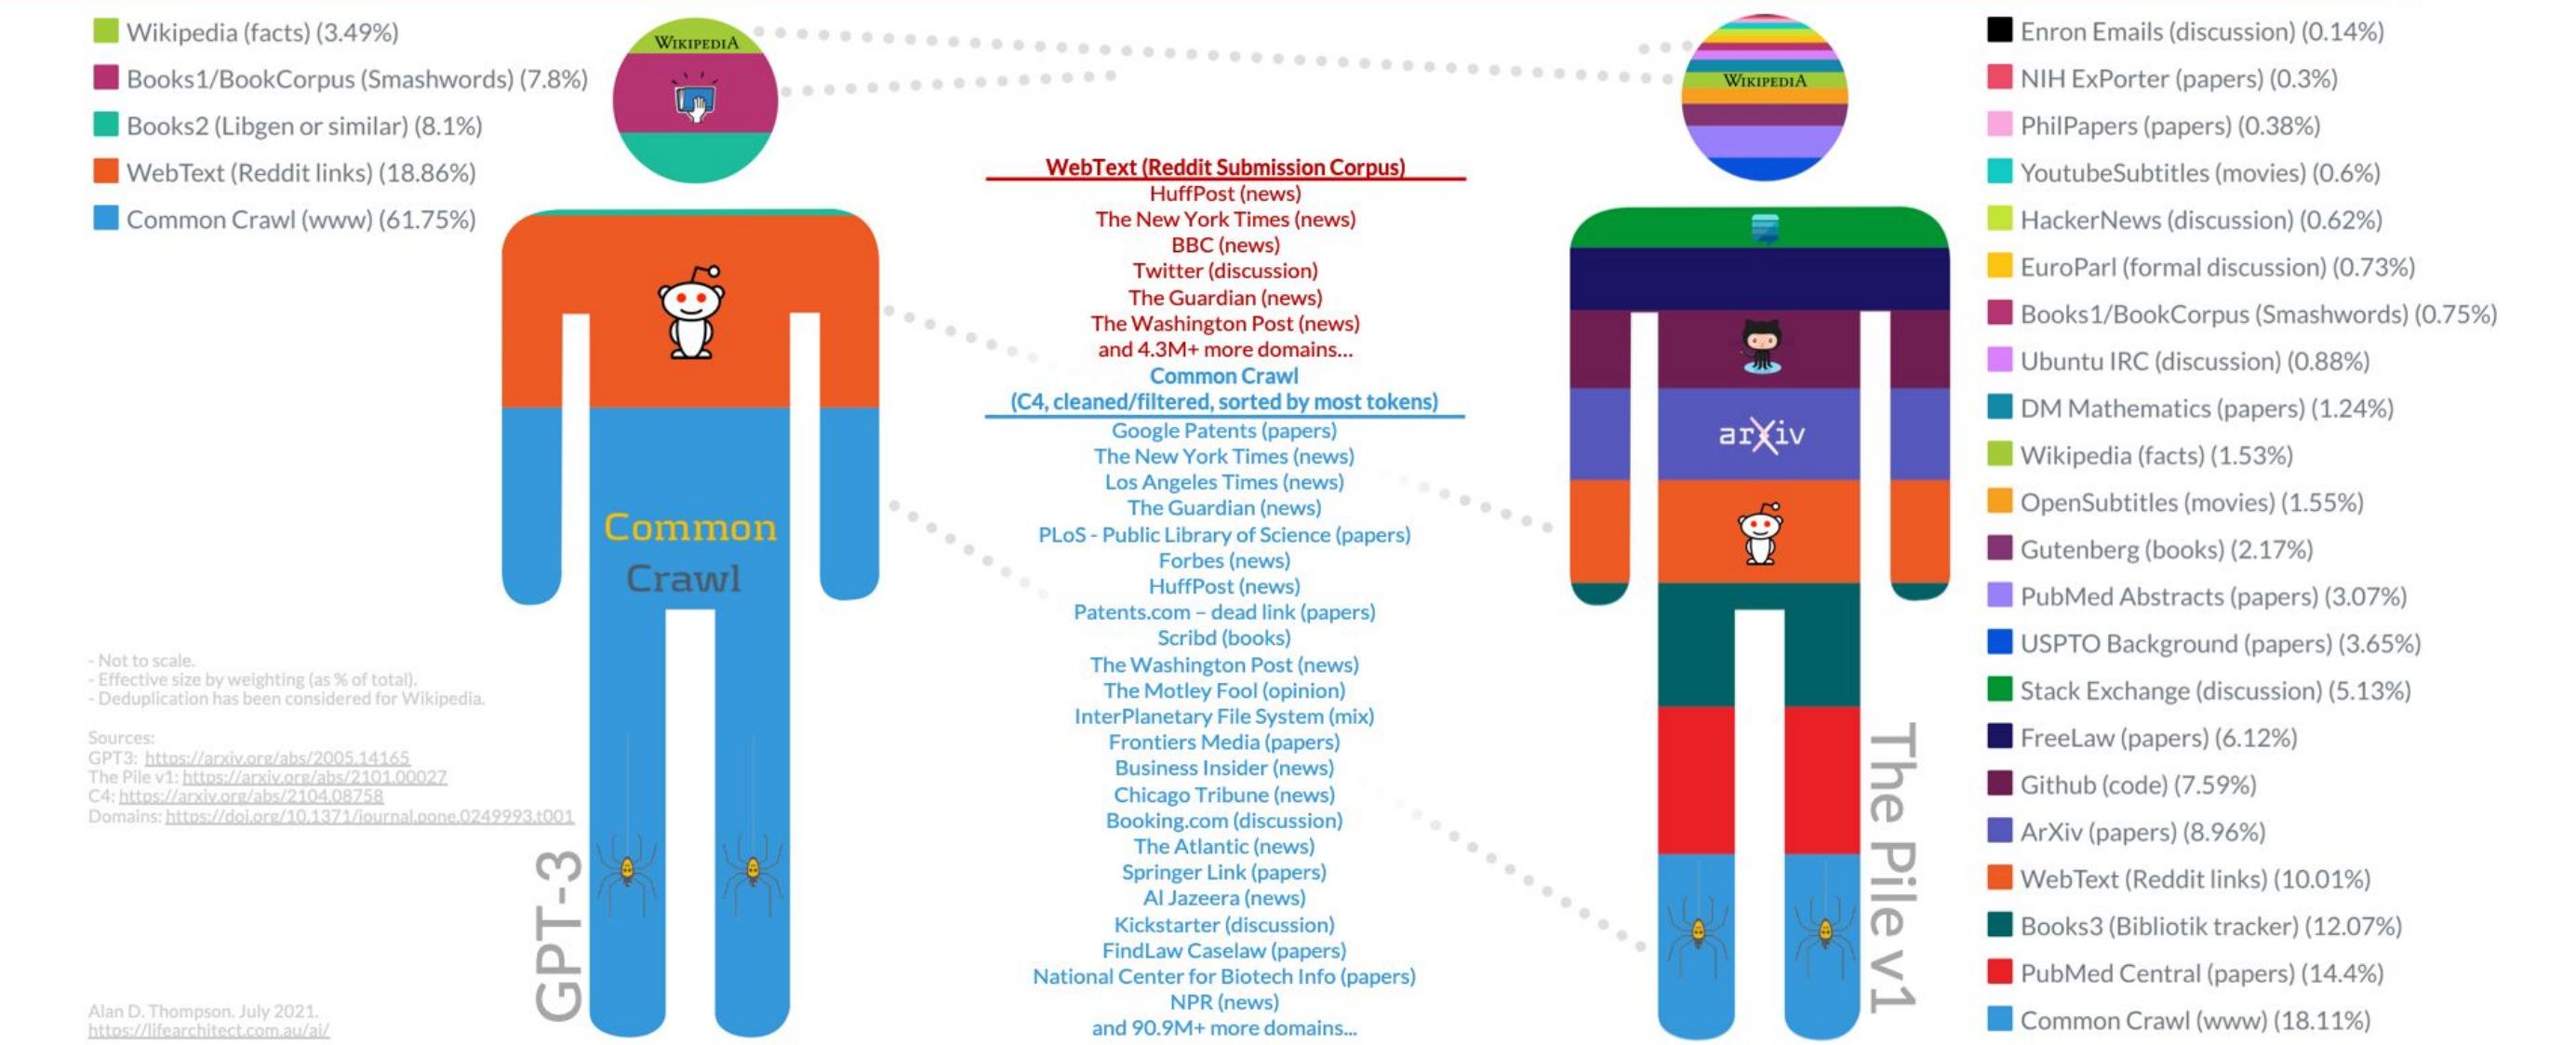
2 files in disk: weights and model

LMs are “just” statistical models of text

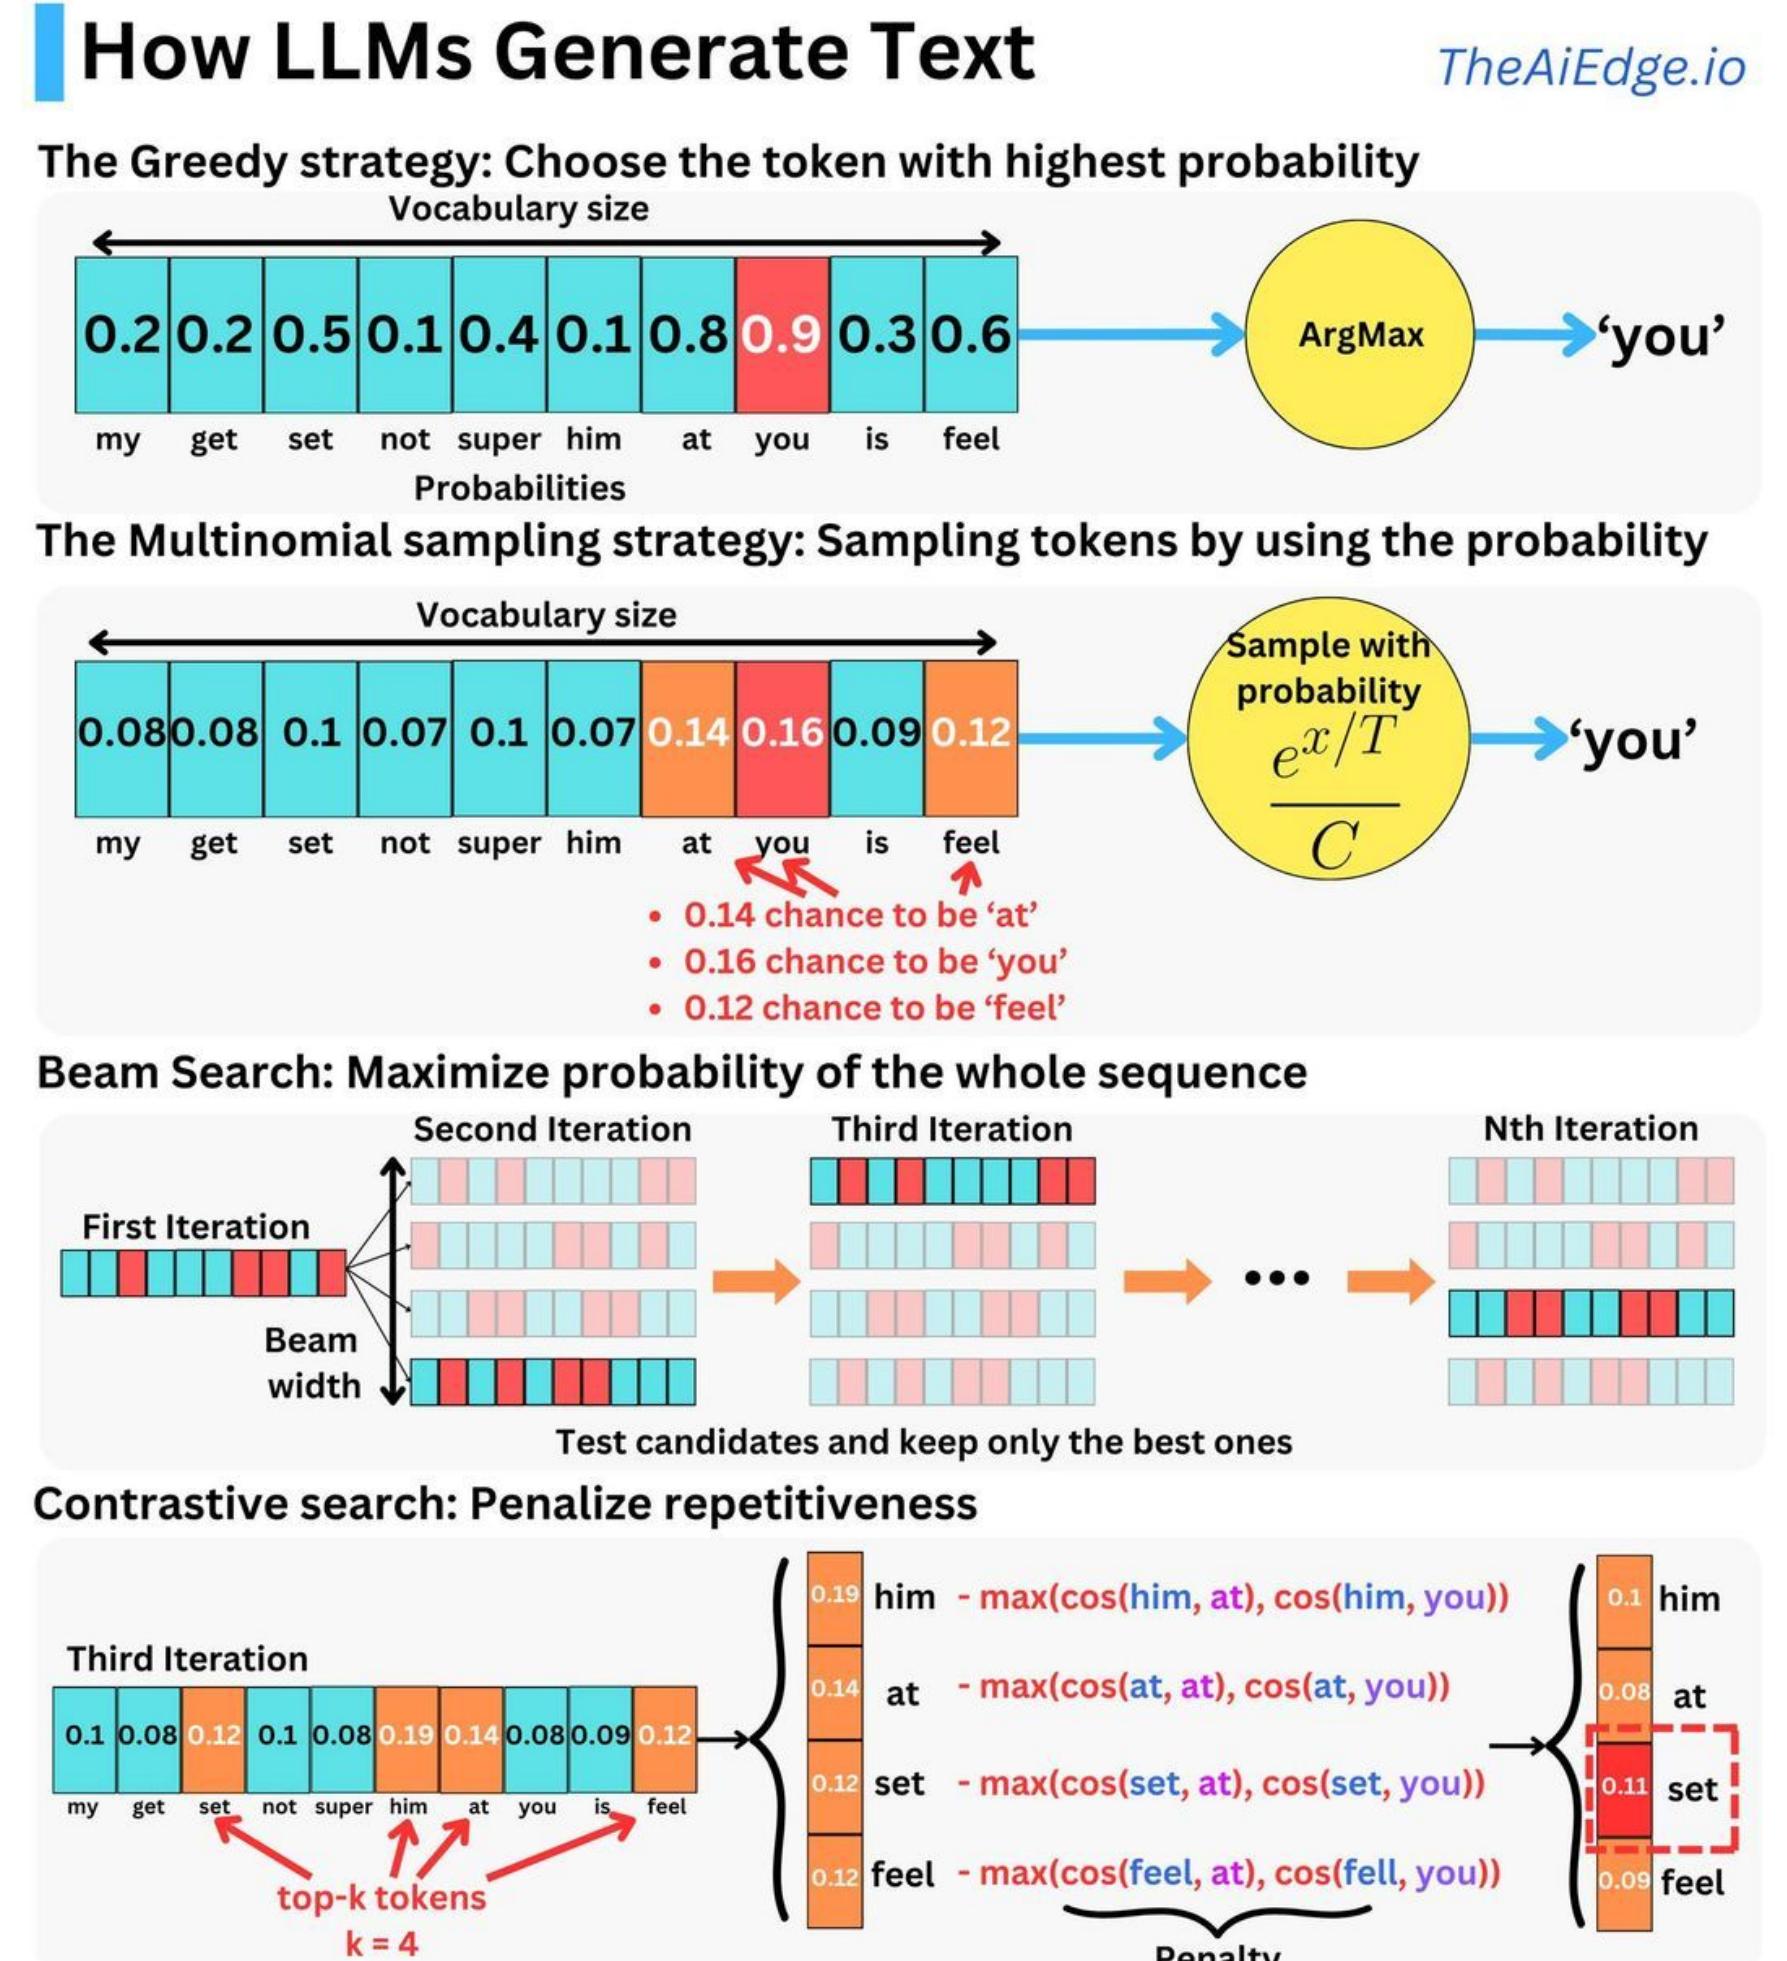
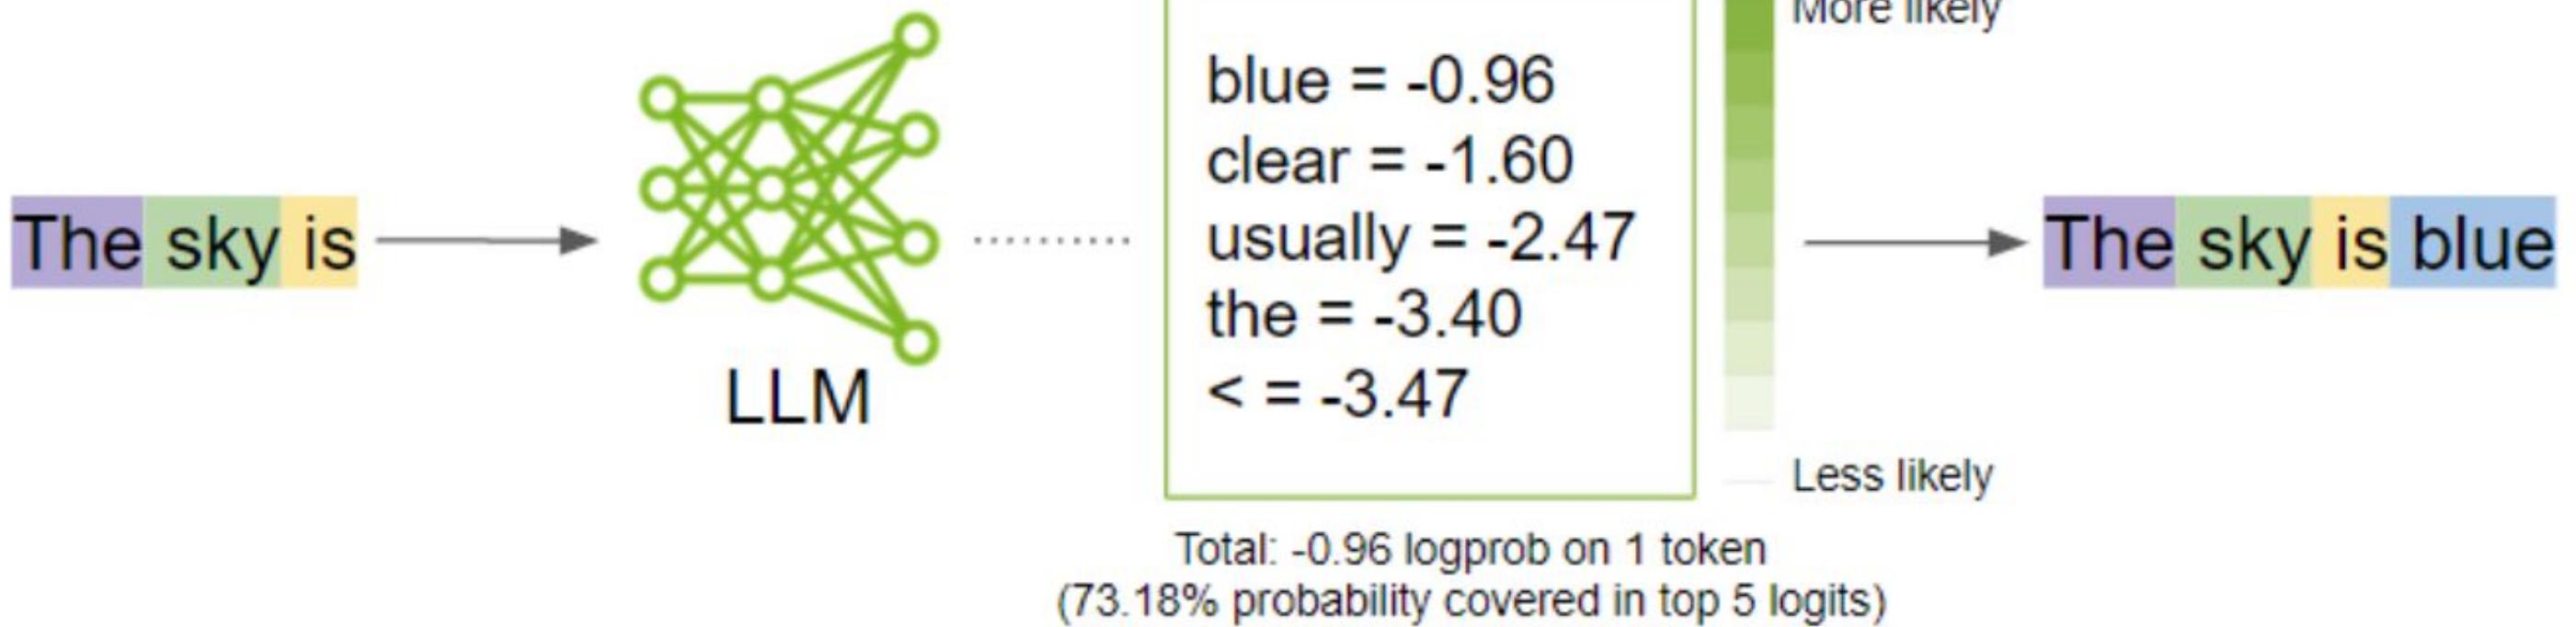


They assign a **probability** to every suffix of a prompt,  
token-by-token.

# How does LIMs work?



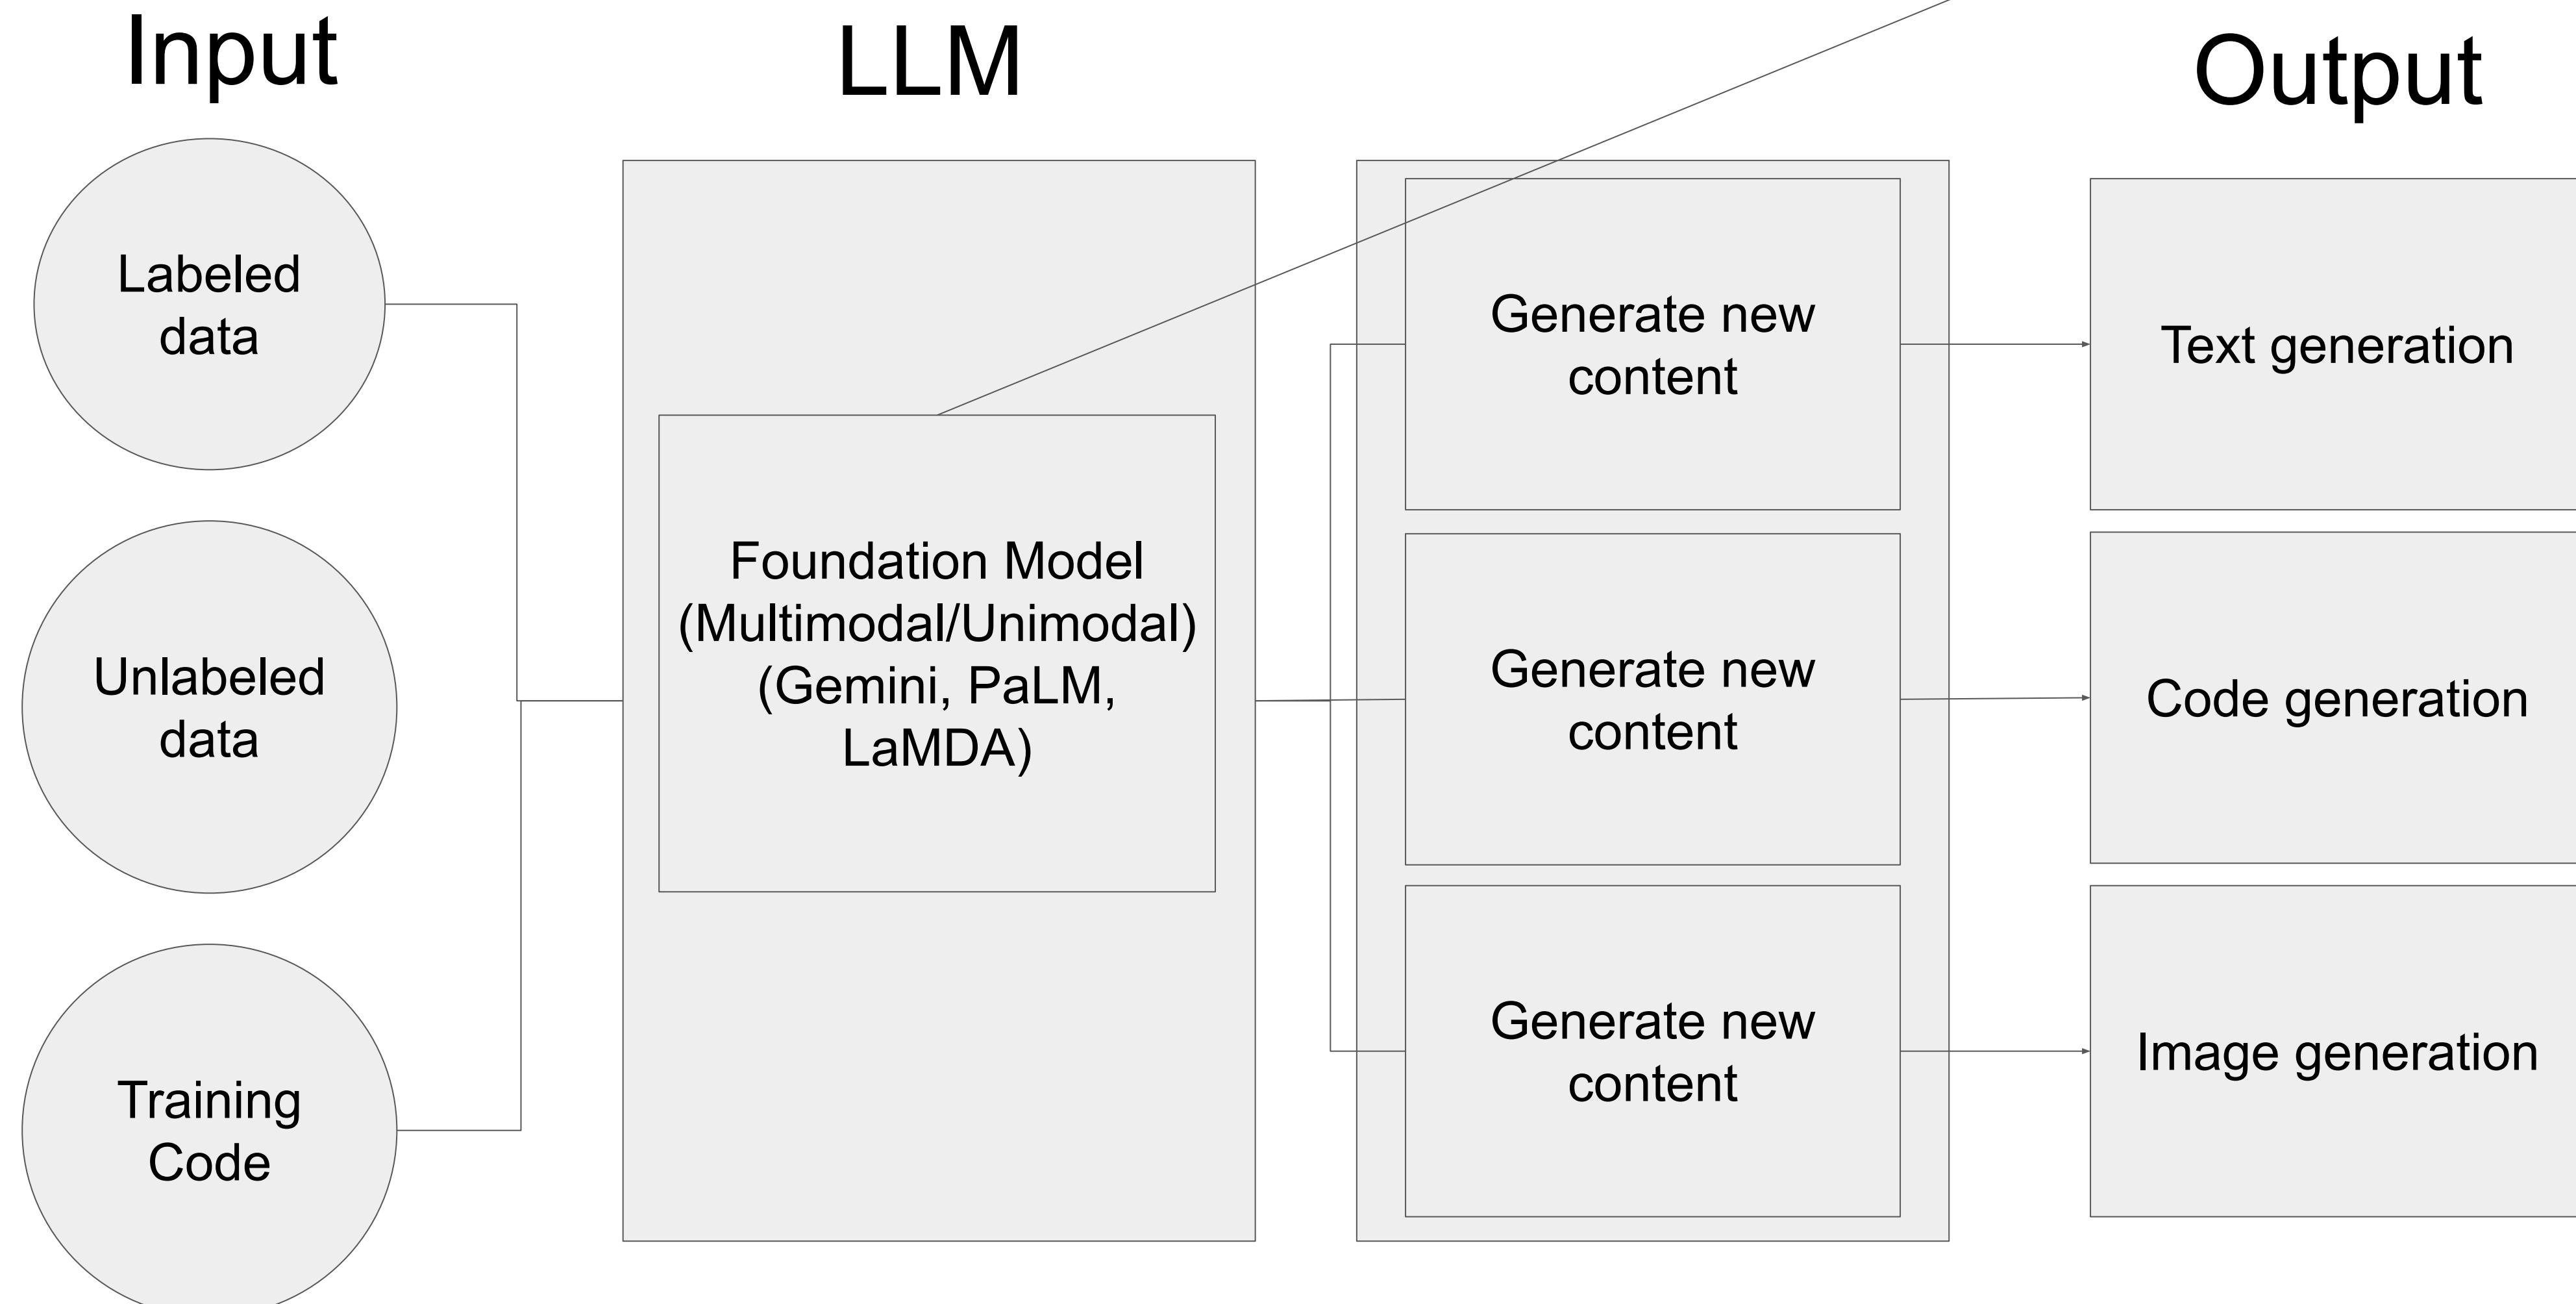
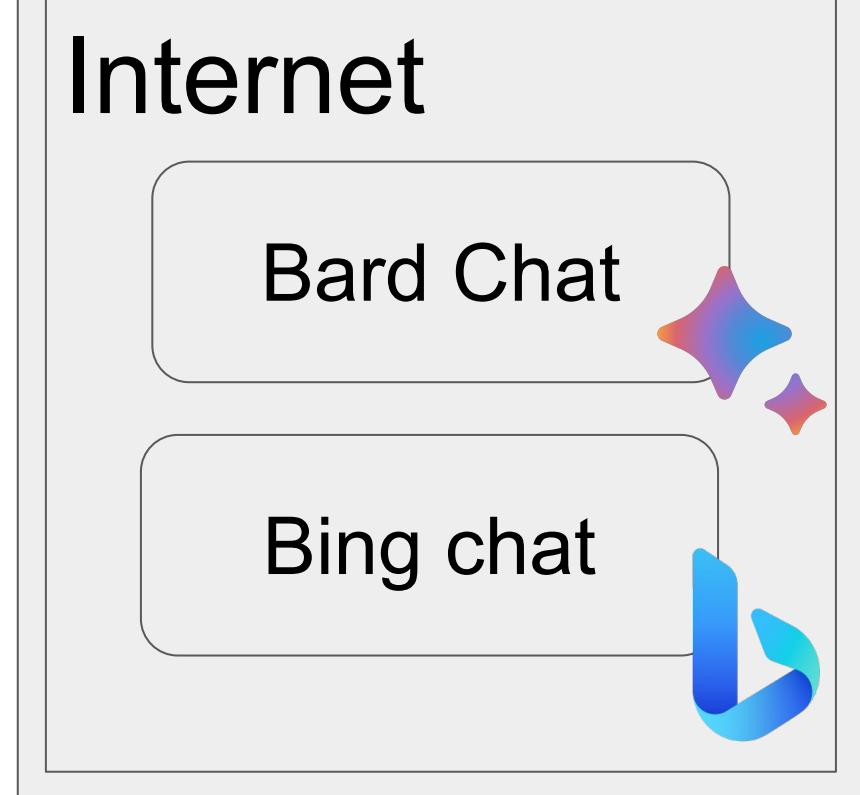
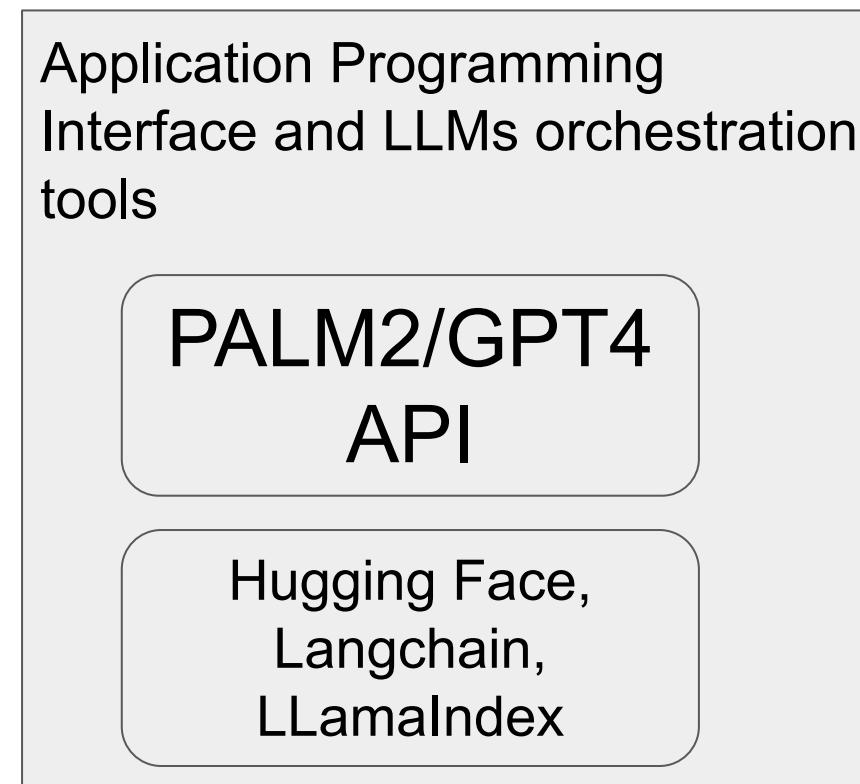
[LifeArchitect.ai/models](#)



Credits to: <https://newsletter.theaiedge.io/about>

# LLM APP pipeline

## Interface



# Google is the industry pioneer in AI. Foundations models evolution



2015

Google DeepMind  
AlphaGo defeats Go champion

2016

Google's DeepMind helps detect eye disease

2017

Google invents Transformer kickstarting LLM revolution

2018

Google's groundbreaking large language model, BERT

2019

Text-to-Text Transfer Transformer  
LLM 10B P Model Open Sourced

2020

Google LaMDA Model Trained to converse

2022

AlphaFold predicts structures of all known proteins

2023

A conversational AI Service powered by PaLM2

2024

Family of multimodal LLMs & products

3,000

Researchers

7,000

Publications

## Responsible AI

✓ Built & Tested for Safety

✓ Privacy in design

✓ Upholds high scientific standards

✓ Accountable to People

✓ Socially Beneficial

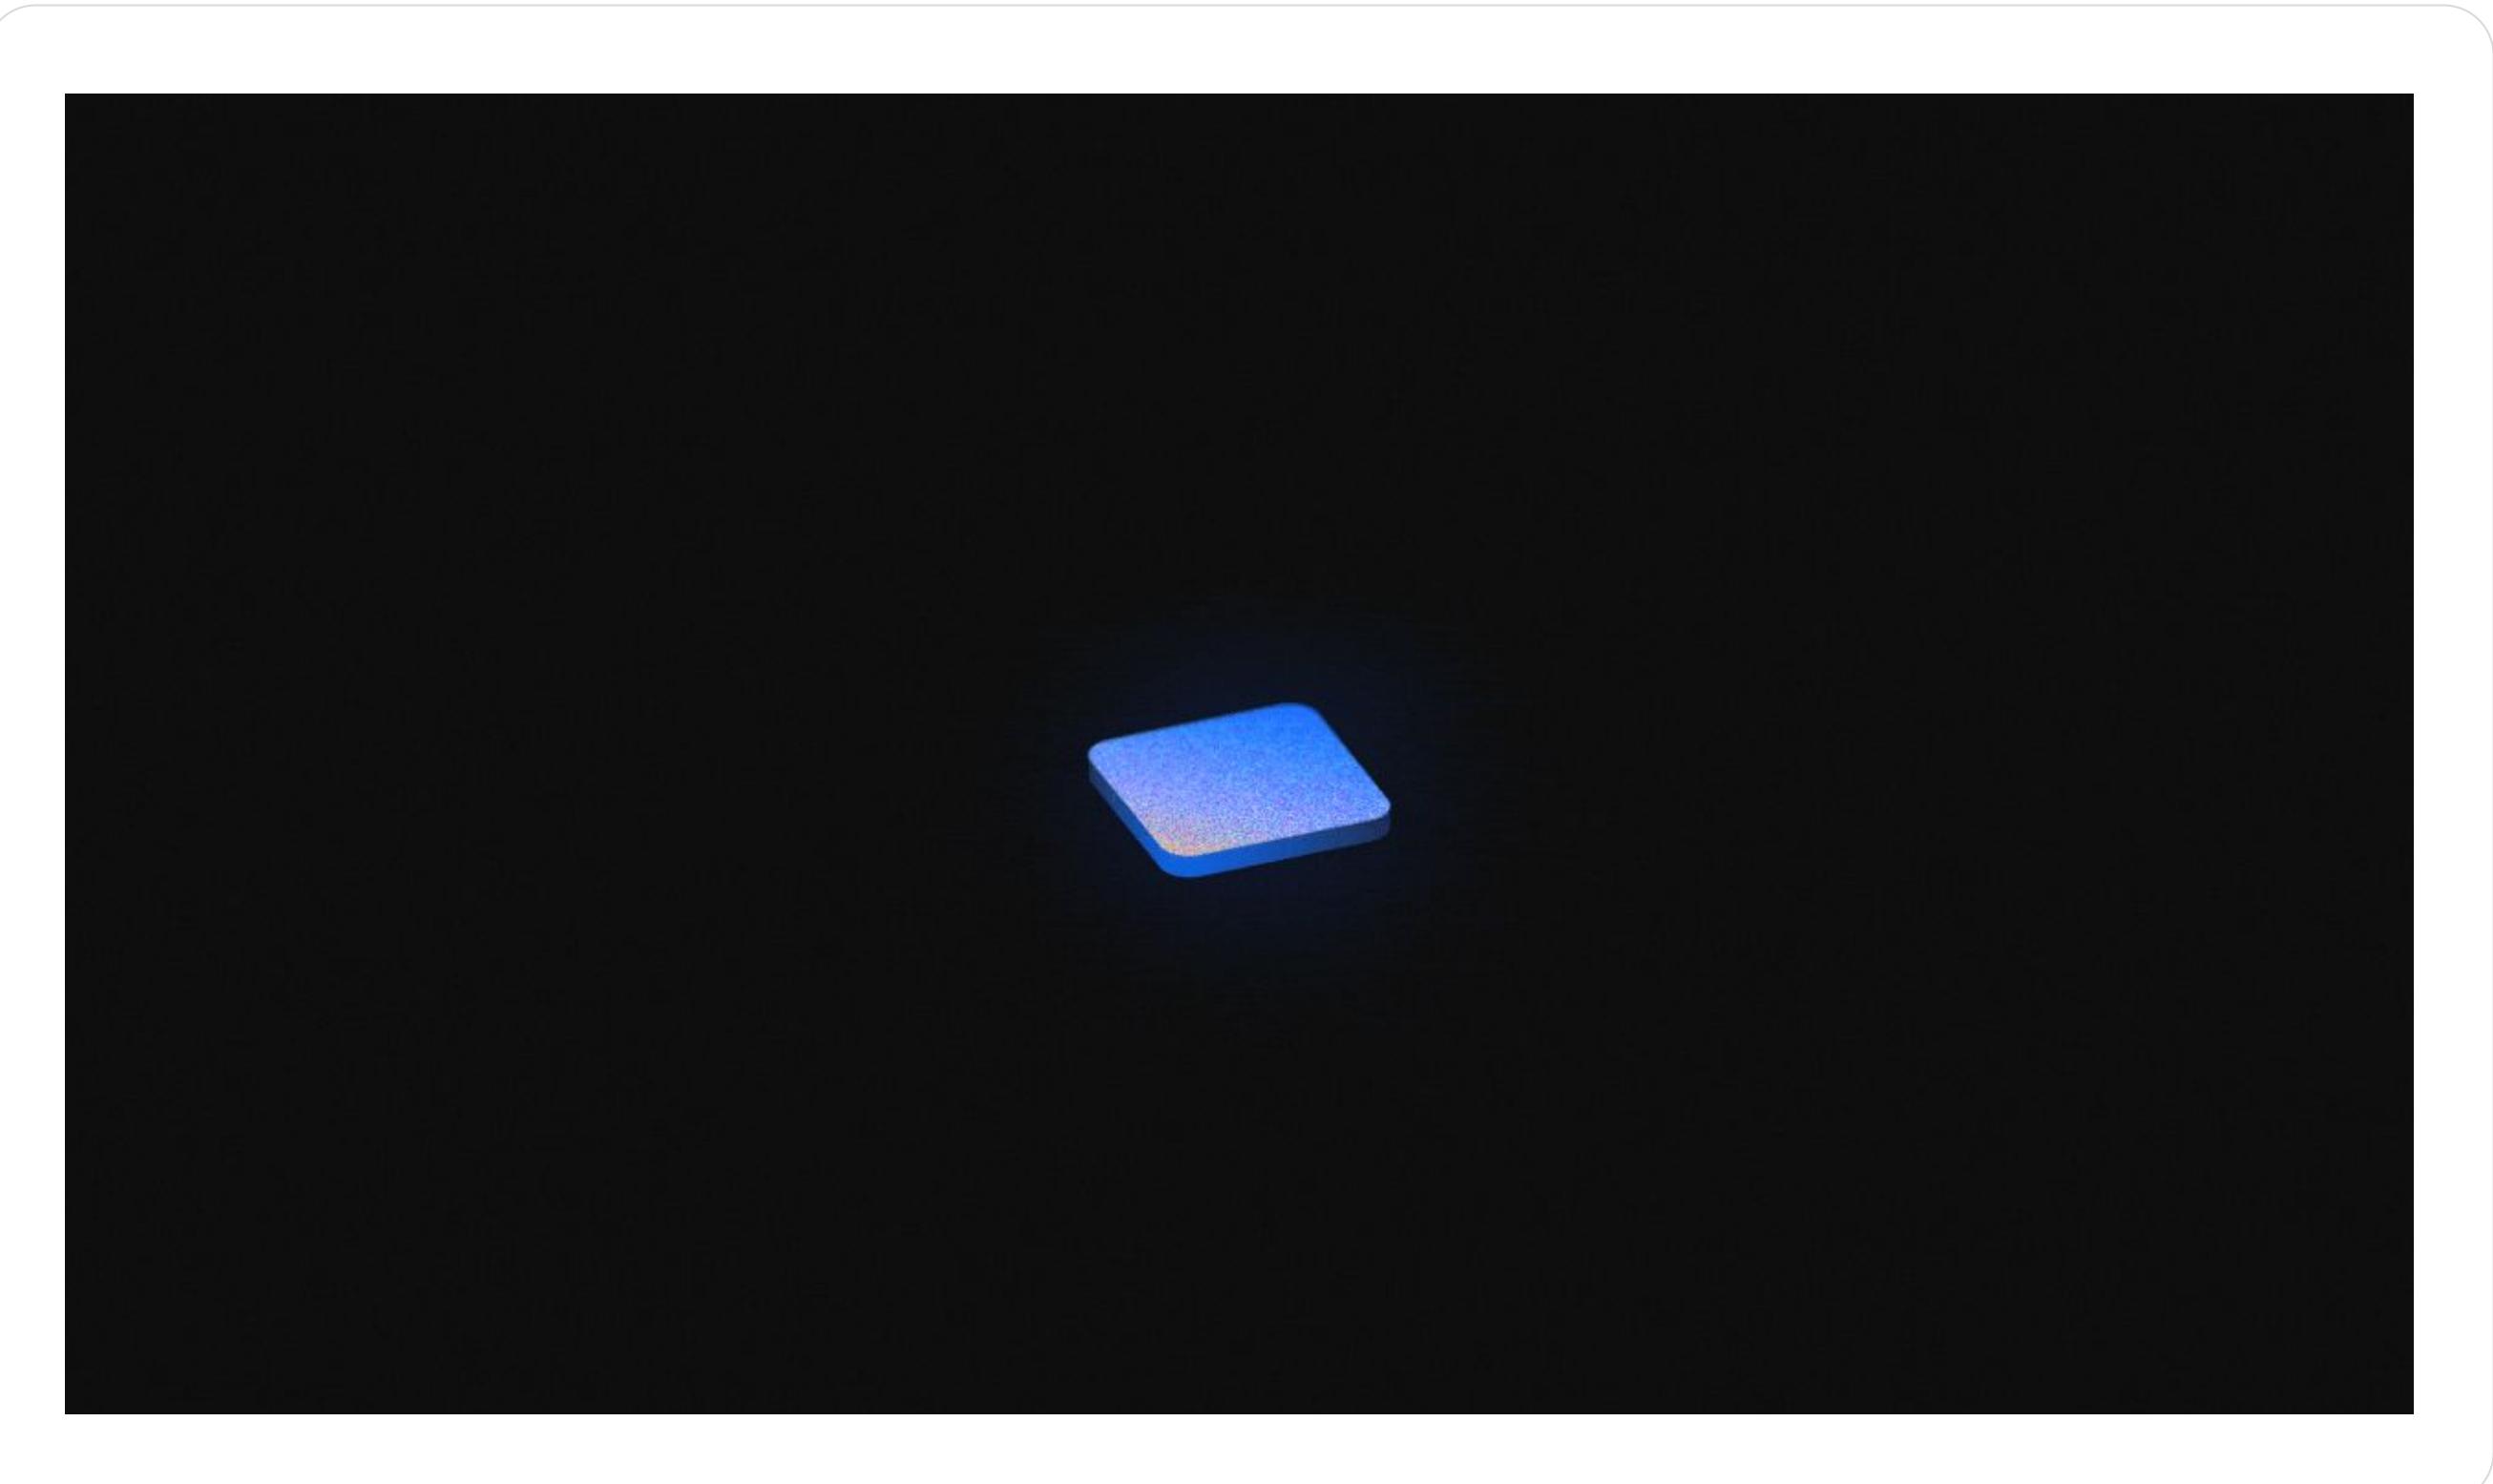
✓ Avoid creating unfair bias

Welcome to  
the Gemini era

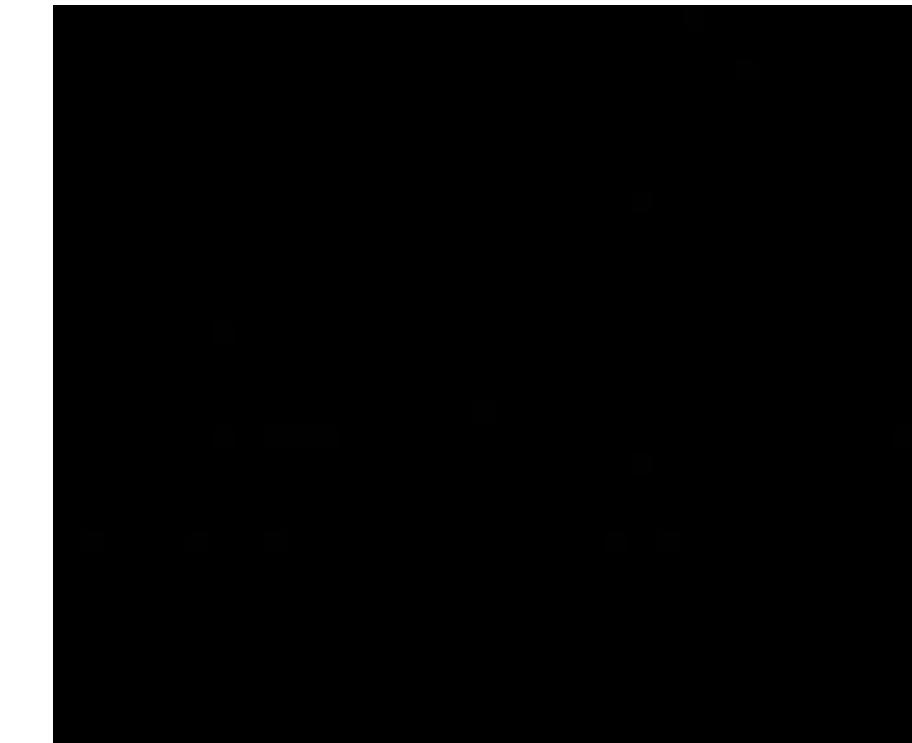
# The next chapter of **Generative AI** innovation



Gemini is the most capable and general model we've ever built, and is the result of a large-scale collaborative effort by teams across Google, including Google DeepMind and Google Research.



# Gemini marks the next phase on our journey to making AI more helpful for everyone and provides



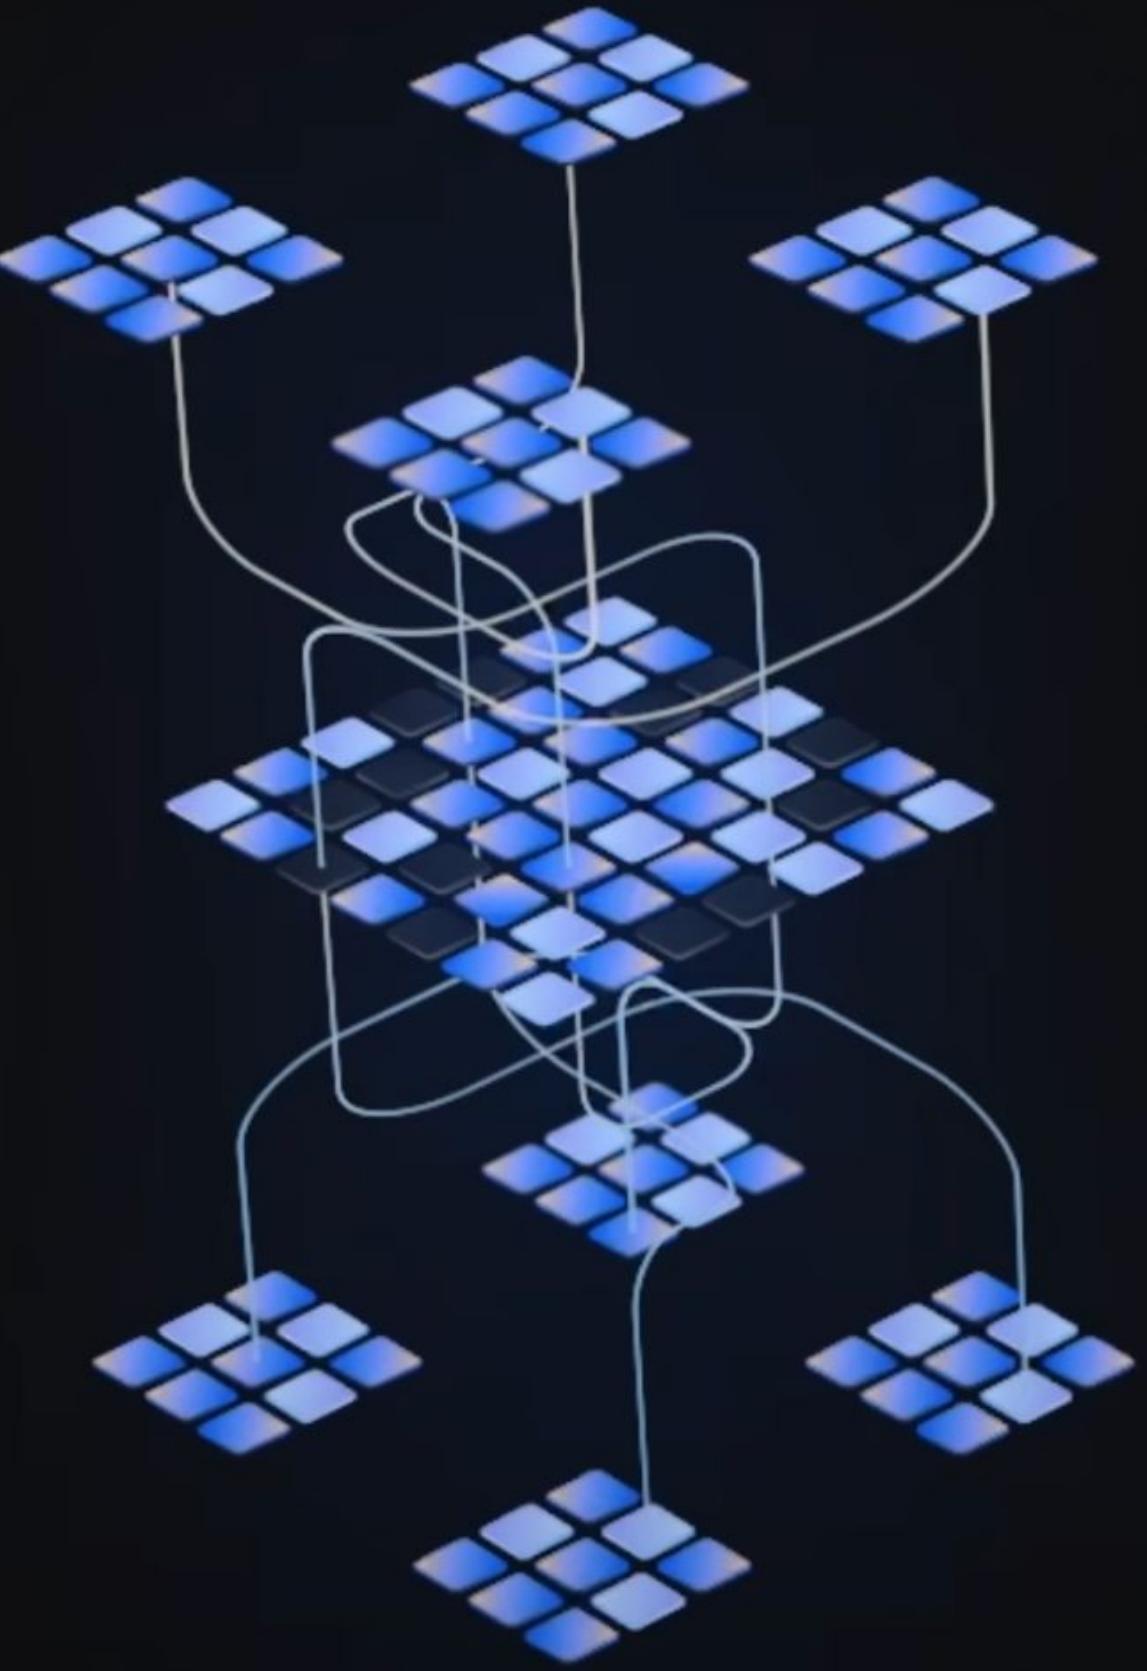
State-of-the-art, natively  
multimodal reasoning  
capabilities



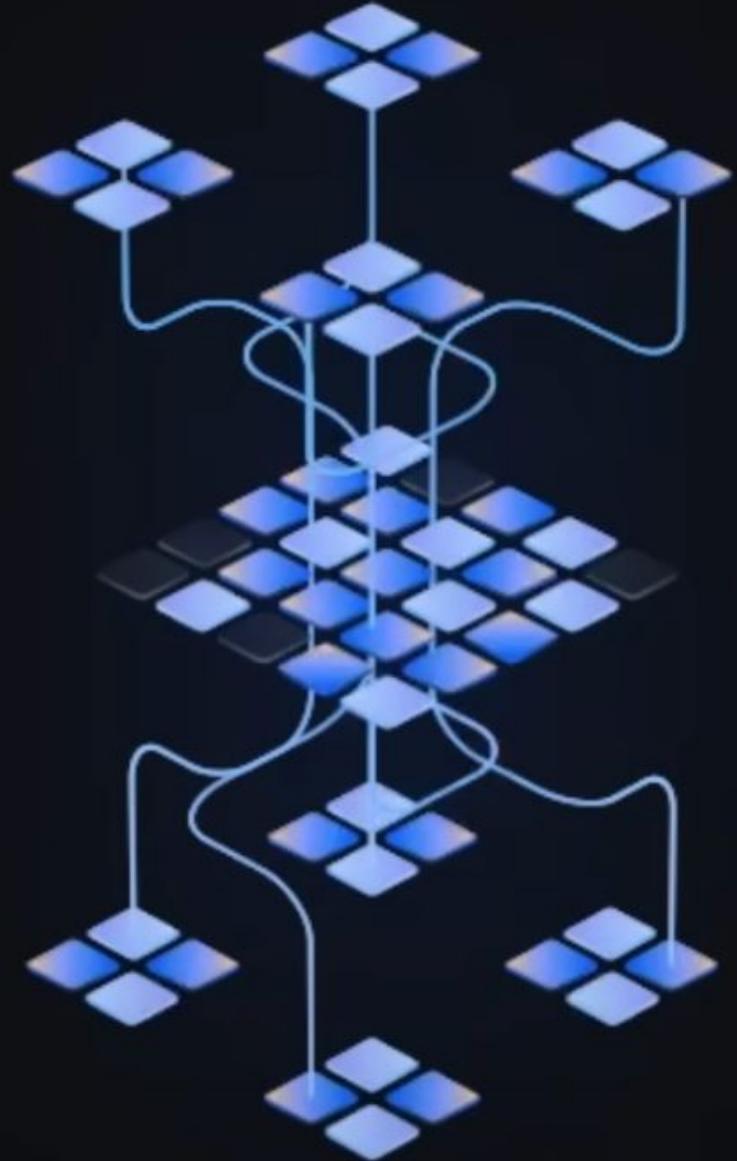
Highly optimized for  
training and inference



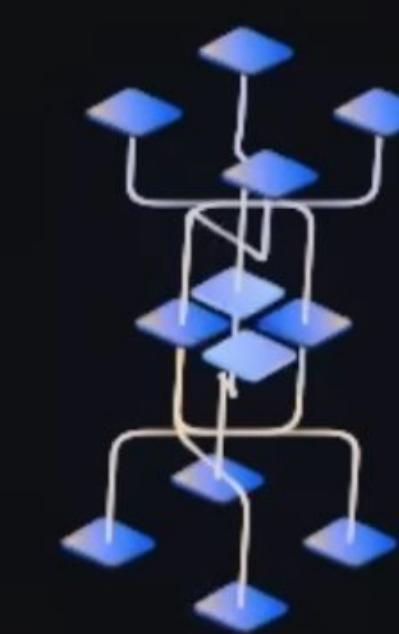
Built with responsibility  
and safety at the core



Ultra



Pro



Nano

(Android AICore)

# Google's Foundation Models on Vertex AI

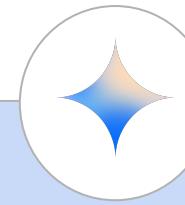
Across a variety of model sizes to address use cases



GA

## Gemini 1.0 Pro

Multimodal reasoning across a wide range of tasks



NEW

## Gemini 1.5 Pro

Multimodal reasoning for longer prompts, 1 million context window



Limited  
Private GA

## Gemini 1.0 Ultra

Largest and most capable model for highly complex tasks



NEW

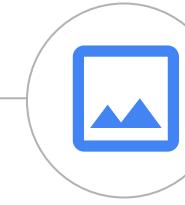
## Gemma 2B and 7B

Family of lightweight, state-of-the-art open models



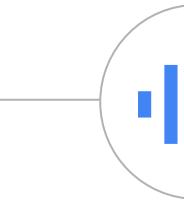
## PaLM for Text / Chat

Custom language tasks and multi-turn conversations



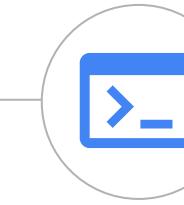
## Imagen 2.0 for Text to Image

Create and edit images from simple prompts



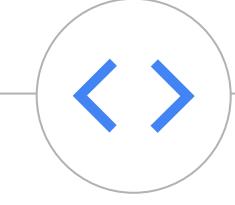
## Chirp for Speech to Text

Build voice enabled applications



## Codey for Code Generation

Improve coding and debugging



## Embeddings API for Text and Image

Extract semantic information from unstructured data



NEW

## Claude on Vertex AI

Claude 2, Instant 1.2, and more



## Open Models on Vertex AI

Mixtral 8x7B, Image Bind, DITO and more



NEW

## Hugging Face Models

Few click deployment to Vertex AI



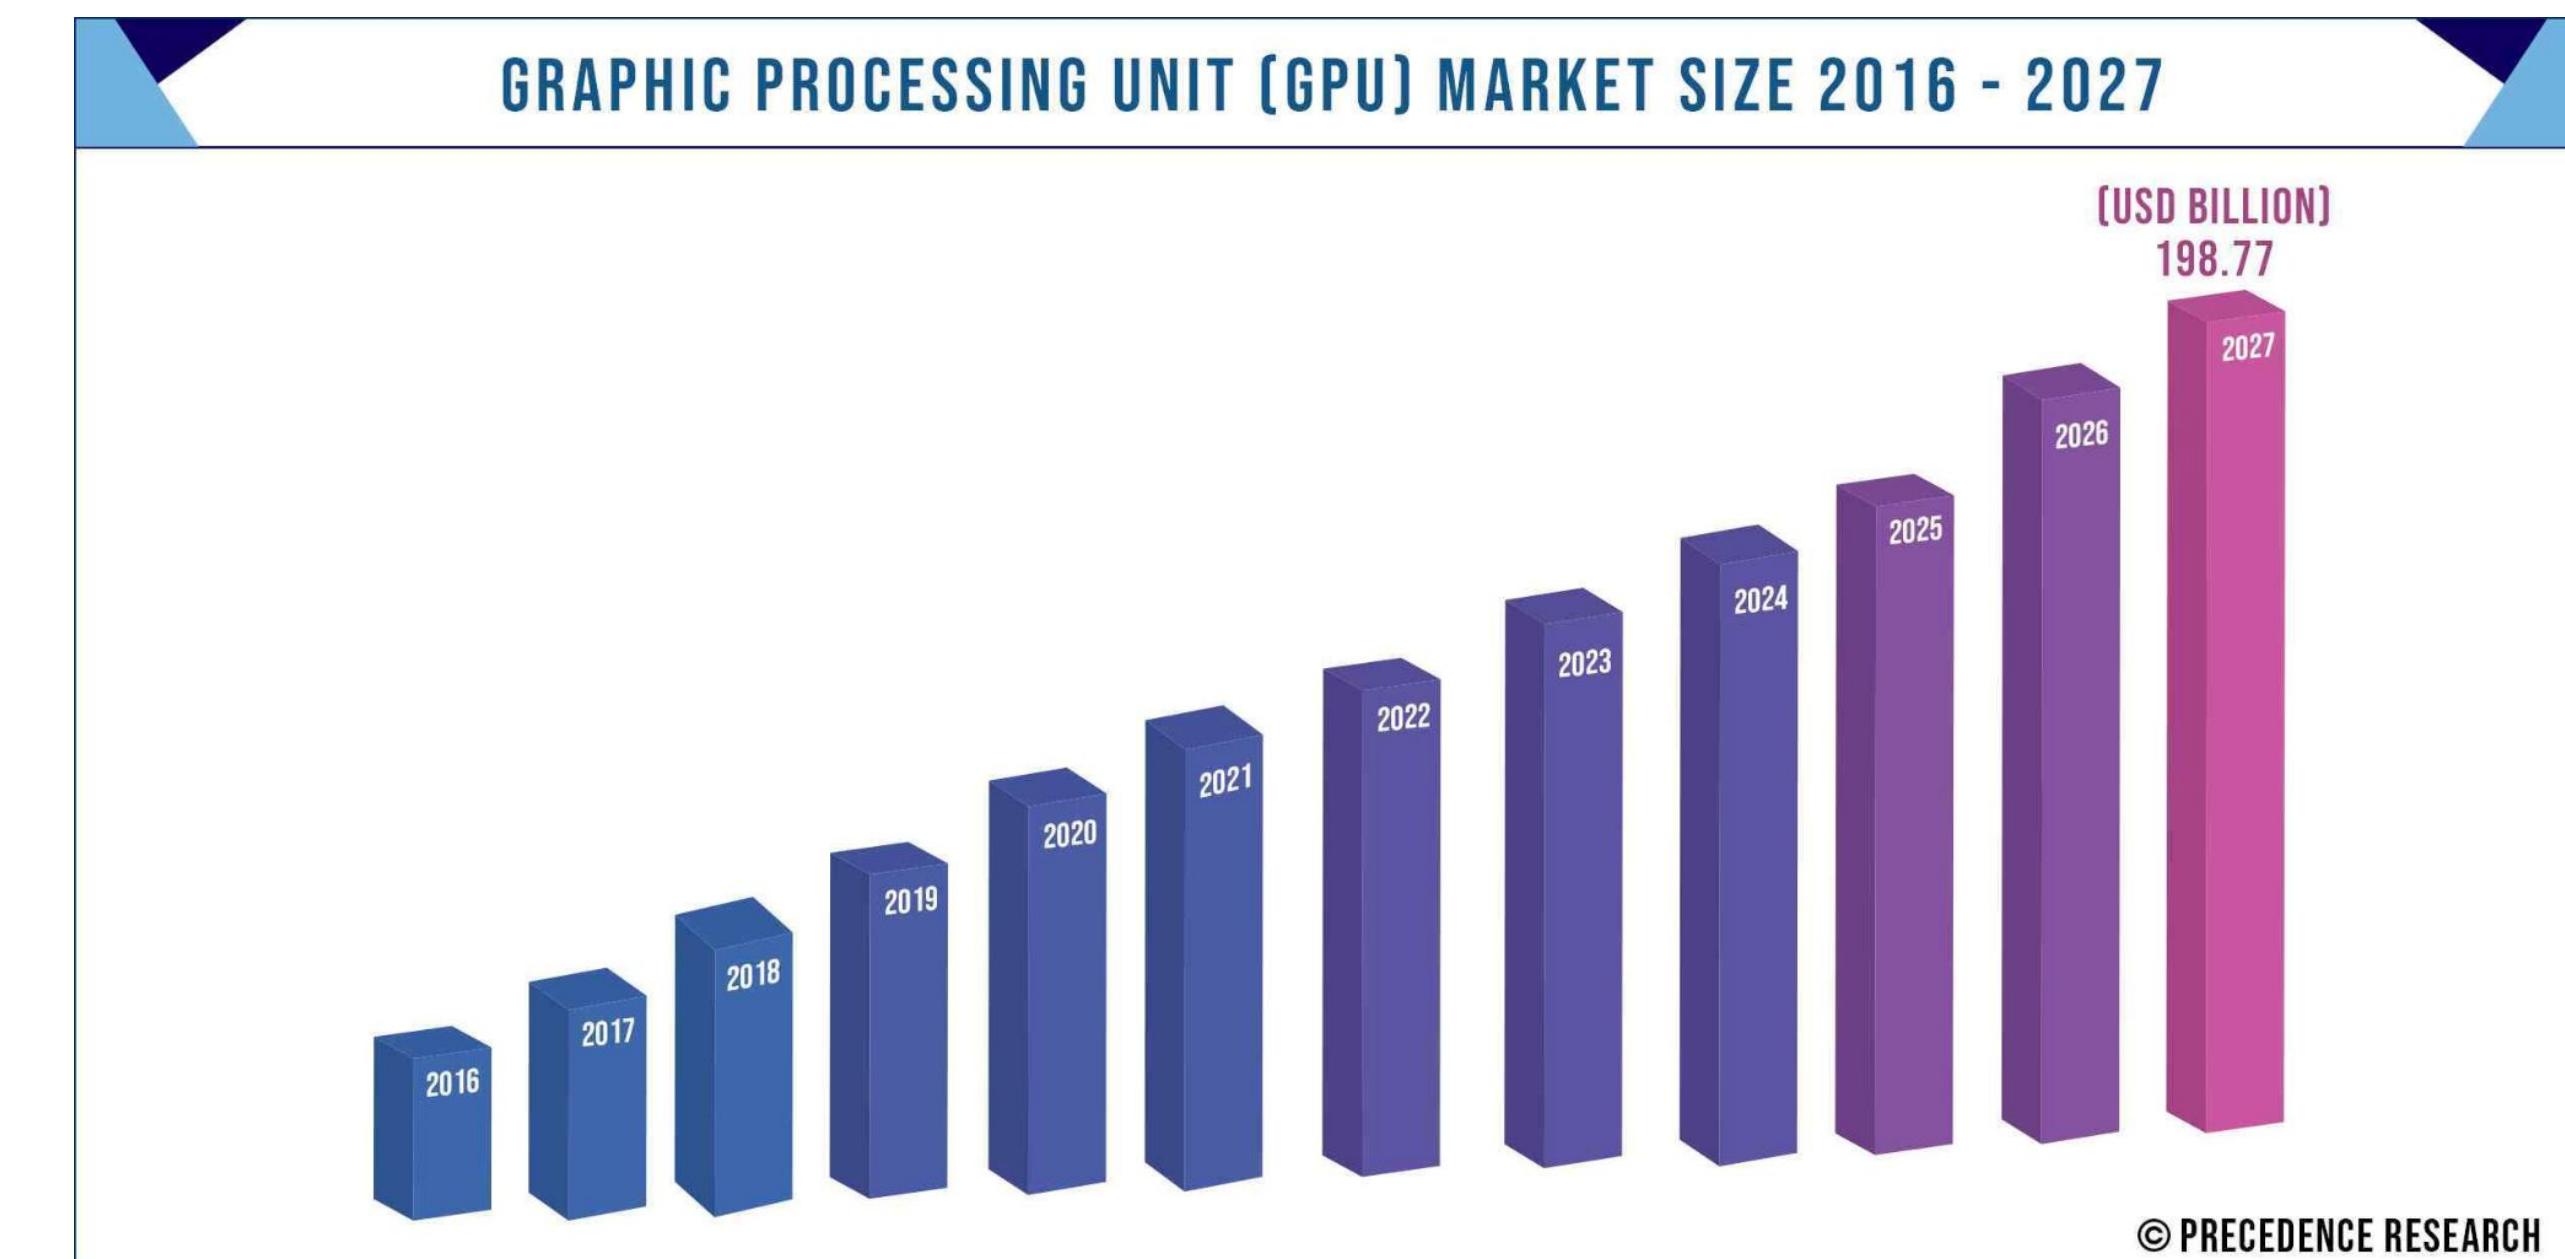


# Gemma Open Models

[ai.google.dev/gemma](https://ai.google.dev/gemma)

# Eight Things to Know about Large Language Models

1. LLMs predictably get more capable with increasing investment, even without targeted innovation.
2. Many important LLM behaviors emerge unpredictably as a byproduct of increasing investment.
3. LLMs often appear to learn and use representations of the outside world.
4. There are no reliable techniques for steering the behavior of LLMs.
5. Experts are not yet able to interpret the inner workings of LLMs.
6. Human performance on a task isn't an upper bound on LLM performance.
7. LLMs need not express the values of their creators nor the values encoded in web text.
8. Brief interactions with LLMs are often misleading

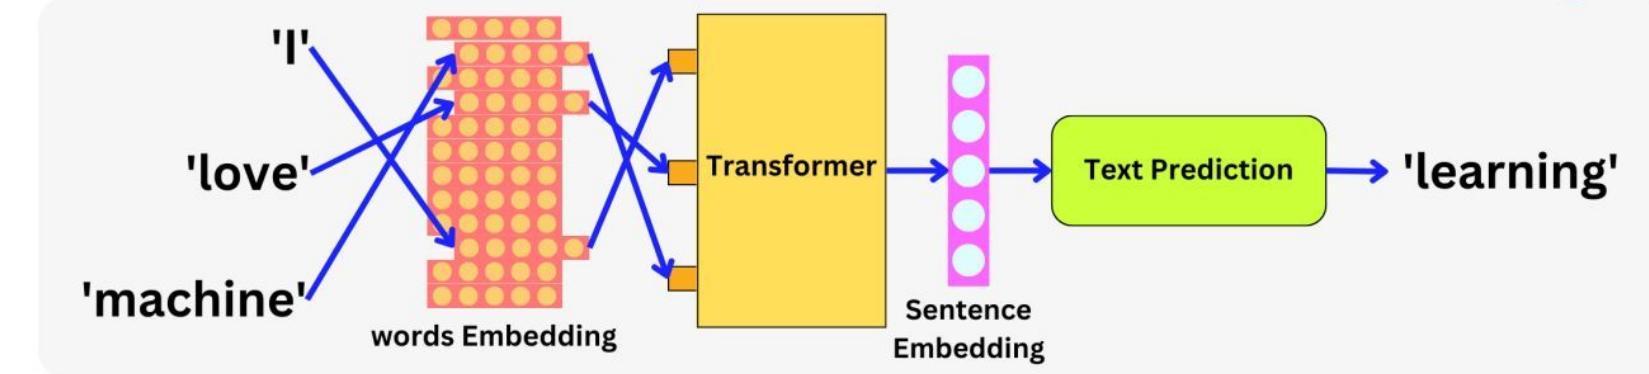


<https://www.precedenceresearch.com/graphic-processing-unit-market>

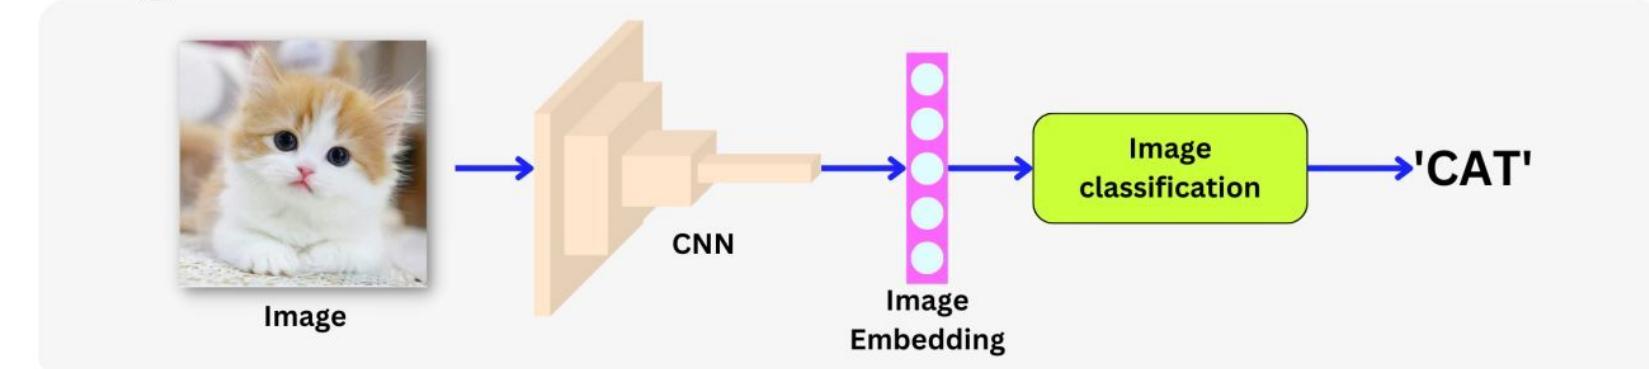
# Modern Deep Learning

## Embeddings: the Superpower of Deep Learning

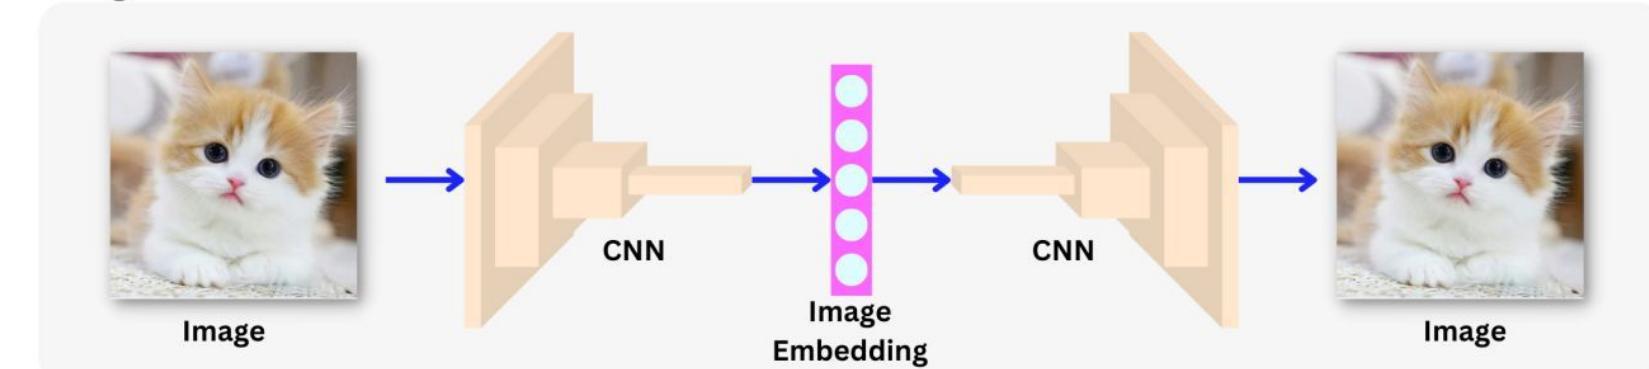
### Text data with Transformers



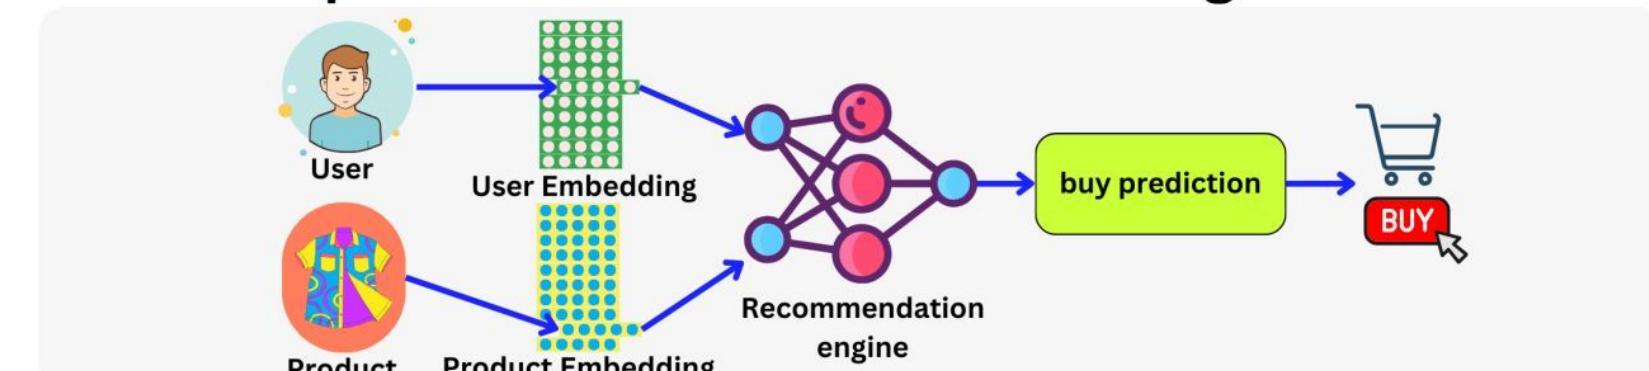
### Image data with ConvNets



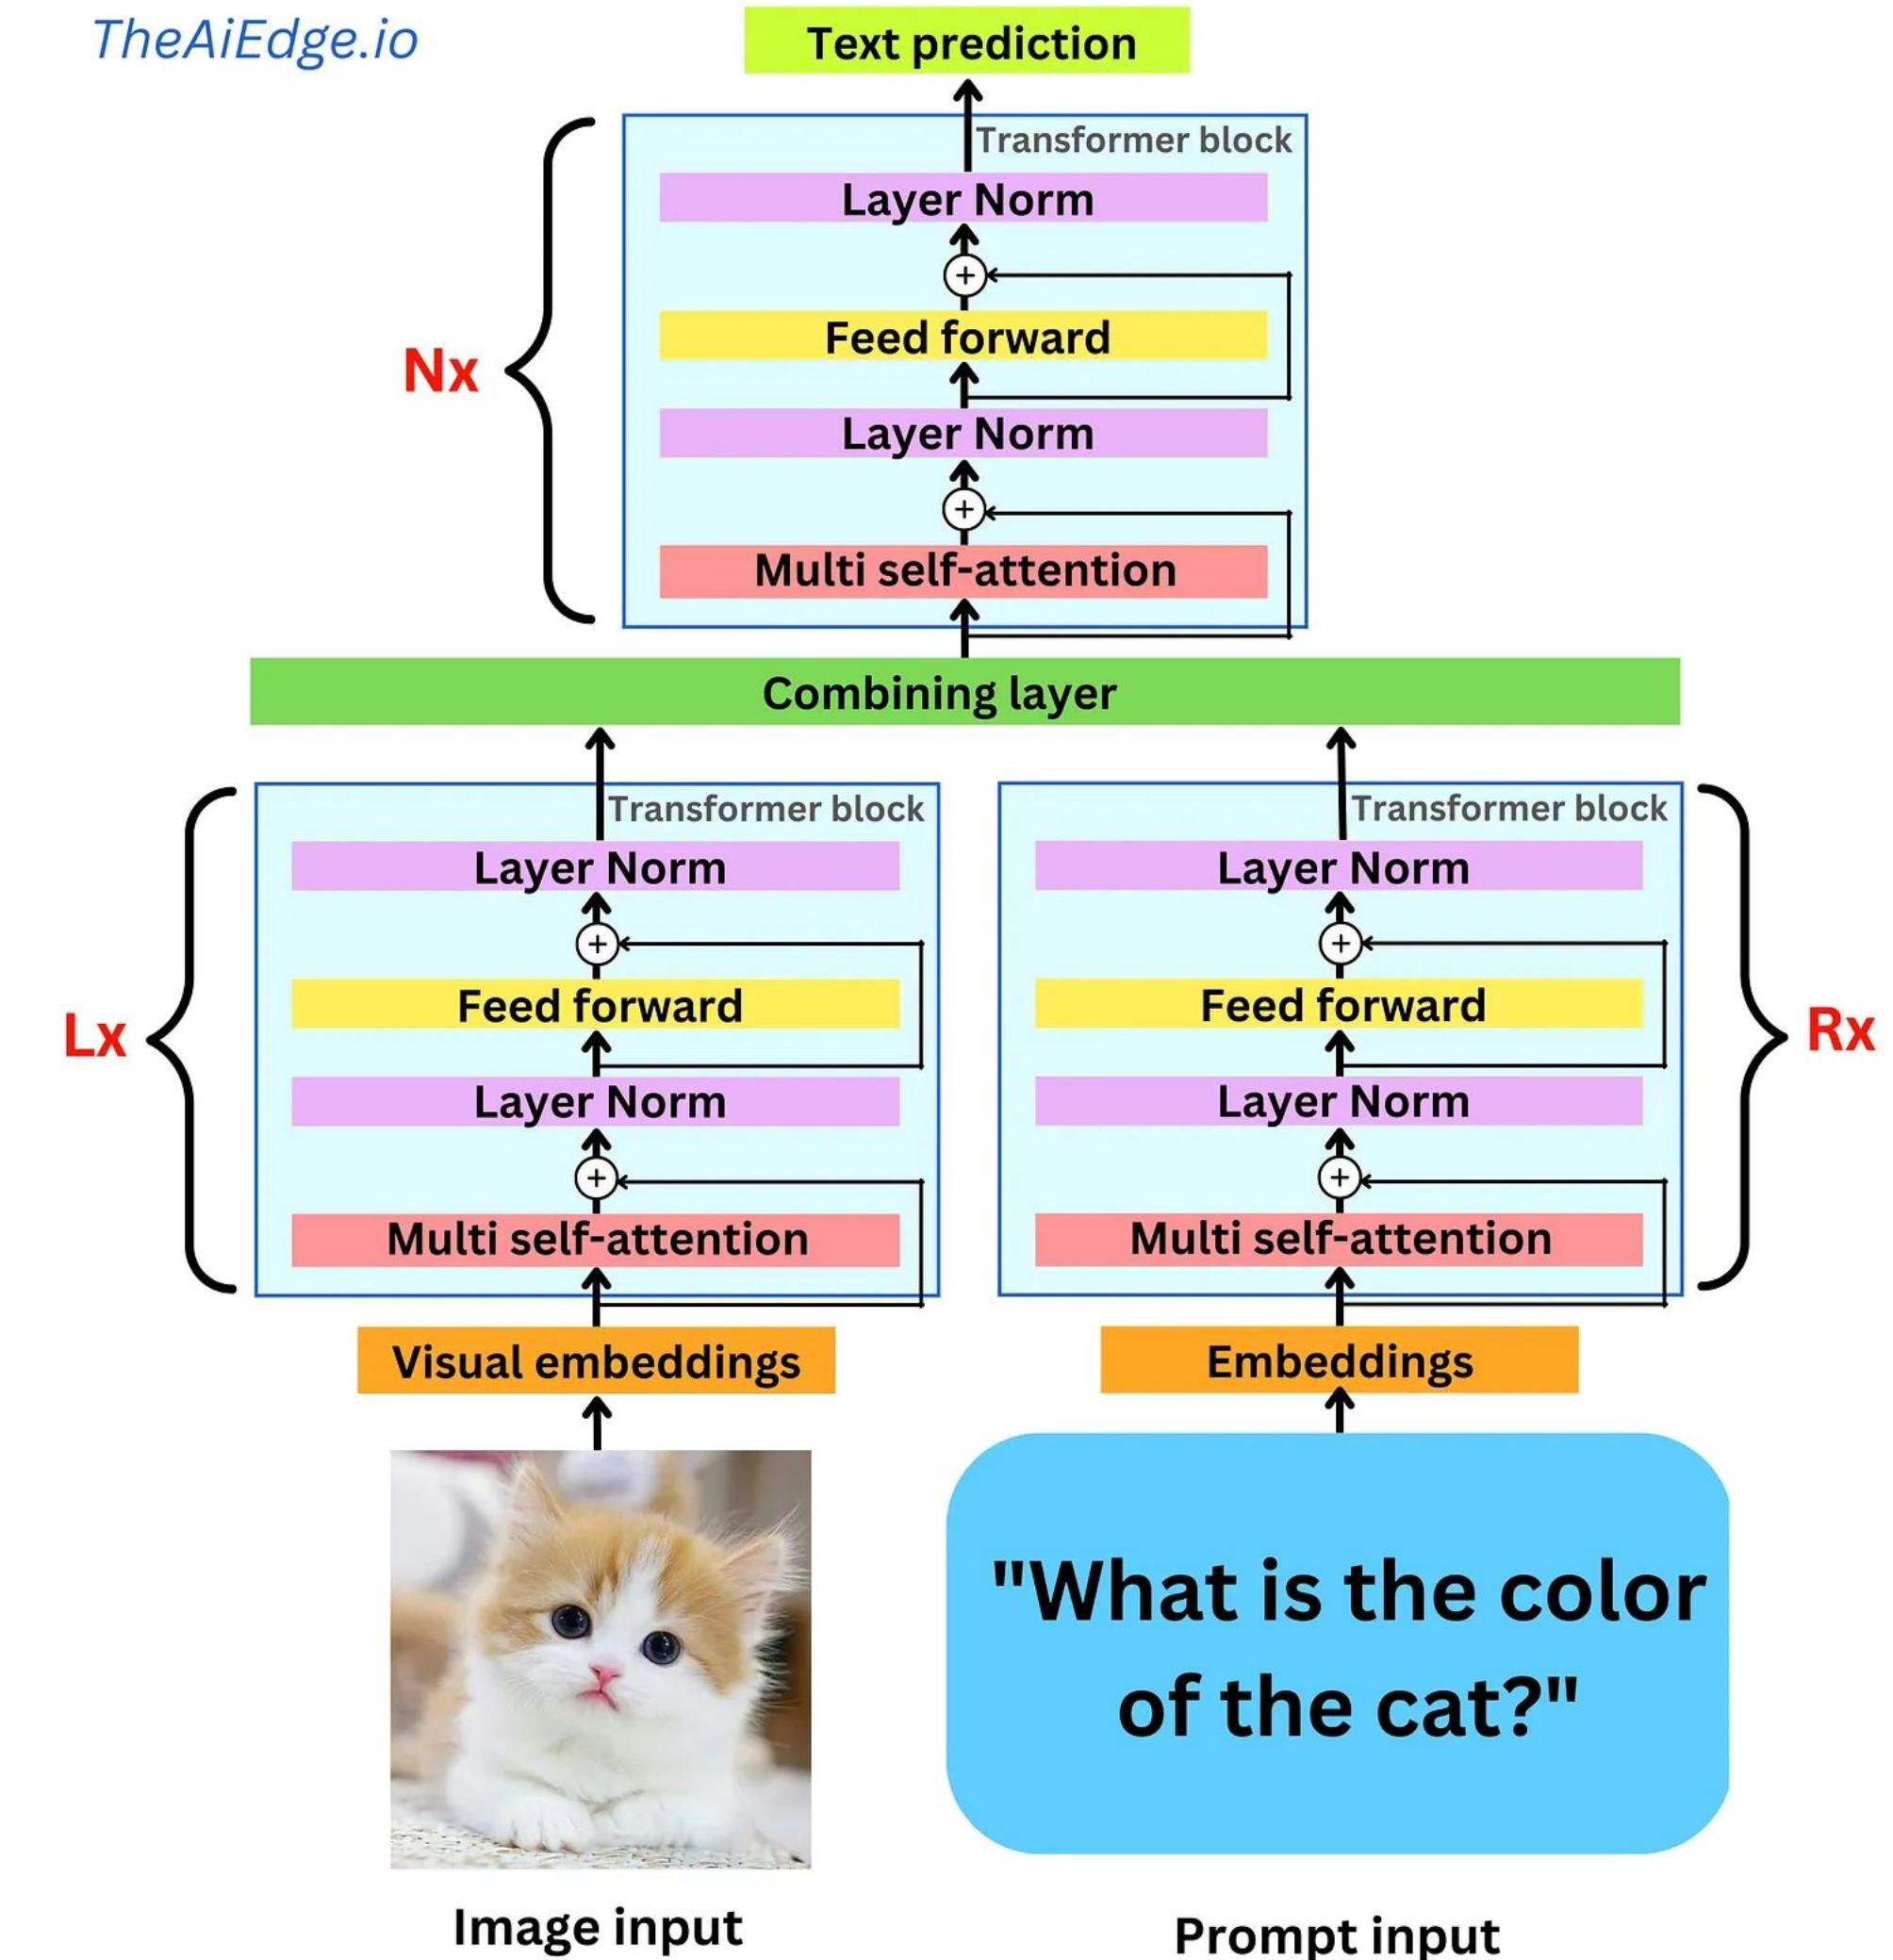
### Any data with variational Auto-encoders



### Users and products with Recommender engines



TheAiEdge.io



# Demo time

